

```
In [1]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os
```

```
In [229]: from google.colab import drive
drive.mount('/content/gdrive')
```

```
-----
ImportError                                Traceback (most recent call last)
<ipython-input-229-6db90e9297aa> in <module>()
----> 1 from google.colab import drive
      2 drive.mount('/content/gdrive')

ImportError: No module named 'google.colab'
```

```

In [230]: #connecting to sqlite db
# con = sqlite3.connect('/content/gdrive/My Drive/Colab Notebooks/Assignment 4/data')
con = sqlite3.connect('database.sqlite')

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 data points
# you can change the number to any other number based on your computing power

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000 """, con)
# for tsne assignment you can take 5k data points

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 """, con)

# Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a negative rating(-1)
def partition(x):
    if x < 3:
        return 0
    else:
        return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)

```

Number of data points in our data (525814, 10)

Out[230]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0
2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1

```
In [231]: display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)
```

```
In [232]: print(display.shape)
display.head()
```

```
(80668, 7)
```

```
Out[232]:
```

	UserId	ProductId	ProfileName	Time	Score	Text	COUNT
0	#oc-R115TNMSPFT9I7	B007Y59HVM	Breyton	1331510400	2	Overall its just OK when considering the price...	2
1	#oc-R11D9D7SHXIJB9	B005HG9ET0	Louis E. Emory "hoppy"	1342396800	5	My wife has recurring extreme muscle spasms, u...	3
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1	This coffee is horrible and unfortunately not ...	2
3	#oc-R11O5J5ZVQE25C	B005HG9ET0	Penguin Chick	1346889600	5	This will be the bottle that you grab from the...	3
4	#oc-R12KPBODL2B5ZD	B007OSBE1U	Christopher P. Presta	1348617600	1	I didnt like this coffee. Instead of telling y...	2

```
In [233]: # Removing duplicate reviews
final=filtered_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"})
print(final.shape)
```

```
(364173, 10)
```

```
In [234]: (final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

```
Out[234]: 69.25890143662969
```



```
100%|██████████████████████████████████████████████████████████████████████████| 364171/364171 [00:00<00:00  
0, 545983.46it/s]
```

```
100%|███████████████████████████████████████████████████████████████████████████
██████████████████████████████████████████████████████████████████████████████| 364171/364171 [02:07<00:
00, 2863.46it/s]
```

364171

```
100%|███████████████████████████████████████████████████████████████████████████  
██████████████████████████████████████████████████████████████████████████████ | 364171/364171 [00:05<00:0  
0, 67496.57it/s]
```

```
In [25]: strip_list = []
         for to_strip in tqdm(decat_lst):
             text = re.sub("\$*\d\$*", "", to_strip).strip()
             strip_list.append(text)
```

```
100%|███████████████████████████████████████████████████████████████████████████  
██████████████████████████████████████████████████████████████████████████████| 364171/364171 [00:19<00:0  
0, 18494.66it/s]
```

```
In [26]: spatial_list = []
         for to_spatial in tqdm(strip_list):
             text = re.sub('[^A-Za-z0-9]+', ' ', to_spatial)
             spatial_list.append(text)
```

```
100%|██████████████████████████████████████████████████████████████████████████|  
██████████████████████████████████████████████████████████████████████████████| 364171/364171 [00:09<00:00  
0, 36464.50it/s]
```

```
In [27]: stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'our  

    'you'll', "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', '  

    'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itsel  

    'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that  

    'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has  

    'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because'  

    'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'th  

    'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off  

    'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all'  

    'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than',  

    's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've  

    've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "di  

    "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma',  

    "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn'  

    'won', "won't", 'wouldn', "wouldn't"])
```

```
In [28]: # Combining all the above students
preprocessed_reviews = []
# tqdm is for printing the status bar
for sentence in tqdm(spatial_list):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
    preprocessed_reviews.append(sentence.strip())
```

```
100%|███████████████████████████████████████████████████████████████████████████
██████████████████████████████████████████████████████████████████████████████| 364171/364171 [02:01<00:
00, 2999.66it/s]
```

```
In [29]: print(len(preprocessed_reviews))
preprocessed_reviews[-1]
```

364171

```
Out[29]: 'great honey satisfied product advertised use cereal raw vinegar general sweetn
er'
```

```
In [30]: final['cleanReview'] = preprocessed_reviews
```

```
In [117]: print(len(final))
final.tail(5)
```

364171

```
Out[117]:
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	He
525809	568450	B001EO7N10	A28KG5XORO54AY	Lettie D. Carter	0	0
525810	568451	B003S1WTCU	A3I8AFVPEE8KI5	R. Sawyer	0	0
525811	568452	B004I613EE	A121AA1GQV751Z	pksd "pk_007"	2	2
525812	568453	B004I613EE	A3IBEVCTXKNOH	Kathy A. Welch "katwel"	1	1
525813	568454	B001LR2CU2	A3LGQPJCZVL9UC	srfell17	0	0

◀ ▶

```
In [118]: final['cleanReview'][0]
```

```
Out[118]: 'good quality dog food bought several vitality canned dog food products found g
ood quality product looks like stew processed meat smells better labrador finic
ky appreciates product better'
```

```
In [119]: final['lengthOfReview'][0]
```

```
Out[119]: 27
```

```
In [2]: dir_path = os.getcwd()
# conn = sqlite3.connect(os.path.join(dir_path, '/content/gdrive/My Drive/Colab N
conn = sqlite3.connect(os.path.join(dir_path, 'final.sqlite'))
# final.to_sql('Reviews', conn, if_exists='replace', index=False)
```

```
In [3]: review_3 = pd.read_sql_query(""" SELECT count(*) FROM Reviews""", conn)
print(review_3)
```

```
count(*)
0      364171
```

```
In [4]: filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews""", conn)
```

```
In [5]: filtered_data.shape
```

```
Out[5]: (364171, 12)
```

```
In [6]: filtered_data["Time"] = pd.to_datetime(filtered_data["Time"], unit = "s")
filtered_data = filtered_data.sort_values(by = "Time")
```


In [7]: `filtered_data.head(5)`

Out[7]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	Hel
117924	150524	0006641040	ACITT7DI6IDDL	shari zychinski	0	0
117901	150501	0006641040	AJ46FKXOVC7NR	Nicholas A Mesiano	2	2
298792	451856	B00004CXX9	AIUWLEQ1ADEG5	Elizabeth Medina	0	0
169281	230285	B00004RYGX	A344SMIA5JECGM	Vincent P. Ross	1	2
298791	451855	B00004CXX9	AJH6LUC1UT1ON	The Phantom of the Opera	0	0



```
In [8]: print(len(filtered_data))
filtered_data.info()
filtered_data = filtered_data.head(100000)
print(len(filtered_data))
```

```
364171
<class 'pandas.core.frame.DataFrame'>
Int64Index: 364171 entries, 117924 to 107253
Data columns (total 12 columns):
Id                364171 non-null int64
ProductId         364171 non-null object
UserId           364171 non-null object
ProfileName       364171 non-null object
HelpfulnessNumerator  364171 non-null int64
HelpfulnessDenominator 364171 non-null int64
Score            364171 non-null int64
Time             364171 non-null datetime64[ns]
Summary          364171 non-null object
Text             364171 non-null object
cleanReview       364171 non-null object
lengthOfReview    364171 non-null int64
dtypes: datetime64[ns](1), int64(5), object(6)
memory usage: 36.1+ MB
100000
```

```
In [9]: filtered_data['Score'].value_counts()
```

```
Out[9]: 1    87729
        0    12271
        Name: Score, dtype: int64
```

```
In [10]: X = filtered_data["cleanReview"]
print(print("shape of X:", X.head(5)))
y = filtered_data["Score"]
print("shape of y:", y.head(5))
X_len = filtered_data['lengthOfReview']
```

```
shape of X: 117924    every book educational witty little book makes...
117901    whole series great way spend time child rememb...
298792    entertainingl funny beetlejuice well written m...
169281    modern day fairy tale twist rumplestiskin capt...
298791    fantastic beetlejuice excellent funny movie ke...
Name: cleanReview, dtype: object
None
shape of y: 117924    1
117901    1
298792    1
169281    1
298791    1
Name: Score, dtype: int64
```

```
In [11]: len(filtered_data['lengthOfReview'])
```

```
Out[11]: 100000
```

```
In [12]: X_train = X[0:100000]
        Y_train = y[0:100000]
```

```
In [13]: print(len(X_train))
        print(len(Y_train))
```

```
100000
100000
```

```
In [14]: print(X_train.shape)
        print(X_train.shape)
```

```
(100000,)
(100000,)
```

Bag of Words

```
In [15]: from sklearn.feature_extraction.text import CountVectorizer
```

```
count_vect = CountVectorizer()
X_train_vect = count_vect.fit_transform(X_train)
X_test_vect = count_vect.transform(X_test)
X_val_vect = count_vect.transform(X_val)
feature_names = count_vect.get_feature_names()
# Bow_dict = {'X_train_vect': X_train_vect, 'X_test_vect': X_test_vect, 'X_val_vect': X_val_vect}
print(X_train_vect.shape)
# print(feature_names)
```

```
(60000, 48270)
```

```
In [16]: X_train_vect.shape
```

```
Out[16]: (60000, 48270)
```

```
In [18]: len(filtered_data['lengthOfReview'])
```

```
Out[18]: 100000
```

```
In [191]: from scipy.sparse import hstack
        # len_review = final['lengthOfReview'].to_sparse()
        concat_data = hstack((X_train_vect, np.array(filtered_data['lengthOfReview'])[0:60000]))
        concat_data_val = hstack((X_val_vect, np.array(filtered_data['lengthOfReview'])[60000:80000]))
        concat_data_test = hstack((X_test_vect, np.array(filtered_data['lengthOfReview'])[80000:100000]))
```

```
In [193]: print(concat_data.shape)
        print(concat_data_val.shape)
        print(concat_data_test.shape)
```

```
(60000, 47536)
(20000, 47536)
(20000, 47536)
```

```
In [194]: print(len(feature_names))
```

```
47535
```

```
In [195]: BoW_dict = {'X_train_vect':concat_data, 'X_test_vect': concat_data_test, 'X_val_v
print(BoW_dict['X_train_vect'].shape)
```

```
(60000, 47536)
```

```
In [196]: import pickle
with open('BoW.pkl', 'wb') as handle:
    pickle.dump(BoW_dict, handle, protocol=pickle.HIGHEST_PROTOCOL)
```

TF-IDF

```
In [15]: tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10)
train_tf_idf = tf_idf_vect.fit_transform(X_train)

print("the shape of out text TFIDF vectorizer ",train_tf_idf.get_shape())
print("the type of count vectorizer ",type(train_tf_idf))
print("the number of unique words including both unigrams and bigrams ", train_tf
```

```
the shape of out text TFIDF vectorizer (100000, 59901)
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the number of unique words including both unigrams and bigrams 59901
```

```
In [210]: tfidf_concat_data_train = hstack((train_tf_idf,np.array(final['lengthOfReview'])[0
tfidf_concat_data_val = hstack((cv_tf_idf,np.array(final['lengthOfReview'])[60000:
tfidf_concat_data_test = hstack((test_tf_idf,np.array(final['lengthOfReview'])[800
```

```
In [211]: tf_idf_dict = {'train_tf_idf': tfidf_concat_data_train, 'cv_tf_idf': tfidf_concat
```



```
In [16]: tf_idf_dict = {'train_tf_idf': train_tf_idf}
```

```
In [17]: import pickle
with open('tf_idf.pkl', 'wb') as handle:
    pickle.dump(tf_idf_dict, handle, protocol=pickle.HIGHEST_PROTOCOL)
```

TruncatedSVD on tfidf

```
In [18]: import pickle
# with open(r"/content/gdrive/My Drive/Colab Notebooks/Assignment 4/tf_idf.pkl",
with open(r"tf_idf.pkl", "rb") as input_file:
    tfidf_dict = pickle.load(input_file)
```

```
In [19]: from sklearn.decomposition import TruncatedSVD
```

```
svd = TruncatedSVD()  
svd.fit(train_tf_idf)
```

```
Out[19]: TruncatedSVD(algorithm='randomized', n_components=2, n_iter=5,  
                    random_state=None, tol=0.0)
```

```
In [20]: idf = tf_idf_vect.idf_  
y = dict(zip(tf_idf_vect.get_feature_names(), idf))
```

```
In [21]: from collections import OrderedDict  
from operator import itemgetter  
  
sorted_dict = OrderedDict(sorted(y.items(), key = itemgetter(1), reverse=True))
```

```
In [22]: top_features = list(sorted_dict)[:2000]
```

```
In [23]: print(len(top_features))  
top_features[0]
```

```
2000
```

```
Out[23]: 'jerky lovers'
```

```
In [24]: #https://codereview.stackexchange.com/questions/188482/generating-a-word-bigram-co-occurrence-matrix

import nltk
from nltk import bigrams

def co_occurrence_matrix(corpus):
    vocab = set(corpus)
    vocab = list(vocab)

    # Key:Value = Word:Index
    vocab_to_index = { word:i for i, word in enumerate(vocab) }

    # Create bigrams from all words in corpus
    bi_grams = list(bigrams(corpus))

    # Frequency distribution of bigrams ((word1, word2), num_occurrences)
    bigram_freq = nltk.FreqDist(bi_grams).most_common(len(bi_grams))

    # Initialise co-occurrence matrix
    # co_occurrence_matrix[current][previous]
    co_occurrence_matrix = np.zeros((len(vocab), len(vocab)))

    # Loop through the bigrams in the frequency distribution, noting the
    # current and previous word, and the number of occurrences of the bigram.
    # Get the vocab index of the current and previous words.
    # Put the number of occurrences into the appropriate element of the array.
    for bigram in bigram_freq:
        current = bigram[0][1]
        previous = bigram[0][0]
        count = bigram[1]
        pos_current = vocab_to_index[current]
        pos_previous = vocab_to_index[previous]
        co_occurrence_matrix[pos_current][pos_previous] = count

    co_occurrence_matrix = np.matrix(co_occurrence_matrix)

    return co_occurrence_matrix

m = co_occurrence_matrix(top_features)
```

```
In [25]: print(m)
m.shape
```

```
[[0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 ...
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 1. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
```

```
Out[25]: (2000, 2000)
```

```
In [26]: type(train_tf_idf)
```

```
Out[26]: scipy.sparse.csr.csr_matrix
```


```
In [27]: from sklearn.decomposition import PCA
from sklearn.preprocessing import MinMaxScaler
```

```
In [85]: #https://towardsdatascience.com/an-approach-to-choosing-the-number-of-components-
#Fitting the Truncated SVD algorithm with our Data
tsvd = TruncatedSVD(n_components=50).fit(m)
```

```
In [86]: tsvd.explained_variance_ratio_
```

```
Out[86]: array([0.00050035, 0.00049886, 0.00050025, 0.00049901, 0.00050005 ,
0.00050034, 0.00050009, 0.00050044, 0.00050049, 0.00050001 ,
0.00050048, 0.00050005 , 0.00050046, 0.00049928, 0.00049991,
0.00050001 , 0.00050003, 0.00050045, 0.00050049, 0.00050034,
0.00049951, 0.00050048, 0.00050047, 0.00050014, 0.00050049,
0.00050048, 0.00050019, 0.00050013, 0.00049987, 0.00050045,
0.00050048, 0.00050005 , 0.00050023, 0.00050046, 0.00050039,
0.00050038, 0.00050012, 0.00050041, 0.00049976, 0.00050033,
0.00050045, 0.00050037, 0.00050025, 0.00050045, 0.00050028,
0.00050035, 0.00050005 , 0.00050032, 0.00049854, 0.00049977])
```

```
In [87]: #Plotting the Cumulative Summation of the Explained Variance
plt.figure()
plt.plot(np.cumsum(tsvd.explained_variance_ratio_))
plt.xlabel('Number of Components')
plt.ylabel('Variance (%)') #for each component
plt.title('Amazon Dataset Explained Variance')
plt.show()
```



```
In [31]: best_svd = TruncatedSVD(n_components=1).fit(train_tf_idf)
```

```
In [ ]: from sklearn.cluster import KMeans
from sklearn.model_selection import GridSearchCV
from tqdm import tqdm
tfidf_k_inertia_train = dict()

for k_val in range(1, 8):
    tfidf_km_clf = KMeans(n_clusters = k_val, n_jobs = -1)
    tfidf_km_clf.fit(train_tf_idf)
    tfidf_k_inertia_train[k_val] = (tfidf_km_clf.inertia_)
```


```
In [33]: tfidf_k_inertia_train
```

```
Out[33]: {1: 98878.82946187379,
2: 98455.39946212494,
3: 98166.69384298568,
4: 97982.62814136724,
5: 97801.75092993751,
6: 97675.24596462022,
7: 97551.88000632105}
```

```
In [34]: with open('cluster_dict.pkl', 'wb') as handle:
    pickle.dump(tfidf_k_inertia_train, handle, protocol=pickle.HIGHEST_PROTOCOL)
```



```
In [35]: plt.figure()
plt.plot(list(tfidf_k_inertia_train.keys()), list(tfidf_k_inertia_train.values()))
plt.scatter(list(tfidf_k_inertia_train.keys()), list(tfidf_k_inertia_train.values))
plt.xlabel("Number of cluster")
plt.ylabel("Inertia")
plt.grid()
plt.show()
```



Training with 7 clusters

```
In [ ]: import datetime

t1 = datetime.datetime.now()
tfidf_final_clf = KMeans(n_clusters = 7, n_jobs = -1)
tfidf_final_clf.fit(train_tf_idf)
print("time required = ",datetime.datetime.now() - t1 )

time required = 0:48:57.536801
```

```
In [44]: from wordcloud import WordCloud
imp_features_tfidf = []
print("Top terms per cluster:")
order_centroids = tfidf_final_clf.cluster_centers_.argsort()[:, :-1]
terms = tf_idf_vect.get_feature_names()
print(len(terms))
for i in range(7):
    for ind in order_centroids[i, :20]:
        imp_features_dict = {}
        imp_features_dict[i] = terms[ind-1]
        imp_features_tfidf.append(imp_features_dict)
```

Top terms per cluster:
59901

```
In [45]: len(imp_features_tfidf)
```

Out[45]: 140

```
In [53]: cl1 = [d[0] for d in imp_features_tfidf if 0 in d]
cl2 = [d[1] for d in imp_features_tfidf if 1 in d]
cl3 = [d[2] for d in imp_features_tfidf if 2 in d]
cl4 = [d[3] for d in imp_features_tfidf if 3 in d]
cl5 = [d[4] for d in imp_features_tfidf if 4 in d]
cl6 = [d[5] for d in imp_features_tfidf if 5 in d]
cl7 = [d[6] for d in imp_features_tfidf if 6 in d]
print(cl1, cl2, cl3, cl4, cl5, cl6)
```

['doesnt taste', 'doggy', 'fong sriracha', 'treatments', 'casual', 'catnip', 'l
overs try', 'nostalgic', 'treasures', 'lousy', 'dog lover', 'dogs little', 'omi
tted', 'greasy salty', 'dog favorite', 'easy yummy', 'lightweight', 'worthy',
'vessel', 'gesture'] ['glutamic acid', 'gluten flour', 'fred', 'brazilian', 'mi
tigate', 'nostalgic', 'cookie would', 'greasy salty', 'besides great', 'past ye
ars', 'whatsoever', 'goo', 'best glad', 'flossies', 'tast', 'lightweight', 'gev
alia signature', 'mak', 'celestial seasons', 'easy yummy'] ['chocolat', 'nostal
gic', 'darjeeling teas', 'dark brown', 'hot chili', 'hosting', 'banned', 'milit
ary', 'lightweight', 'coco', 'goo', 'besides great', 'tast', 'barry tea', 'cook
ie would', 'greasy salty', 'flavonoids', 'lousy', 'delicioso', 'sweep'] ['nosta
lgic', 'goo', 'lightweight', 'tast', 'greasy salty', 'flavonoids', 'lousy', 'be
sides great', 'omitted', 'worthy', 'realizing', 'delicioso', 'nj', 'sufficientl
y', 'litters', 'usda organic', 'sweep', 'gesture', 'easy yummy', 'mak'] ['produ
cing', 'greasy salty', 'amazingly tasty', 'prey', 'nostalgic', 'financial', 'sh
ipper', 'goo', 'orchards', 'buttons', 'storage not', 'great problem', 'excellen
ce', 'stored refrigerator', 'lousy', 'tim tams', 'grocers', 'lobster meat', 'wo
rthy', 'besides great'] ['coffe', 'cumin', 'nostalgic', 'greasy salty', 'goo',
'flavonoids', 'stroller', 'cup cocoa', 'lightweight', 'roaring', 'boils', 'besi
des great', 'pod system', 'great coconut', 'tast', 'espressione', 'omitted', 's
ense smell', 'darjeeling teas', 'cuppa tea']

```
In [54]: cl1_string = ' '.join(cl1)
cl2_string = ' '.join(cl2)
cl3_string = ' '.join(cl3)
cl4_string = ' '.join(cl4)
cl5_string = ' '.join(cl5)
cl6_string = ' '.join(cl6)
cl7_string = ' '.join(cl7)
```


```
In [55]: from wordcloud import WordCloud
wordcloud_cluster1 = WordCloud(width = 500, height = 500, background_color = 'white')
wordcloud_cluster2 = WordCloud(width = 500, height = 500, background_color = 'white')
wordcloud_cluster3 = WordCloud(width = 500, height = 500, background_color = 'white')
wordcloud_cluster4 = WordCloud(width = 500, height = 500, background_color = 'white')
wordcloud_cluster5 = WordCloud(width = 500, height = 500, background_color = 'white')
wordcloud_cluster6 = WordCloud(width = 500, height = 500, background_color = 'white')
wordcloud_cluster7 = WordCloud(width = 500, height = 500, background_color = 'white')
# plot the WordCloud image
plt.imshow(wordcloud_cluster1)
plt.tight_layout(pad = 0)
plt.show()
```




```
In [56]: plt.imshow(wordcloud_cluster2)  
plt.tight_layout(pad = 0)  
plt.show()
```



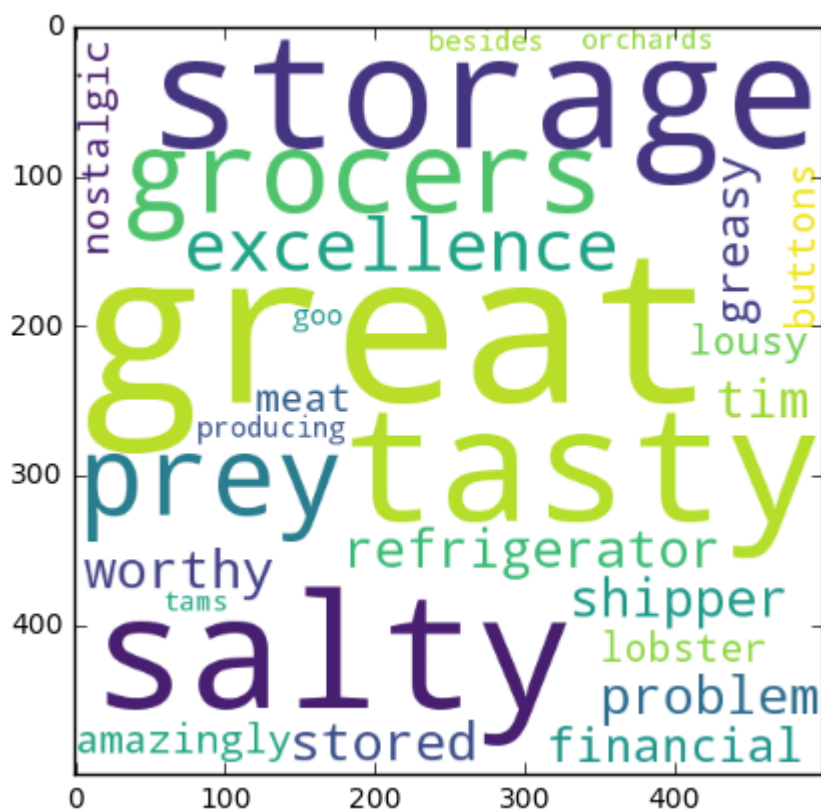
```
In [57]: plt.imshow(wordcloud_cluster3)
plt.tight_layout(pad = 0)
plt.show()
```




```
plt.imshow(wordcloud_cluster4)
plt.tight_layout(pad = 0)
plt.show()
```




```
In [59]: plt.imshow(wordcloud_cluster5)  
plt.tight_layout(pad = 0)  
plt.show()
```



```
plt.imshow(wordcloud_cluster6)
plt.tight_layout(pad = 0)
plt.show()
```




```
In [61]: plt.imshow(wordcloud_cluster7)
plt.tight_layout(pad = 0)
plt.show()
```



```
In [62]: #function that takes a word and returns the most similar words using cosine similarity
def cosine_similarity(a,b):
    return dot(a,b) / ( (dot(a,a) **.5) * (dot(b,b) **.5) )
```

Observations

cluster 1 represents related to dog food eg. dog, doggy, yummy
cluster 2 represents type of food eg. cookie, salty, brazilian
cluster 3 represents taste of food eg. salty, chili, sweet, coco
cluster 4 represents organic food eg. organic, great, worthy
cluster 5 represents shipping related information eg. shipping, financial, refrigerator
cluster 6 represents taste of coffee eg. coffee, cocoa, smell, stronger
cluster 7 represents tea related information eg. tea, green, tasty