

Keras -- MLPs on MNIST

```
In [2]: # if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow"
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

Using TensorFlow backend.

```
In [3]: %matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

```
In [4]: # the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz> (<https://s3.amazonaws.com/img-datasets/mnist.npz>)
 11493376/11490434 [=====] - 5s 0us/step

```
In [5]: print("Number of training examples :", X_train.shape[0], "and each image is of shape", X_train.shape[1]*X_train.shape[2])
print("Number of training examples :", X_test.shape[0], "and each image is of shape", X_test.shape[1]*X_test.shape[2])
```

Number of training examples : 60000 and each image is of shape (28, 28)
 Number of training examples : 10000 and each image is of shape (28, 28)

```
In [6]: # if you observe the input shape its 3 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

```
In [7]: # after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape", X_train.shape[1])
print("Number of training examples :", X_test.shape[0], "and each image is of shape", X_test.shape[1])
```

Number of training examples : 60000 and each image is of shape (784)
 Number of training examples : 10000 and each image is of shape (784)

```
In [8]: # An example data point
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18  126  136  175  26  166  255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94  154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  249 253 249  64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  23  66 213 253 253 253 253 198  81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  18 171 219 253 253 253 253 195
 80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 55 172 226 253 253 253 253 244 133  11  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  136 253 253 253 212 135 132  16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0]
```

```
In [9]: # if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 
```

```
X_train = X_train/255
X_test = X_test/255
```

```
In [10]: # example data point after normlizing
print(X_train[0])
```

```
[0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.]
```

```
In [11]: # here we are having a class number for each image
print("Class label of first image :", y_train[0])

# Lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

Class label of first image : 5

After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

Softmax classifier

```

In [12]: # https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

####

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias)
# activation is the element-wise activation function passed as the activation argument
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True)

# output = activation(dot(input, kernel) + bias) => y = activation(WT.X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation Layer, or through the activation argument of a layer

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation

```

```
In [18]: # some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

```
In [19]: print(X_train.shape[1])
```

784

MLP + ReLU + ADAM with 2 layers without Dropout and Batch Normalisation

```
In [20]: model_relu = Sequential()
model_relu.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_relu.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epochs)
```

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 364)	285740
dense_5 (Dense)	(None, 52)	18980
dense_6 (Dense)	(None, 10)	530

=====
 Total params: 305,250
 Trainable params: 305,250
 Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20
 60000/60000 [=====] - 11s 180us/step - loss: 0.2688 - acc: 0.9216 - val_loss: 0.1258 - val_acc: 0.9609

Epoch 2/20
 60000/60000 [=====] - 8s 137us/step - loss: 0.1032 - acc: 0.9697 - val_loss: 0.0881 - val_acc: 0.9721

Epoch 3/20
 60000/60000 [=====] - 8s 126us/step - loss: 0.0660 - acc: 0.9800 - val_loss: 0.0776 - val_acc: 0.9750

Epoch 4/20
 60000/60000 [=====] - 8s 128us/step - loss: 0.0476 - acc: 0.9859 - val_loss: 0.0919 - val_acc: 0.9724

Epoch 5/20
 60000/60000 [=====] - 8s 127us/step - loss: 0.0331 - acc: 0.9893 - val_loss: 0.0683 - val_acc: 0.9802

Epoch 6/20
 60000/60000 [=====] - 8s 133us/step - loss: 0.0244 - acc: 0.9926 - val_loss: 0.0721 - val_acc: 0.9792

Epoch 7/20
 60000/60000 [=====] - 8s 140us/step - loss: 0.0201 - acc: 0.9938 - val_loss: 0.0777 - val_acc: 0.9787

Epoch 8/20
 60000/60000 [=====] - 8s 132us/step - loss: 0.0152 - acc: 0.9957 - val_loss: 0.0777 - val_acc: 0.9803

Epoch 9/20
 60000/60000 [=====] - 9s 142us/step - loss: 0.0130 - acc: 0.9958 - val_loss: 0.0799 - val_acc: 0.9778

Epoch 10/20
 60000/60000 [=====] - 10s 170us/step - loss: 0.0138 - acc: 0.9954 - val_loss: 0.0830 - val_acc: 0.9779

```
Epoch 11/20
60000/60000 [=====] - 10s 165us/step - loss: 0.0126 -
acc: 0.9956 - val_loss: 0.0758 - val_acc: 0.9814
Epoch 12/20
60000/60000 [=====] - 10s 168us/step - loss: 0.0104 -
acc: 0.9967 - val_loss: 0.0785 - val_acc: 0.9816
Epoch 13/20
60000/60000 [=====] - 10s 161us/step - loss: 0.0090 -
acc: 0.9970 - val_loss: 0.0907 - val_acc: 0.9801
Epoch 14/20
60000/60000 [=====] - 10s 160us/step - loss: 0.0059 -
acc: 0.9984 - val_loss: 0.1034 - val_acc: 0.9770
Epoch 15/20
60000/60000 [=====] - 10s 162us/step - loss: 0.0106 -
acc: 0.9966 - val_loss: 0.1226 - val_acc: 0.9741
Epoch 16/20
60000/60000 [=====] - 10s 163us/step - loss: 0.0099 -
acc: 0.9965 - val_loss: 0.0945 - val_acc: 0.9796
Epoch 17/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0052 - a
cc: 0.9983 - val_loss: 0.0903 - val_acc: 0.9823
Epoch 18/20
60000/60000 [=====] - 10s 159us/step - loss: 0.0049 -
acc: 0.9986 - val_loss: 0.0967 - val_acc: 0.9813
Epoch 19/20
60000/60000 [=====] - 9s 158us/step - loss: 0.0053 - a
cc: 0.9985 - val_loss: 0.1076 - val_acc: 0.9792
Epoch 20/20
60000/60000 [=====] - 10s 159us/step - loss: 0.0085 -
acc: 0.9971 - val_loss: 0.1080 - val_acc: 0.9790
```

```

In [21]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epochs)

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.10797379065210243

Test accuracy: 0.979




```

In [23]: w_after = model_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



C:\Program Files\Anaconda3\lib\site-packages\scipy\stats\stats.py:1626: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error

nor or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

MLP + Batch-Norm on 2 hidden Layers + AdamOptimizer

```
In [26]: # Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condition
#  $h1 \Rightarrow \sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(\theta, \sigma) = N(\theta, 0.039)$ 
#  $h2 \Rightarrow \sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(\theta, \sigma) = N(\theta, 0.055)$ 
#  $h1 \Rightarrow \sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(\theta, \sigma) = N(\theta, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0, std=0.039)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0, std=0.055)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Layer (type)	Output Shape	Param #
dense_11 (Dense)	(None, 364)	285740
batch_normalization_3 (Batch Normalization)	(None, 364)	1456
dense_12 (Dense)	(None, 52)	18980
batch_normalization_4 (Batch Normalization)	(None, 52)	208
dense_13 (Dense)	(None, 10)	530
Total params: 306,914		
Trainable params: 306,082		
Non-trainable params: 832		

```
In [27]: model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoc

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 13s 219us/step - loss: 0.2254 -
acc: 0.9364 - val_loss: 0.1244 - val_acc: 0.9634
Epoch 2/20
60000/60000 [=====] - 10s 163us/step - loss: 0.0861 -
acc: 0.9747 - val_loss: 0.0873 - val_acc: 0.9741
Epoch 3/20
60000/60000 [=====] - 11s 182us/step - loss: 0.0549 -
acc: 0.9836 - val_loss: 0.0879 - val_acc: 0.9731
Epoch 4/20
60000/60000 [=====] - 8s 141us/step - loss: 0.0409 - a
cc: 0.9876 - val_loss: 0.0817 - val_acc: 0.9756
Epoch 5/20
60000/60000 [=====] - 8s 138us/step - loss: 0.0312 - a
cc: 0.9903 - val_loss: 0.0816 - val_acc: 0.9761
Epoch 6/20
60000/60000 [=====] - 8s 139us/step - loss: 0.0247 - a
cc: 0.9924 - val_loss: 0.0798 - val_acc: 0.9749
Epoch 7/20
60000/60000 [=====] - 8s 139us/step - loss: 0.0211 - a
cc: 0.9933 - val_loss: 0.0911 - val_acc: 0.9756
Epoch 8/20
60000/60000 [=====] - 8s 139us/step - loss: 0.0176 - a
cc: 0.9941 - val_loss: 0.0896 - val_acc: 0.9746
Epoch 9/20
60000/60000 [=====] - 9s 144us/step - loss: 0.0167 - a
cc: 0.9949 - val_loss: 0.0852 - val_acc: 0.9769
Epoch 10/20
60000/60000 [=====] - 9s 143us/step - loss: 0.0139 - a
cc: 0.9955 - val_loss: 0.0793 - val_acc: 0.9768
Epoch 11/20
60000/60000 [=====] - 8s 140us/step - loss: 0.0123 - a
cc: 0.9960 - val_loss: 0.0973 - val_acc: 0.9749
Epoch 12/20
60000/60000 [=====] - 8s 140us/step - loss: 0.0103 - a
cc: 0.9970 - val_loss: 0.0956 - val_acc: 0.9766
Epoch 13/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0090 - a
cc: 0.9972 - val_loss: 0.0840 - val_acc: 0.9782
Epoch 14/20
60000/60000 [=====] - 11s 183us/step - loss: 0.0094 -
acc: 0.9969 - val_loss: 0.0880 - val_acc: 0.9771
Epoch 15/20
60000/60000 [=====] - 11s 179us/step - loss: 0.0113 -
acc: 0.9960 - val_loss: 0.0901 - val_acc: 0.9781
Epoch 16/20
60000/60000 [=====] - 11s 176us/step - loss: 0.0088 -
acc: 0.9976 - val_loss: 0.0810 - val_acc: 0.9780
Epoch 17/20
60000/60000 [=====] - 11s 179us/step - loss: 0.0092 -
acc: 0.9970 - val_loss: 0.0838 - val_acc: 0.9776
Epoch 18/20
```

```
60000/60000 [=====] - 11s 183us/step - loss: 0.0069 -  
acc: 0.9979 - val_loss: 0.0957 - val_acc: 0.9776  
Epoch 19/20  
60000/60000 [=====] - 11s 187us/step - loss: 0.0062 -  
acc: 0.9980 - val_loss: 0.0830 - val_acc: 0.9778  
Epoch 20/20  
60000/60000 [=====] - 11s 183us/step - loss: 0.0076 -  
acc: 0.9974 - val_loss: 0.0870 - val_acc: 0.9798
```

```

In [28]: score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch)

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.08703972299850357

Test accuracy: 0.9798



```

In [29]: w_after = model_batch.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



C:\Program Files\Anaconda3\lib\site-packages\scipy\stats\stats.py:1626: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error

ror or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

MLP + Dropout + AdamOptimizer with 2 hidden layers

In [35]: *# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormaliza*

```
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_20 (Dense)	(None, 364)	285740
dropout_5 (Dropout)	(None, 364)	0
dense_21 (Dense)	(None, 52)	18980
dropout_6 (Dropout)	(None, 52)	0
dense_22 (Dense)	(None, 10)	530

=====
 Total params: 305,250
 Trainable params: 305,250
 Non-trainable params: 0


```
In [36]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['a
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 11s 175us/step - loss: 0.9381 -
acc: 0.7072 - val_loss: 0.2438 - val_acc: 0.9337
Epoch 2/20
60000/60000 [=====] - 10s 161us/step - loss: 0.4452 -
acc: 0.8702 - val_loss: 0.1901 - val_acc: 0.9459
Epoch 3/20
60000/60000 [=====] - 11s 179us/step - loss: 0.3585 -
acc: 0.8967 - val_loss: 0.1622 - val_acc: 0.9531
Epoch 4/20
60000/60000 [=====] - 10s 174us/step - loss: 0.3074 -
acc: 0.9121 - val_loss: 0.1551 - val_acc: 0.9569
Epoch 5/20
60000/60000 [=====] - 10s 173us/step - loss: 0.2725 -
acc: 0.9233 - val_loss: 0.1363 - val_acc: 0.9606
Epoch 6/20
60000/60000 [=====] - 10s 169us/step - loss: 0.2420 -
acc: 0.9331 - val_loss: 0.1302 - val_acc: 0.9638
Epoch 7/20
60000/60000 [=====] - 11s 181us/step - loss: 0.2228 -
acc: 0.9382 - val_loss: 0.1279 - val_acc: 0.9649
Epoch 8/20
60000/60000 [=====] - 11s 181us/step - loss: 0.2064 -
acc: 0.9423 - val_loss: 0.1191 - val_acc: 0.9656
Epoch 9/20
60000/60000 [=====] - 11s 191us/step - loss: 0.1984 -
acc: 0.9451 - val_loss: 0.1094 - val_acc: 0.9698
Epoch 10/20
60000/60000 [=====] - 12s 193us/step - loss: 0.1887 -
acc: 0.9477 - val_loss: 0.1118 - val_acc: 0.9687
Epoch 11/20
60000/60000 [=====] - 11s 182us/step - loss: 0.1756 -
acc: 0.9519 - val_loss: 0.1076 - val_acc: 0.9705
Epoch 12/20
60000/60000 [=====] - 11s 180us/step - loss: 0.1716 -
acc: 0.9527 - val_loss: 0.1048 - val_acc: 0.9720
Epoch 13/20
60000/60000 [=====] - 11s 184us/step - loss: 0.1570 -
acc: 0.9562 - val_loss: 0.0986 - val_acc: 0.9726
Epoch 14/20
60000/60000 [=====] - 10s 159us/step - loss: 0.1540 -
acc: 0.9575 - val_loss: 0.1034 - val_acc: 0.9731
Epoch 15/20
60000/60000 [=====] - 9s 152us/step - loss: 0.1478 - a
cc: 0.9589 - val_loss: 0.1007 - val_acc: 0.9727
Epoch 16/20
60000/60000 [=====] - 9s 153us/step - loss: 0.1406 - a
cc: 0.9611 - val_loss: 0.0963 - val_acc: 0.9740
Epoch 17/20
60000/60000 [=====] - 9s 158us/step - loss: 0.1377 - a
cc: 0.9629 - val_loss: 0.0984 - val_acc: 0.9743
Epoch 18/20
```

```
60000/60000 [=====] - 10s 159us/step - loss: 0.1328 -  
acc: 0.9622 - val_loss: 0.1004 - val_acc: 0.9746  
Epoch 19/20  
60000/60000 [=====] - 10s 158us/step - loss: 0.1289 -  
acc: 0.9637 - val_loss: 0.1005 - val_acc: 0.9753  
Epoch 20/20  
60000/60000 [=====] - 10s 159us/step - loss: 0.1230 -  
acc: 0.9650 - val_loss: 0.1064 - val_acc: 0.9744
```

```

In [37]: score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch)

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.106384716608499

Test accuracy: 0.9744



```

In [38]: w_after = model_drop.get_weights()

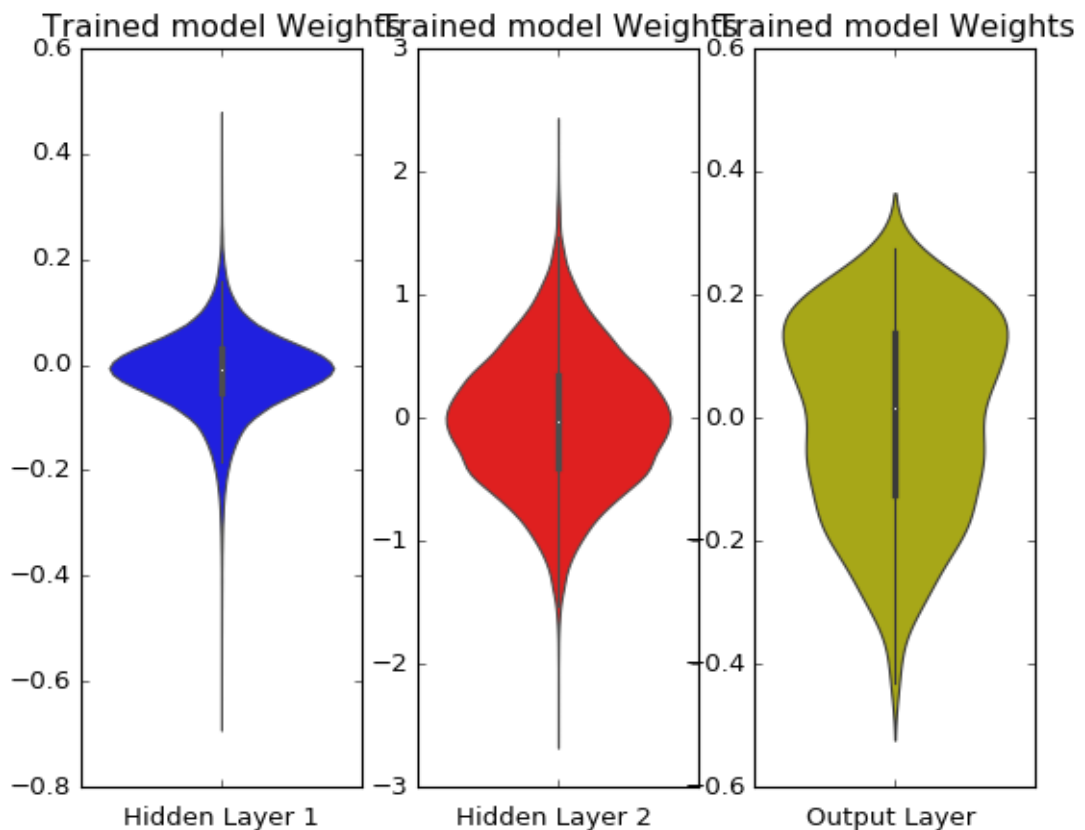
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



C:\Program Files\Anaconda3\lib\site-packages\scipy\stats\stats.py:1626: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error

ror or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

MLP + BatchNormalization + Dropout + AdamOptimizer with 2 hidden layers

```
In [34]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormaliza

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_17 (Dense)	(None, 364)	285740
dropout_3 (Dropout)	(None, 364)	0
dense_18 (Dense)	(None, 52)	18980
dropout_4 (Dropout)	(None, 52)	0
dense_19 (Dense)	(None, 10)	530
=====	=====	=====
Total params: 305,250		
Trainable params: 305,250		
Non-trainable params: 0		

```
In [31]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['a
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 14s 228us/step - loss: 0.5853 -
acc: 0.8241 - val_loss: 0.1880 - val_acc: 0.9449
Epoch 2/20
60000/60000 [=====] - 11s 188us/step - loss: 0.3184 -
acc: 0.9081 - val_loss: 0.1446 - val_acc: 0.9551
Epoch 3/20
60000/60000 [=====] - 11s 178us/step - loss: 0.2581 -
acc: 0.9246 - val_loss: 0.1272 - val_acc: 0.9601
Epoch 4/20
60000/60000 [=====] - 11s 182us/step - loss: 0.2201 -
acc: 0.9359 - val_loss: 0.1106 - val_acc: 0.9650
Epoch 5/20
60000/60000 [=====] - 11s 176us/step - loss: 0.1999 -
acc: 0.9418 - val_loss: 0.0952 - val_acc: 0.9705
Epoch 6/20
60000/60000 [=====] - 11s 176us/step - loss: 0.1816 -
acc: 0.9477 - val_loss: 0.0952 - val_acc: 0.9706
Epoch 7/20
60000/60000 [=====] - 12s 207us/step - loss: 0.1663 -
acc: 0.9518 - val_loss: 0.0877 - val_acc: 0.9737
Epoch 8/20
60000/60000 [=====] - 12s 207us/step - loss: 0.1548 -
acc: 0.9540 - val_loss: 0.0847 - val_acc: 0.9727
Epoch 9/20
60000/60000 [=====] - 13s 216us/step - loss: 0.1463 -
acc: 0.9576 - val_loss: 0.0868 - val_acc: 0.9749
Epoch 10/20
60000/60000 [=====] - 10s 166us/step - loss: 0.1379 -
acc: 0.9588 - val_loss: 0.0833 - val_acc: 0.9750
Epoch 11/20
60000/60000 [=====] - 9s 158us/step - loss: 0.1301 - a
cc: 0.9624 - val_loss: 0.0784 - val_acc: 0.9766
Epoch 12/20
60000/60000 [=====] - 10s 169us/step - loss: 0.1260 -
acc: 0.9633 - val_loss: 0.0753 - val_acc: 0.9774
Epoch 13/20
60000/60000 [=====] - 10s 172us/step - loss: 0.1242 -
acc: 0.9637 - val_loss: 0.0737 - val_acc: 0.9787
Epoch 14/20
60000/60000 [=====] - 12s 194us/step - loss: 0.1136 -
acc: 0.9663 - val_loss: 0.0753 - val_acc: 0.9788
Epoch 15/20
60000/60000 [=====] - 12s 195us/step - loss: 0.1106 -
acc: 0.9675 - val_loss: 0.0766 - val_acc: 0.9791
Epoch 16/20
60000/60000 [=====] - 11s 189us/step - loss: 0.1061 -
acc: 0.9687 - val_loss: 0.0718 - val_acc: 0.9800
Epoch 17/20
60000/60000 [=====] - 11s 186us/step - loss: 0.1067 -
acc: 0.9676 - val_loss: 0.0714 - val_acc: 0.9795
Epoch 18/20
```

```
60000/60000 [=====] - 11s 186us/step - loss: 0.0991 -  
acc: 0.9708 - val_loss: 0.0681 - val_acc: 0.9800  
Epoch 19/20  
60000/60000 [=====] - 11s 190us/step - loss: 0.0999 -  
acc: 0.9705 - val_loss: 0.0729 - val_acc: 0.9797  
Epoch 20/20  
60000/60000 [=====] - 11s 190us/step - loss: 0.0945 -  
acc: 0.9721 - val_loss: 0.0676 - val_acc: 0.9808
```

```

In [32]: score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch)

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.06761971237978433

Test accuracy: 0.9808




```

In [33]: w_after = model_drop.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



C:\Program Files\Anaconda3\lib\site-packages\scipy\stats\stats.py:1626: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error

nor or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

MLP + ReLU + ADAM with 3 layers without Dropout and Batch Normalisation

```
In [55]: model_relu = Sequential()
model_relu.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_relu.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epochs)
```

Layer (type)	Output Shape	Param #
dense_35 (Dense)	(None, 364)	285740
dense_36 (Dense)	(None, 128)	46720
dense_37 (Dense)	(None, 52)	6708
dense_38 (Dense)	(None, 10)	530

=====
 Total params: 339,698
 Trainable params: 339,698
 Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 14s 236us/step - loss: 0.2435 - acc: 0.9263 - val_loss: 0.1237 - val_acc: 0.9618

Epoch 2/20

60000/60000 [=====] - 11s 183us/step - loss: 0.0923 - acc: 0.9721 - val_loss: 0.0868 - val_acc: 0.9727

Epoch 3/20

60000/60000 [=====] - 11s 182us/step - loss: 0.0601 - acc: 0.9819 - val_loss: 0.0794 - val_acc: 0.9760

Epoch 4/20

60000/60000 [=====] - 12s 204us/step - loss: 0.0412 - acc: 0.9870 - val_loss: 0.0709 - val_acc: 0.9785

Epoch 5/20

60000/60000 [=====] - 11s 190us/step - loss: 0.0318 - acc: 0.9901 - val_loss: 0.0782 - val_acc: 0.9753

Epoch 6/20

60000/60000 [=====] - 11s 183us/step - loss: 0.0255 - acc: 0.9914 - val_loss: 0.0837 - val_acc: 0.9768

Epoch 7/20

60000/60000 [=====] - 11s 182us/step - loss: 0.0218 - acc: 0.9930 - val_loss: 0.0668 - val_acc: 0.9804

Epoch 8/20

60000/60000 [=====] - 11s 188us/step - loss: 0.0208 - acc: 0.9930 - val_loss: 0.0912 - val_acc: 0.9766

Epoch 9/20

60000/60000 [=====] - 11s 184us/step - loss: 0.0141 - acc: 0.9955 - val_loss: 0.0870 - val_acc: 0.9775

```
Epoch 10/20
60000/60000 [=====] - 12s 201us/step - loss: 0.0179 -
acc: 0.9938 - val_loss: 0.0842 - val_acc: 0.9783
Epoch 11/20
60000/60000 [=====] - 12s 196us/step - loss: 0.0120 -
acc: 0.9962 - val_loss: 0.0815 - val_acc: 0.9794
Epoch 12/20
60000/60000 [=====] - 11s 184us/step - loss: 0.0132 -
acc: 0.9958 - val_loss: 0.1108 - val_acc: 0.9761
Epoch 13/20
60000/60000 [=====] - 11s 183us/step - loss: 0.0171 -
acc: 0.9945 - val_loss: 0.0812 - val_acc: 0.9820
Epoch 14/20
60000/60000 [=====] - 11s 191us/step - loss: 0.0106 -
acc: 0.9967 - val_loss: 0.0926 - val_acc: 0.9809
Epoch 15/20
60000/60000 [=====] - 11s 190us/step - loss: 0.0094 -
acc: 0.9969 - val_loss: 0.0898 - val_acc: 0.9808
Epoch 16/20
60000/60000 [=====] - 11s 189us/step - loss: 0.0088 -
acc: 0.9970 - val_loss: 0.1011 - val_acc: 0.9790
Epoch 17/20
60000/60000 [=====] - 11s 188us/step - loss: 0.0109 -
acc: 0.9965 - val_loss: 0.0928 - val_acc: 0.9819
Epoch 18/20
60000/60000 [=====] - 11s 187us/step - loss: 0.0103 -
acc: 0.9963 - val_loss: 0.0966 - val_acc: 0.9814
Epoch 19/20
60000/60000 [=====] - 12s 198us/step - loss: 0.0099 -
acc: 0.9966 - val_loss: 0.1025 - val_acc: 0.9804
Epoch 20/20
60000/60000 [=====] - 11s 191us/step - loss: 0.0090 -
acc: 0.9972 - val_loss: 0.1190 - val_acc: 0.9787
```

```

In [56]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epochs)

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.11902276019332958

Test accuracy: 0.9787




```
In [60]: w_after = model_relu.get_weights()

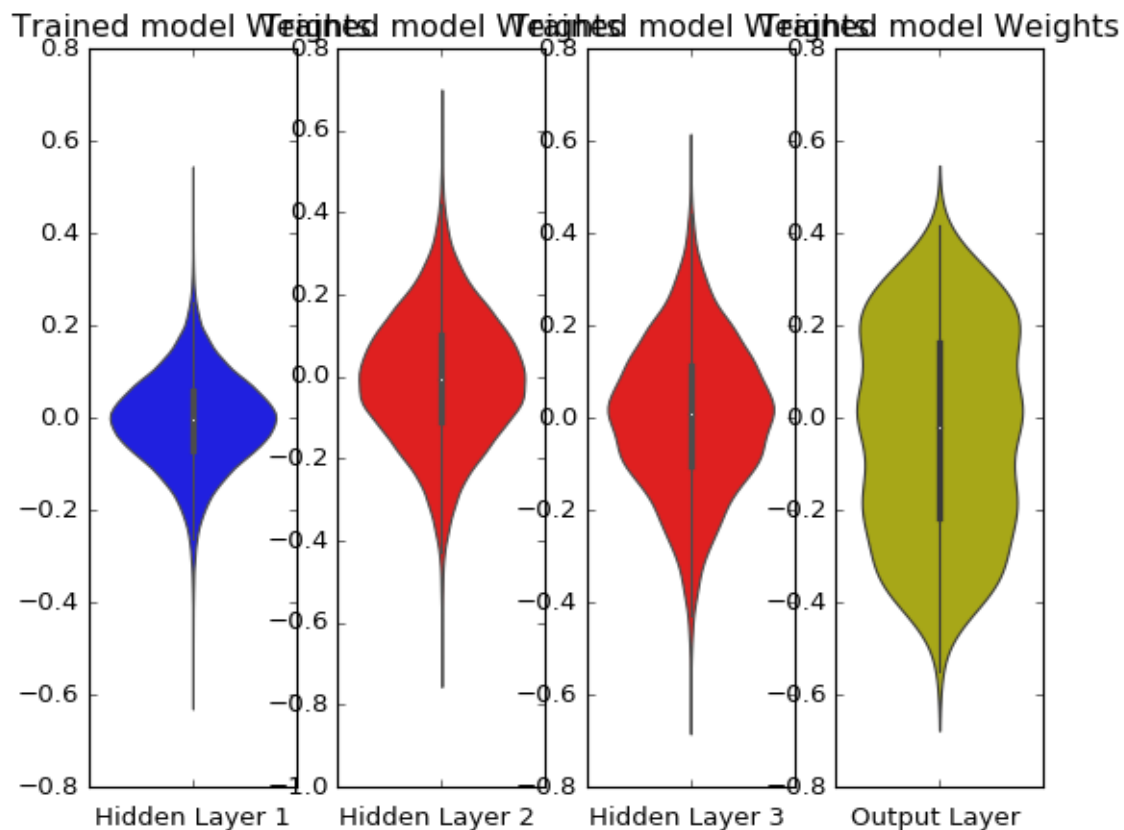
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



C:\Program Files\Anaconda3\lib\site-packages\scipy\stats\stats.py:1626: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

MLP + Batch-Norm on 3 hidden Layers + AdamOptimizer


```

In [61]: # Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition
#  $h1 \Rightarrow \sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(0, \sigma) = N(0, 0.039)$ 
#  $h2 \Rightarrow \sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(0, \sigma) = N(0, 0.055)$ 
#  $h1 \Rightarrow \sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0, std=0.039)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0, std=0.055)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0, std=0.120)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()

```

Layer (type)	Output Shape	Param #
dense_39 (Dense)	(None, 364)	285740
batch_normalization_10 (Batch Normalization)	(None, 364)	1456
dense_40 (Dense)	(None, 128)	46720
batch_normalization_11 (Batch Normalization)	(None, 128)	512
dense_41 (Dense)	(None, 52)	6708
batch_normalization_12 (Batch Normalization)	(None, 52)	208
dense_42 (Dense)	(None, 10)	530
Total params: 341,874		
Trainable params: 340,786		
Non-trainable params: 1,088		

```
In [62]: model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoc

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 14s 227us/step - loss: 0.2393 -
acc: 0.9323 - val_loss: 0.1212 - val_acc: 0.9651
Epoch 2/20
60000/60000 [=====] - 12s 193us/step - loss: 0.0855 -
acc: 0.9751 - val_loss: 0.0913 - val_acc: 0.9732
Epoch 3/20
60000/60000 [=====] - 12s 201us/step - loss: 0.0564 -
acc: 0.9828 - val_loss: 0.0801 - val_acc: 0.9749
Epoch 4/20
60000/60000 [=====] - 12s 201us/step - loss: 0.0421 -
acc: 0.9868 - val_loss: 0.0864 - val_acc: 0.9741
Epoch 5/20
60000/60000 [=====] - 13s 222us/step - loss: 0.0326 -
acc: 0.9897 - val_loss: 0.0771 - val_acc: 0.9774
Epoch 6/20
60000/60000 [=====] - 13s 223us/step - loss: 0.0269 -
acc: 0.9915 - val_loss: 0.0807 - val_acc: 0.9780
Epoch 7/20
60000/60000 [=====] - 14s 233us/step - loss: 0.0220 -
acc: 0.9928 - val_loss: 0.0776 - val_acc: 0.9773
Epoch 8/20
60000/60000 [=====] - 14s 229us/step - loss: 0.0186 -
acc: 0.9938 - val_loss: 0.0755 - val_acc: 0.9806
Epoch 9/20
60000/60000 [=====] - 12s 205us/step - loss: 0.0172 -
acc: 0.9946 - val_loss: 0.0756 - val_acc: 0.9787
Epoch 10/20
60000/60000 [=====] - 13s 211us/step - loss: 0.0158 -
acc: 0.9948 - val_loss: 0.0811 - val_acc: 0.9774
Epoch 11/20
60000/60000 [=====] - 12s 206us/step - loss: 0.0164 -
acc: 0.9942 - val_loss: 0.0984 - val_acc: 0.9764
Epoch 12/20
60000/60000 [=====] - 12s 200us/step - loss: 0.0149 -
acc: 0.9949 - val_loss: 0.0768 - val_acc: 0.9802
Epoch 13/20
60000/60000 [=====] - 13s 214us/step - loss: 0.0122 -
acc: 0.9959 - val_loss: 0.0810 - val_acc: 0.9809
Epoch 14/20
60000/60000 [=====] - 13s 221us/step - loss: 0.0130 -
acc: 0.9958 - val_loss: 0.0767 - val_acc: 0.9787
Epoch 15/20
60000/60000 [=====] - 13s 210us/step - loss: 0.0095 -
acc: 0.9968 - val_loss: 0.0832 - val_acc: 0.9797
Epoch 16/20
60000/60000 [=====] - 12s 201us/step - loss: 0.0077 -
acc: 0.9975 - val_loss: 0.0787 - val_acc: 0.9794
Epoch 17/20
60000/60000 [=====] - 12s 194us/step - loss: 0.0123 -
acc: 0.9958 - val_loss: 0.0974 - val_acc: 0.9765
Epoch 18/20
```

```
60000/60000 [=====] - 12s 194us/step - loss: 0.0102 -  
acc: 0.9966 - val_loss: 0.0778 - val_acc: 0.9806  
Epoch 19/20  
60000/60000 [=====] - 12s 204us/step - loss: 0.0056 -  
acc: 0.9982 - val_loss: 0.0787 - val_acc: 0.9812  
Epoch 20/20  
60000/60000 [=====] - 12s 193us/step - loss: 0.0086 -  
acc: 0.9972 - val_loss: 0.1038 - val_acc: 0.9759
```

```

In [63]: score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch)

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.10380259412173763

Test accuracy: 0.9759



```

In [64]: w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

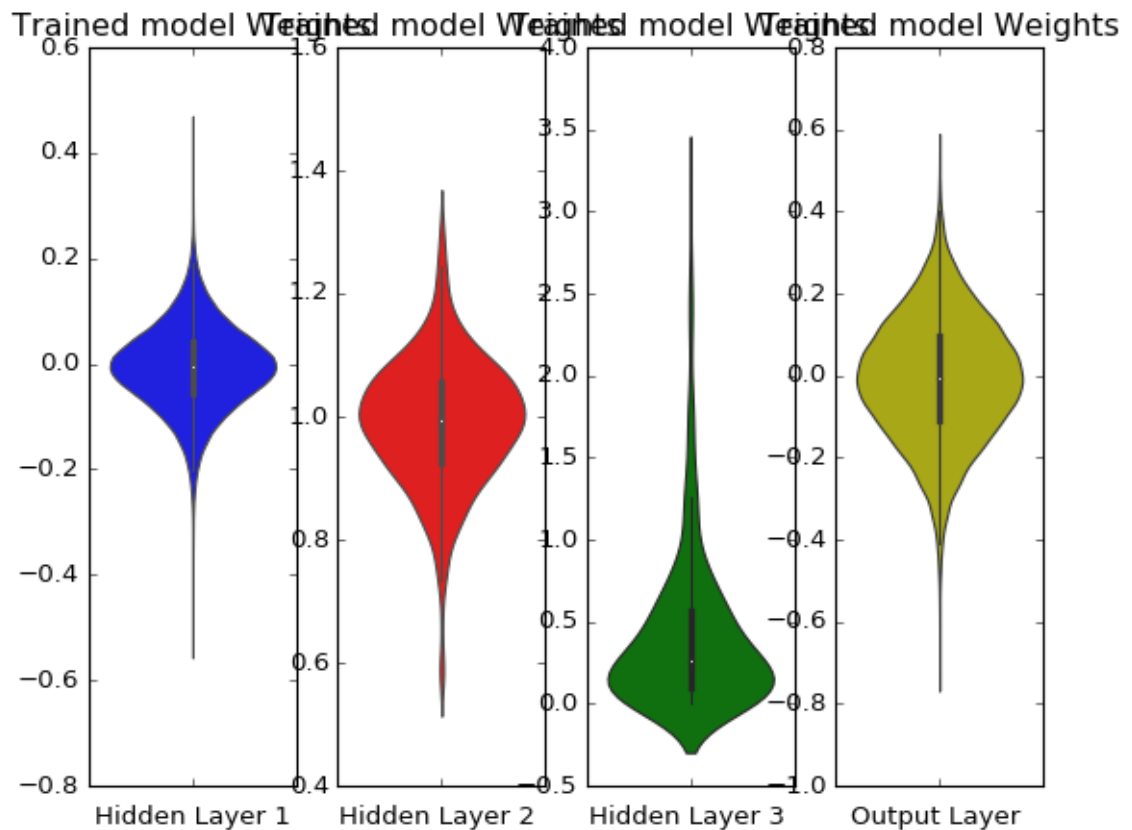
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



C:\Program Files\Anaconda3\lib\site-packages\scipy\stats\stats.py:1626: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

MLP + Dropout + AdamOptimizer with 3 hidden layers

```
In [65]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormaliza

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_43 (Dense)	(None, 364)	285740
dropout_10 (Dropout)	(None, 364)	0
dense_44 (Dense)	(None, 128)	46720
dropout_11 (Dropout)	(None, 128)	0
dense_45 (Dense)	(None, 52)	6708
dropout_12 (Dropout)	(None, 52)	0
dense_46 (Dense)	(None, 10)	530
Total params: 339,698		
Trainable params: 339,698		
Non-trainable params: 0		

```
In [66]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['a
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 12s 205us/step - loss: 1.3565 -
acc: 0.5857 - val_loss: 0.3252 - val_acc: 0.9171
Epoch 2/20
60000/60000 [=====] - 9s 155us/step - loss: 0.5646 - a
cc: 0.8396 - val_loss: 0.2378 - val_acc: 0.9411
Epoch 3/20
60000/60000 [=====] - 9s 152us/step - loss: 0.4256 - a
cc: 0.8865 - val_loss: 0.1949 - val_acc: 0.9506
Epoch 4/20
60000/60000 [=====] - 10s 162us/step - loss: 0.3604 -
acc: 0.9068 - val_loss: 0.1762 - val_acc: 0.9568
Epoch 5/20
60000/60000 [=====] - 10s 167us/step - loss: 0.3266 -
acc: 0.9173 - val_loss: 0.1547 - val_acc: 0.9612
Epoch 6/20
60000/60000 [=====] - 9s 155us/step - loss: 0.2920 - a
cc: 0.9252 - val_loss: 0.1450 - val_acc: 0.9621
Epoch 7/20
60000/60000 [=====] - 9s 152us/step - loss: 0.2694 - a
cc: 0.9332 - val_loss: 0.1352 - val_acc: 0.9662
Epoch 8/20
60000/60000 [=====] - 11s 179us/step - loss: 0.2495 -
acc: 0.9369 - val_loss: 0.1357 - val_acc: 0.9651
Epoch 9/20
60000/60000 [=====] - 12s 206us/step - loss: 0.2358 -
acc: 0.9419 - val_loss: 0.1256 - val_acc: 0.9682
Epoch 10/20
60000/60000 [=====] - 10s 160us/step - loss: 0.2235 -
acc: 0.9438 - val_loss: 0.1291 - val_acc: 0.9685
Epoch 11/20
60000/60000 [=====] - 10s 159us/step - loss: 0.2175 -
acc: 0.9446 - val_loss: 0.1233 - val_acc: 0.9691
Epoch 12/20
60000/60000 [=====] - 9s 153us/step - loss: 0.2040 - a
cc: 0.9483 - val_loss: 0.1175 - val_acc: 0.9708
Epoch 13/20
60000/60000 [=====] - 10s 166us/step - loss: 0.1926 -
acc: 0.9515 - val_loss: 0.1191 - val_acc: 0.9714
Epoch 14/20
60000/60000 [=====] - 10s 170us/step - loss: 0.1826 -
acc: 0.9534 - val_loss: 0.1161 - val_acc: 0.9722
Epoch 15/20
60000/60000 [=====] - 11s 188us/step - loss: 0.1841 -
acc: 0.9549 - val_loss: 0.1190 - val_acc: 0.9714
Epoch 16/20
60000/60000 [=====] - 11s 188us/step - loss: 0.1730 -
acc: 0.9566 - val_loss: 0.1142 - val_acc: 0.9737
Epoch 17/20
60000/60000 [=====] - 11s 190us/step - loss: 0.1637 -
acc: 0.9591 - val_loss: 0.1182 - val_acc: 0.9732
Epoch 18/20
```

```
60000/60000 [=====] - 11s 184us/step - loss: 0.1679 -  
acc: 0.9587 - val_loss: 0.1139 - val_acc: 0.9738  
Epoch 19/20  
60000/60000 [=====] - 12s 192us/step - loss: 0.1586 -  
acc: 0.9605 - val_loss: 0.1078 - val_acc: 0.9762  
Epoch 20/20  
60000/60000 [=====] - 11s 188us/step - loss: 0.1546 -  
acc: 0.9600 - val_loss: 0.1115 - val_acc: 0.9738
```



```

In [67]: score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch)

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.1115148579780769

Test accuracy: 0.9738



```
In [68]: w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)



fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

Figure 1



C:\Program Files\Anaconda3\lib\site-packages\scipy\stats\stats.py:1626: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

MLP + BatchNormalization + Dropout + AdamOptimizer with 3 hidden layers

```
In [69]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormaliza

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(364, activation='relu', input_shape=(input_dim,)), kernel_initializer=RandomNormal(mean=0, stddev=0.01))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_47 (Dense)	(None, 364)	285740
batch_normalization_13 (Batch Normalization)	(None, 364)	1456
dropout_13 (Dropout)	(None, 364)	0
dense_48 (Dense)	(None, 128)	46720
batch_normalization_14 (Batch Normalization)	(None, 128)	512
dropout_14 (Dropout)	(None, 128)	0
dense_49 (Dense)	(None, 52)	6708
batch_normalization_15 (Batch Normalization)	(None, 52)	208
dropout_15 (Dropout)	(None, 52)	0
dense_50 (Dense)	(None, 10)	530
Total params: 341,874		
Trainable params: 340,786		
Non-trainable params: 1,088		

```
In [70]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['a
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 15s 244us/step - loss: 0.9654 -
acc: 0.6985 - val_loss: 0.2492 - val_acc: 0.9254
Epoch 2/20
60000/60000 [=====] - 12s 196us/step - loss: 0.4556 -
acc: 0.8676 - val_loss: 0.1930 - val_acc: 0.9420
Epoch 3/20
60000/60000 [=====] - 12s 199us/step - loss: 0.3565 -
acc: 0.8991 - val_loss: 0.1581 - val_acc: 0.9530
Epoch 4/20
60000/60000 [=====] - 11s 181us/step - loss: 0.2967 -
acc: 0.9176 - val_loss: 0.1319 - val_acc: 0.9612
Epoch 5/20
60000/60000 [=====] - 14s 228us/step - loss: 0.2620 -
acc: 0.9276 - val_loss: 0.1246 - val_acc: 0.9618
Epoch 6/20
60000/60000 [=====] - 13s 220us/step - loss: 0.2385 -
acc: 0.9355 - val_loss: 0.1116 - val_acc: 0.9678
Epoch 7/20
60000/60000 [=====] - 13s 218us/step - loss: 0.2194 -
acc: 0.9394 - val_loss: 0.1086 - val_acc: 0.9696
Epoch 8/20
60000/60000 [=====] - 12s 207us/step - loss: 0.1987 -
acc: 0.9453 - val_loss: 0.0983 - val_acc: 0.9711
Epoch 9/20
60000/60000 [=====] - 11s 185us/step - loss: 0.1889 -
acc: 0.9483 - val_loss: 0.0982 - val_acc: 0.9726
Epoch 10/20
60000/60000 [=====] - 11s 178us/step - loss: 0.1772 -
acc: 0.9508 - val_loss: 0.0983 - val_acc: 0.9715
Epoch 11/20
60000/60000 [=====] - 11s 187us/step - loss: 0.1645 -
acc: 0.9548 - val_loss: 0.0912 - val_acc: 0.9749
Epoch 12/20
60000/60000 [=====] - 13s 217us/step - loss: 0.1559 -
acc: 0.9561 - val_loss: 0.0884 - val_acc: 0.9749
Epoch 13/20
60000/60000 [=====] - 12s 207us/step - loss: 0.1531 -
acc: 0.9581 - val_loss: 0.0874 - val_acc: 0.9769
Epoch 14/20
60000/60000 [=====] - 13s 217us/step - loss: 0.1472 -
acc: 0.9598 - val_loss: 0.0812 - val_acc: 0.9767
Epoch 15/20
60000/60000 [=====] - 13s 216us/step - loss: 0.1385 -
acc: 0.9629 - val_loss: 0.0811 - val_acc: 0.9777
Epoch 16/20
60000/60000 [=====] - 13s 211us/step - loss: 0.1311 -
acc: 0.9638 - val_loss: 0.0839 - val_acc: 0.9772
Epoch 17/20
60000/60000 [=====] - 13s 213us/step - loss: 0.1308 -
acc: 0.9638 - val_loss: 0.0787 - val_acc: 0.9774
Epoch 18/20
```

```
60000/60000 [=====] - 12s 207us/step - loss: 0.1242 -  
acc: 0.9658 - val_loss: 0.0790 - val_acc: 0.9783  
Epoch 19/20  
60000/60000 [=====] - 13s 214us/step - loss: 0.1183 -  
acc: 0.9678 - val_loss: 0.0730 - val_acc: 0.9809  
Epoch 20/20  
60000/60000 [=====] - 12s 202us/step - loss: 0.1146 -  
acc: 0.9688 - val_loss: 0.0778 - val_acc: 0.9784
```

```

In [71]: score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch)

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.07784969780180837

Test accuracy: 0.9784



```

In [72]: w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```




```
C:\Program Files\Anaconda3\lib\site-packages\scipy\stats\stats.py:1626: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.
```

```
    return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

In []:

MLP + ReLU + ADAM with 5 layers without Dropout and Batch Normalisation

```
In [73]: model_relu = Sequential()
model_relu.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_relu.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_relu.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_relu.add(Dense(16, activation='relu', kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epochs)
```

Layer (type)	Output Shape	Param #
dense_51 (Dense)	(None, 256)	200960
dense_52 (Dense)	(None, 128)	32896
dense_53 (Dense)	(None, 64)	8256
dense_54 (Dense)	(None, 32)	2080
dense_55 (Dense)	(None, 16)	528
dense_56 (Dense)	(None, 10)	170
Total params: 244,890		
Trainable params: 244,890		
Non-trainable params: 0		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 10s 161us/step - loss: 0.3539 - acc: 0.8935 - val_loss: 0.1465 - val_acc: 0.9570

Epoch 2/20

60000/60000 [=====] - 7s 119us/step - loss: 0.1118 - acc: 0.9668 - val_loss: 0.1290 - val_acc: 0.9642

Epoch 3/20

60000/60000 [=====] - 7s 118us/step - loss: 0.0786 - acc: 0.9766 - val_loss: 0.0887 - val_acc: 0.9746

Epoch 4/20

60000/60000 [=====] - 7s 108us/step - loss: 0.0568 - acc: 0.9831 - val_loss: 0.0948 - val_acc: 0.9729

Epoch 5/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0438 - acc: 0.9861 - val_loss: 0.0930 - val_acc: 0.9731

Epoch 6/20

60000/60000 [=====] - 6s 107us/step - loss: 0.0366 - acc: 0.9878 - val_loss: 0.0779 - val_acc: 0.9767

Epoch 7/20

60000/60000 [=====] - 7s 119us/step - loss: 0.0285 - acc: 0.9909 - val_loss: 0.0853 - val_acc: 0.9760

```
Epoch 8/20
60000/60000 [=====] - 6s 108us/step - loss: 0.0253 - a
cc: 0.9914 - val_loss: 0.0911 - val_acc: 0.9763
Epoch 9/20
60000/60000 [=====] - 6s 104us/step - loss: 0.0187 - a
cc: 0.9937 - val_loss: 0.1003 - val_acc: 0.9752
Epoch 10/20
60000/60000 [=====] - 6s 99us/step - loss: 0.0200 - ac
c: 0.9936 - val_loss: 0.1284 - val_acc: 0.9674
Epoch 11/20
60000/60000 [=====] - 6s 103us/step - loss: 0.0191 - a
cc: 0.9936 - val_loss: 0.0905 - val_acc: 0.9775
Epoch 12/20
60000/60000 [=====] - 6s 99us/step - loss: 0.0159 - ac
c: 0.9945 - val_loss: 0.0951 - val_acc: 0.9773
Epoch 13/20
60000/60000 [=====] - 6s 97us/step - loss: 0.0119 - ac
c: 0.9962 - val_loss: 0.0942 - val_acc: 0.9801
Epoch 14/20
60000/60000 [=====] - 6s 97us/step - loss: 0.0162 - ac
c: 0.9947 - val_loss: 0.0961 - val_acc: 0.9775
Epoch 15/20
60000/60000 [=====] - 6s 98us/step - loss: 0.0127 - ac
c: 0.9960 - val_loss: 0.1063 - val_acc: 0.9775
Epoch 16/20
60000/60000 [=====] - 6s 103us/step - loss: 0.0139 - a
cc: 0.9956 - val_loss: 0.1199 - val_acc: 0.9762
Epoch 17/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0128 - a
cc: 0.9960 - val_loss: 0.0911 - val_acc: 0.9807
Epoch 18/20
60000/60000 [=====] - 6s 101us/step - loss: 0.0104 - a
cc: 0.9970 - val_loss: 0.1075 - val_acc: 0.9775
Epoch 19/20
60000/60000 [=====] - 6s 100us/step - loss: 0.0100 - a
cc: 0.9970 - val_loss: 0.1037 - val_acc: 0.9791
Epoch 20/20
60000/60000 [=====] - 6s 97us/step - loss: 0.0137 - ac
c: 0.9957 - val_loss: 0.0950 - val_acc: 0.9809
```

```

In [74]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epochs)

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09499137327706776

Test accuracy: 0.9809




```
In [78]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='cyan')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='brown')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



C:\Program Files\Anaconda3\lib\site-packages\scipy\stats\stats.py:1626: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

MLP + Batch-Norm on 5 hidden Layers + AdamOptimizer

```

In [79]: # Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition
#  $h1 \Rightarrow \sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(0, \sigma) = N(0, 0.039)$ 
#  $h2 \Rightarrow \sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(0, \sigma) = N(0, 0.055)$ 
#  $h1 \Rightarrow \sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0, std=0.039)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(132, activation='relu', kernel_initializer=RandomNormal(mean=0, std=0.055)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0, std=0.120)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0, std=0.120)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(16, activation='relu', kernel_initializer=RandomNormal(mean=0, std=0.120)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()

```

Layer (type)	Output Shape	Param #
=====		
dense_57 (Dense)	(None, 256)	200960
batch_normalization_16 (Batch Normalization)	(None, 256)	1024
dense_58 (Dense)	(None, 132)	33924
batch_normalization_17 (Batch Normalization)	(None, 132)	528
dense_59 (Dense)	(None, 64)	8512
batch_normalization_18 (Batch Normalization)	(None, 64)	256
dense_60 (Dense)	(None, 32)	2080
batch_normalization_19 (Batch Normalization)	(None, 32)	128
dense_61 (Dense)	(None, 16)	528
batch_normalization_20 (Batch Normalization)	(None, 16)	64
dense_62 (Dense)	(None, 10)	170


```
=====
Total params: 248,174
Trainable params: 247,174
Non-trainable params: 1,000

```

```
In [80]: model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoc

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 13s 215us/step - loss: 0.3855 -
acc: 0.9065 - val_loss: 0.1549 - val_acc: 0.9567
Epoch 2/20
60000/60000 [=====] - 8s 138us/step - loss: 0.1210 - a
cc: 0.9662 - val_loss: 0.1150 - val_acc: 0.9672
Epoch 3/20
60000/60000 [=====] - 9s 153us/step - loss: 0.0813 - a
cc: 0.9761 - val_loss: 0.1012 - val_acc: 0.9709
Epoch 4/20
60000/60000 [=====] - 9s 145us/step - loss: 0.0631 - a
cc: 0.9806 - val_loss: 0.0808 - val_acc: 0.9769
Epoch 5/20
60000/60000 [=====] - 9s 144us/step - loss: 0.0512 - a
cc: 0.9840 - val_loss: 0.1057 - val_acc: 0.9674
Epoch 6/20
60000/60000 [=====] - 8s 138us/step - loss: 0.0426 - a
cc: 0.9870 - val_loss: 0.0886 - val_acc: 0.9744
Epoch 7/20
60000/60000 [=====] - 8s 139us/step - loss: 0.0331 - a
cc: 0.9896 - val_loss: 0.0946 - val_acc: 0.9729
Epoch 8/20
60000/60000 [=====] - 8s 141us/step - loss: 0.0330 - a
cc: 0.9895 - val_loss: 0.0868 - val_acc: 0.9762
Epoch 9/20
60000/60000 [=====] - 8s 139us/step - loss: 0.0313 - a
cc: 0.9899 - val_loss: 0.0887 - val_acc: 0.9759
Epoch 10/20
60000/60000 [=====] - 10s 160us/step - loss: 0.0246 -
acc: 0.9916 - val_loss: 0.0869 - val_acc: 0.9760
Epoch 11/20
60000/60000 [=====] - 9s 145us/step - loss: 0.0260 - a
cc: 0.9913 - val_loss: 0.0831 - val_acc: 0.9779
Epoch 12/20
60000/60000 [=====] - 8s 135us/step - loss: 0.0216 - a
cc: 0.9928 - val_loss: 0.0854 - val_acc: 0.9783
Epoch 13/20
60000/60000 [=====] - 8s 132us/step - loss: 0.0201 - a
cc: 0.9937 - val_loss: 0.0865 - val_acc: 0.9784
Epoch 14/20
60000/60000 [=====] - 8s 133us/step - loss: 0.0176 - a
cc: 0.9945 - val_loss: 0.0784 - val_acc: 0.9794
Epoch 15/20
60000/60000 [=====] - 9s 153us/step - loss: 0.0165 - a
cc: 0.9946 - val_loss: 0.0824 - val_acc: 0.9781
Epoch 16/20
60000/60000 [=====] - 8s 141us/step - loss: 0.0150 - a
cc: 0.9950 - val_loss: 0.0869 - val_acc: 0.9804
Epoch 17/20
60000/60000 [=====] - 10s 166us/step - loss: 0.0164 -
acc: 0.9944 - val_loss: 0.0667 - val_acc: 0.9820
Epoch 18/20
```

```
60000/60000 [=====] - 9s 158us/step - loss: 0.0136 - a
cc: 0.9954 - val_loss: 0.0785 - val_acc: 0.9811
Epoch 19/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0158 - a
cc: 0.9947 - val_loss: 0.0868 - val_acc: 0.9791
Epoch 20/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0139 - a
cc: 0.9956 - val_loss: 0.0819 - val_acc: 0.9814
```

```
In [84]: score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

```
Test score: 0.08187933229937626
Test accuracy: 0.9814
```

In [85]:

```

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch)

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```



```
In [86]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='cyan')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='brown')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



C:\Program Files\Anaconda3\lib\site-packages\scipy\stats\stats.py:1626: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

MLP + Dropout + AdamOptimizer with 5 hidden layers

In [87]: `# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormaliza`

```
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(256, activation='relu', input_shape=(input_dim,)), kernel_initializer=RandomNormal(mean=0, stddev=0.01))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(16, activation='relu', kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_63 (Dense)	(None, 256)	200960
dropout_16 (Dropout)	(None, 256)	0
dense_64 (Dense)	(None, 128)	32896
dropout_17 (Dropout)	(None, 128)	0
dense_65 (Dense)	(None, 64)	8256
dropout_18 (Dropout)	(None, 64)	0
dense_66 (Dense)	(None, 32)	2080
dropout_19 (Dropout)	(None, 32)	0
dense_67 (Dense)	(None, 16)	528
dropout_20 (Dropout)	(None, 16)	0
dense_68 (Dense)	(None, 10)	170
Total params: 244,890		
Trainable params: 244,890		
Non-trainable params: 0		

```
In [88]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['a
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 12s 200us/step - loss: 2.1425 -
acc: 0.2220 - val_loss: 1.5689 - val_acc: 0.5284
Epoch 2/20
60000/60000 [=====] - 10s 160us/step - loss: 1.6576 -
acc: 0.4145 - val_loss: 1.3220 - val_acc: 0.5546
Epoch 3/20
60000/60000 [=====] - 8s 138us/step - loss: 1.4285 - a
cc: 0.5034 - val_loss: 1.0335 - val_acc: 0.6939
Epoch 4/20
60000/60000 [=====] - 8s 136us/step - loss: 1.2482 - a
cc: 0.5584 - val_loss: 0.8864 - val_acc: 0.7516
Epoch 5/20
60000/60000 [=====] - 8s 137us/step - loss: 1.1219 - a
cc: 0.6141 - val_loss: 0.7616 - val_acc: 0.8076
Epoch 6/20
60000/60000 [=====] - 8s 139us/step - loss: 1.0295 - a
cc: 0.6447 - val_loss: 0.6917 - val_acc: 0.8053
Epoch 7/20
60000/60000 [=====] - 8s 135us/step - loss: 0.9681 - a
cc: 0.6627 - val_loss: 0.6602 - val_acc: 0.8036
Epoch 8/20
60000/60000 [=====] - 8s 137us/step - loss: 0.9128 - a
cc: 0.6807 - val_loss: 0.6569 - val_acc: 0.7593
Epoch 9/20
60000/60000 [=====] - 8s 131us/step - loss: 0.8891 - a
cc: 0.6880 - val_loss: 0.6218 - val_acc: 0.7908
Epoch 10/20
60000/60000 [=====] - 8s 133us/step - loss: 0.8644 - a
cc: 0.6985 - val_loss: 0.5864 - val_acc: 0.8127
Epoch 11/20
60000/60000 [=====] - 8s 134us/step - loss: 0.8475 - a
cc: 0.7026 - val_loss: 0.5951 - val_acc: 0.7848
Epoch 12/20
60000/60000 [=====] - 8s 138us/step - loss: 0.8229 - a
cc: 0.7129 - val_loss: 0.5810 - val_acc: 0.8239
Epoch 13/20
60000/60000 [=====] - 8s 139us/step - loss: 0.7966 - a
cc: 0.7257 - val_loss: 0.5449 - val_acc: 0.8237
Epoch 14/20
60000/60000 [=====] - 8s 139us/step - loss: 0.7932 - a
cc: 0.7276 - val_loss: 0.5404 - val_acc: 0.8323
Epoch 15/20
60000/60000 [=====] - 9s 146us/step - loss: 0.7958 - a
cc: 0.7321 - val_loss: 0.5297 - val_acc: 0.8271
Epoch 16/20
60000/60000 [=====] - 9s 150us/step - loss: 0.7734 - a
cc: 0.7401 - val_loss: 0.5279 - val_acc: 0.8400
Epoch 17/20
60000/60000 [=====] - 9s 152us/step - loss: 0.7516 - a
cc: 0.7449 - val_loss: 0.5365 - val_acc: 0.8304
Epoch 18/20
```



```
60000/60000 [=====] - 9s 144us/step - loss: 0.7421 - a
cc: 0.7506 - val_loss: 0.5041 - val_acc: 0.8411
Epoch 19/20
60000/60000 [=====] - 8s 141us/step - loss: 0.7363 - a
cc: 0.7537 - val_loss: 0.4806 - val_acc: 0.8496
Epoch 20/20
60000/60000 [=====] - 9s 150us/step - loss: 0.7230 - a
cc: 0.7584 - val_loss: 0.4595 - val_acc: 0.8580
```

```
In [89]: score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

```
Test score: 0.4594558026790619
Test accuracy: 0.858
```

```

In [90]: fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epochs)

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```



```
In [91]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='cyan')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='brown')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



C:\Program Files\Anaconda3\lib\site-packages\scipy\stats\stats.py:1626: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

MLP + BatchNormalization + Dropout + AdamOptimizer with 3 hidden layers

In [92]: `# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormaliza`

```
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(16, activation='relu', kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_69 (Dense)	(None, 256)	200960
batch_normalization_21 (Batch Normalization)	(None, 256)	1024
dropout_21 (Dropout)	(None, 256)	0
dense_70 (Dense)	(None, 128)	32896
batch_normalization_22 (Batch Normalization)	(None, 128)	512
dropout_22 (Dropout)	(None, 128)	0
dense_71 (Dense)	(None, 64)	8256
batch_normalization_23 (Batch Normalization)	(None, 64)	256
dropout_23 (Dropout)	(None, 64)	0
dense_72 (Dense)	(None, 32)	2080
batch_normalization_24 (Batch Normalization)	(None, 32)	128

dropout_24 (Dropout)	(None, 32)	0
dense_73 (Dense)	(None, 16)	528
batch_normalization_25 (Batch Normalization)	(None, 16)	64
dropout_25 (Dropout)	(None, 16)	0
dense_74 (Dense)	(None, 10)	170
=====		
Total params: 246,874		
Trainable params: 245,882		
Non-trainable params: 992		

```
In [93]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['a
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 16s 271us/step - loss: 2.1278 -
acc: 0.2514 - val_loss: 1.1959 - val_acc: 0.7221
Epoch 2/20
60000/60000 [=====] - 11s 176us/step - loss: 1.4928 -
acc: 0.4529 - val_loss: 0.7785 - val_acc: 0.7865
Epoch 3/20
60000/60000 [=====] - 11s 184us/step - loss: 1.1906 -
acc: 0.5626 - val_loss: 0.5776 - val_acc: 0.8514
Epoch 4/20
60000/60000 [=====] - 11s 185us/step - loss: 1.0123 -
acc: 0.6299 - val_loss: 0.4885 - val_acc: 0.8413
Epoch 5/20
60000/60000 [=====] - 11s 185us/step - loss: 0.9047 -
acc: 0.6725 - val_loss: 0.4304 - val_acc: 0.8450
Epoch 6/20
60000/60000 [=====] - 11s 181us/step - loss: 0.8163 -
acc: 0.7062 - val_loss: 0.3703 - val_acc: 0.8852
Epoch 7/20
60000/60000 [=====] - 9s 154us/step - loss: 0.7419 - a
cc: 0.7370 - val_loss: 0.3397 - val_acc: 0.8845
Epoch 8/20
60000/60000 [=====] - 10s 171us/step - loss: 0.6941 -
acc: 0.7580 - val_loss: 0.3168 - val_acc: 0.9081
Epoch 9/20
60000/60000 [=====] - 11s 191us/step - loss: 0.6566 -
acc: 0.7774 - val_loss: 0.2958 - val_acc: 0.9234
Epoch 10/20
60000/60000 [=====] - 10s 168us/step - loss: 0.6195 -
acc: 0.7901 - val_loss: 0.2795 - val_acc: 0.9047
Epoch 11/20
60000/60000 [=====] - 10s 174us/step - loss: 0.5903 -
acc: 0.8044 - val_loss: 0.2705 - val_acc: 0.9022
Epoch 12/20
60000/60000 [=====] - 10s 167us/step - loss: 0.5624 -
acc: 0.8158 - val_loss: 0.2521 - val_acc: 0.9353
Epoch 13/20
60000/60000 [=====] - 12s 192us/step - loss: 0.5452 -
acc: 0.8250 - val_loss: 0.2466 - val_acc: 0.9243
Epoch 14/20
60000/60000 [=====] - 12s 200us/step - loss: 0.5280 -
acc: 0.8308 - val_loss: 0.2431 - val_acc: 0.9507
Epoch 15/20
60000/60000 [=====] - 12s 200us/step - loss: 0.5180 -
acc: 0.8363 - val_loss: 0.2258 - val_acc: 0.9539
Epoch 16/20
60000/60000 [=====] - 12s 198us/step - loss: 0.5051 -
acc: 0.8412 - val_loss: 0.2227 - val_acc: 0.9540
Epoch 17/20
60000/60000 [=====] - 12s 197us/step - loss: 0.4821 -
acc: 0.8486 - val_loss: 0.2090 - val_acc: 0.9575
Epoch 18/20
```

```
60000/60000 [=====] - 12s 202us/step - loss: 0.4715 -  
acc: 0.8532 - val_loss: 0.2110 - val_acc: 0.9592  
Epoch 19/20  
60000/60000 [=====] - 12s 197us/step - loss: 0.4581 -  
acc: 0.8579 - val_loss: 0.2048 - val_acc: 0.9599  
Epoch 20/20  
60000/60000 [=====] - 11s 176us/step - loss: 0.4610 -  
acc: 0.8600 - val_loss: 0.2019 - val_acc: 0.9590
```

```
In [95]: score = model_drop.evaluate(X_test, Y_test, verbose=0)  
print('Test score:', score[0])  
print('Test accuracy:', score[1])
```

```
Test score: 0.2019003223091364  
Test accuracy: 0.959
```



```

In [96]: fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch)

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```



```
In [97]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='cyan')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='brown')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



C:\Program Files\Anaconda3\lib\site-packages\scipy\stats\stats.py:1626: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

In [98]: *#Comparisons*

In [99]: **from** prettytable **import** PrettyTable

```
x = PrettyTable()
x.field_names = ["Architecture", "parameters", "Accuracy"]

x.add_row(["2 layer", "without Dropout and Batch Normalization", 97.9])
x.add_row(["2 layer", "with Dropuot", 97.44])
x.add_row(["2 layer", "with Batch Normalization", 97.98])
x.add_row(["2 layer", "with Dropuot and Batch Normalization ", 98.08])

x.add_row(["3 layer", "without Dropout and Batch Normalization", 97.87])
x.add_row(["3 layer", "with Dropuot", 97.38])
x.add_row(["3 layer", "with Batch Normalization", 97.59])
x.add_row(["3 layer", "with Dropuot and Batch Normalization ", 97.84])

x.add_row(["5 layer", "without Dropout and Batch Normalization", 98.09])
x.add_row(["5 layer", "with Dropuot", 85.8])
x.add_row(["5 layer", "with Batch Normalization", 98.14])
x.add_row(["5 layer", "with Dropuot and Batch Normalization ", 95.9])

print(x)
```

Architecture	parameters	Accuracy
2 layer	without Dropout and Batch Normalization	97.9
2 layer	with Dropuot	97.44
2 layer	with Batch Normalization	97.98
2 layer	with Dropuot and Batch Normalization	98.08
3 layer	without Dropout and Batch Normalization	97.87
3 layer	with Dropuot	97.38
3 layer	with Batch Normalization	97.59
3 layer	with Dropuot and Batch Normalization	97.84
5 layer	without Dropout and Batch Normalization	98.09
5 layer	with Dropuot	85.8
5 layer	with Batch Normalization	98.14
5 layer	with Dropuot and Batch Normalization	95.9

In []: