

Keras -- MLPs on MNIST

In [1]:

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this c
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

Using TensorFlow backend.

In [2]:

```
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [3]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz> (<http://s3.amazonaws.com/img-datasets/mnist.npz>)
11493376/11490434 [=====] - 10s 1us/step

In [4]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)" % (X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)" % (X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)

In [5]:

```
# if you observe the input shape its 3 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [6]:

```
# after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)%"
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)%(

Number of training examples : 60000 and each image is of shape (784)
Number of training examples : 10000 and each image is of shape (784)
```

In [7]:

```
# An example data point
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  249 253 249  64  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  23  66 213 253 253 253 253 198  81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  18 171 219 253 253 253 195
 80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 55 172 226 253 253 253 253 244 133  11  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 136 253 253 253 212 135 132  16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0]
```

In [8]:

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 
```

```
X_train = X_train/255
X_test = X_test/255
```

In [9]:

```
# example data point after normlizing
print(X_train[0])
```

```
[0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.]
```

In [10]:

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# Lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs
```

```
Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

Class label of first image : 5

After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

Softmax classifier

In [11]:

```
# https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT.X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument of a layer.

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation
```

In [12]:

```
# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

In [13]:

```
print(X_train.shape[1])
```

784

MLP + ReLU + ADAM with 2 layers without Dropout and Batch Normalisation

In [14]:

```

model_relu = Sequential()
model_relu.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_relu.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)

```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 364)	285740
dense_2 (Dense)	(None, 52)	18980
dense_3 (Dense)	(None, 10)	530

=====
 Total params: 305,250
 Trainable params: 305,250
 Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 95us/step - loss: 0.2546 - acc: 0.9248 - val_loss: 0.1287 - val_acc: 0.9603

Epoch 2/20

60000/60000 [=====] - 6s 96us/step - loss: 0.0997 - acc: 0.9703 - val_loss: 0.0878 - val_acc: 0.9728

Epoch 3/20

60000/60000 [=====] - 5s 85us/step - loss: 0.0641 - acc: 0.9806 - val_loss: 0.0810 - val_acc: 0.9757

Epoch 4/20

60000/60000 [=====] - 5s 89us/step - loss: 0.0466 - acc: 0.9857 - val_loss: 0.0864 - val_acc: 0.9748

Epoch 5/20

60000/60000 [=====] - 5s 89us/step - loss: 0.0356 - acc: 0.9884 - val_loss: 0.0704 - val_acc: 0.9781

Epoch 6/20

60000/60000 [=====] - 6s 93us/step - loss: 0.0241 - acc: 0.9930 - val_loss: 0.0732 - val_acc: 0.9783

Epoch 7/20

60000/60000 [=====] - 6s 95us/step - loss: 0.0189 - acc: 0.9943 - val_loss: 0.0737 - val_acc: 0.9782

Epoch 8/20

60000/60000 [=====] - 8s 125us/step - loss: 0.0150 - acc: 0.9952 - val_loss: 0.0699 - val_acc: 0.9809

Epoch 9/20

60000/60000 [=====] - 7s 115us/step - loss: 0.0121 - acc: 0.9963 - val_loss: 0.0692 - val_acc: 0.9797

Epoch 10/20

60000/60000 [=====] - 7s 113us/step - loss: 0.0101 - acc: 0.9969 - val_loss: 0.0756 - val_acc: 0.9805

Epoch 11/20

60000/60000 [=====] - 6s 108us/step - loss: 0.0115

```
- acc: 0.9963 - val_loss: 0.1052 - val_acc: 0.9741
Epoch 12/20
60000/60000 [=====] - 7s 117us/step - loss: 0.0122
- acc: 0.9961 - val_loss: 0.0850 - val_acc: 0.9799
Epoch 13/20
60000/60000 [=====] - 8s 127us/step - loss: 0.0122
- acc: 0.9958 - val_loss: 0.0838 - val_acc: 0.9796
Epoch 14/20
60000/60000 [=====] - 7s 115us/step - loss: 0.0072
- acc: 0.9978 - val_loss: 0.0788 - val_acc: 0.9816
Epoch 15/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0057 -
acc: 0.9983 - val_loss: 0.0999 - val_acc: 0.9783
Epoch 16/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0091 -
acc: 0.9970 - val_loss: 0.0921 - val_acc: 0.9797
Epoch 17/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0055 -
acc: 0.9982 - val_loss: 0.1007 - val_acc: 0.9792
Epoch 18/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0102 -
acc: 0.9968 - val_loss: 0.1081 - val_acc: 0.9788
Epoch 19/20
60000/60000 [=====] - 5s 80us/step - loss: 0.0067 -
acc: 0.9978 - val_loss: 0.0874 - val_acc: 0.9807
Epoch 20/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0039 -
acc: 0.9989 - val_loss: 0.0861 - val_acc: 0.9821
```

In [15]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

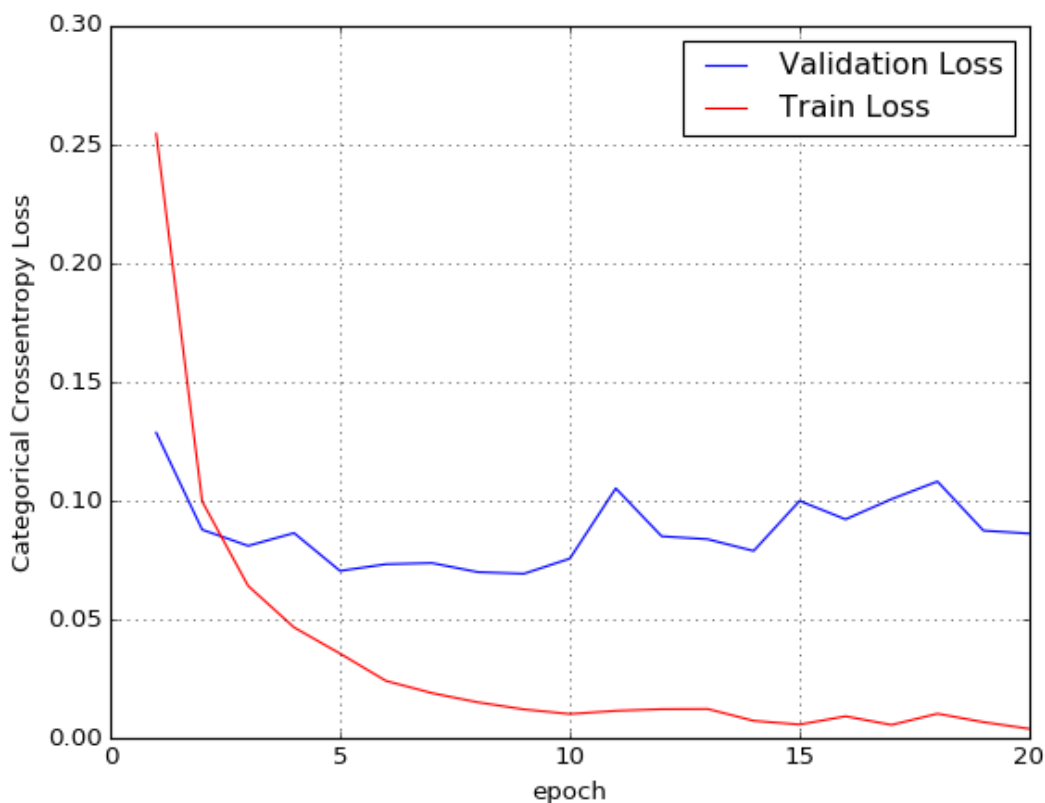
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# Loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08614989673080777

Test accuracy: 0.9821



In [16]:

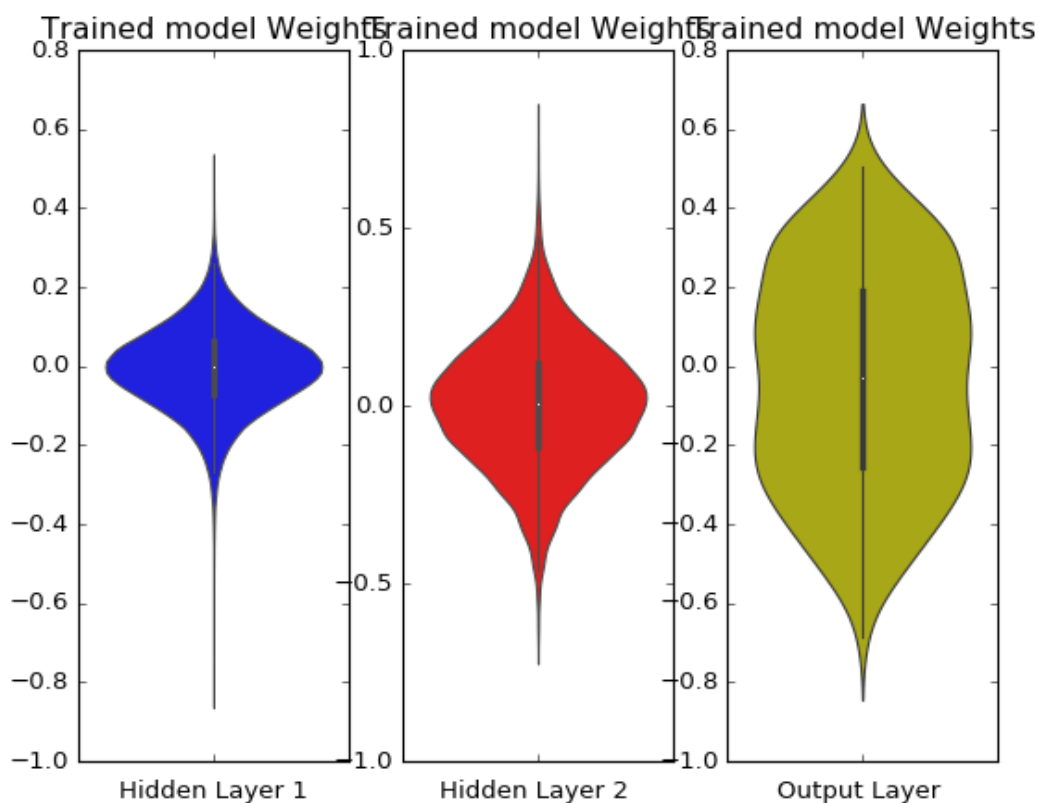
```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Batch-Norm on 2 hidden Layers + AdamOptimizer

In [17]:

```

# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i + n_{i+1})}$ 
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(\theta, \sigma) = N(\theta, 0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(\theta, \sigma) = N(\theta, 0.055)$ 
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(\theta, \sigma) = N(\theta, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, std=0.039)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std=0.055)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()

```

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 364)	285740
batch_normalization_1 (Batch Normalization)	(None, 364)	1456
dense_5 (Dense)	(None, 52)	18980
batch_normalization_2 (Batch Normalization)	(None, 52)	208
dense_6 (Dense)	(None, 10)	530
Total params: 306,914		
Trainable params: 306,082		
Non-trainable params: 832		

In [18]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 101us/step - loss: 0.2357
- acc: 0.9333 - val_loss: 0.1332 - val_acc: 0.9596

Epoch 2/20

60000/60000 [=====] - 5s 85us/step - loss: 0.0887 -
acc: 0.9737 - val_loss: 0.1011 - val_acc: 0.9693

Epoch 3/20

60000/60000 [=====] - 5s 83us/step - loss: 0.0585 -
acc: 0.9826 - val_loss: 0.0838 - val_acc: 0.9760

Epoch 4/20

60000/60000 [=====] - 5s 83us/step - loss: 0.0427 -
acc: 0.9871 - val_loss: 0.0793 - val_acc: 0.9749

Epoch 5/20

60000/60000 [=====] - 5s 82us/step - loss: 0.0306 -
acc: 0.9904 - val_loss: 0.0888 - val_acc: 0.9748

Epoch 6/20

60000/60000 [=====] - 5s 84us/step - loss: 0.0266 -
acc: 0.9915 - val_loss: 0.0750 - val_acc: 0.9794

Epoch 7/20

60000/60000 [=====] - 5s 83us/step - loss: 0.0224 -
acc: 0.9931 - val_loss: 0.0798 - val_acc: 0.9771

Epoch 8/20

60000/60000 [=====] - 5s 84us/step - loss: 0.0169 -
acc: 0.9952 - val_loss: 0.0846 - val_acc: 0.9746

Epoch 9/20

60000/60000 [=====] - 5s 84us/step - loss: 0.0143 -
acc: 0.9957 - val_loss: 0.0946 - val_acc: 0.9752

Epoch 10/20

60000/60000 [=====] - 5s 87us/step - loss: 0.0162 -
acc: 0.9947 - val_loss: 0.0827 - val_acc: 0.9780

Epoch 11/20

60000/60000 [=====] - 9s 149us/step - loss: 0.0115
- acc: 0.9962 - val_loss: 0.0796 - val_acc: 0.9778

Epoch 12/20

60000/60000 [=====] - 6s 107us/step - loss: 0.0108
- acc: 0.9969 - val_loss: 0.0961 - val_acc: 0.9764

Epoch 13/20

60000/60000 [=====] - 6s 95us/step - loss: 0.0103 -
acc: 0.9967 - val_loss: 0.0934 - val_acc: 0.9760

Epoch 14/20

60000/60000 [=====] - 6s 102us/step - loss: 0.0094
- acc: 0.9972 - val_loss: 0.0873 - val_acc: 0.9788

Epoch 15/20

60000/60000 [=====] - 8s 126us/step - loss: 0.0103
- acc: 0.9966 - val_loss: 0.0959 - val_acc: 0.9762

Epoch 16/20

60000/60000 [=====] - 6s 103us/step - loss: 0.0094
- acc: 0.9970 - val_loss: 0.0901 - val_acc: 0.9789

Epoch 17/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0074
- acc: 0.9977 - val_loss: 0.0829 - val_acc: 0.9795

Epoch 18/20

60000/60000 [=====] - 5s 88us/step - loss: 0.0088 -
acc: 0.9971 - val_loss: 0.0912 - val_acc: 0.9773

Epoch 19/20

60000/60000 [=====] - 5s 91us/step - loss: 0.0088 -

acc: 0.9970 - val_loss: 0.0904 - val_acc: 0.9776

Epoch 20/20

60000/60000 [=====] - 5s 87us/step - loss: 0.0058 -

acc: 0.9981 - val_loss: 0.0895 - val_acc: 0.9792

In [20]:

```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

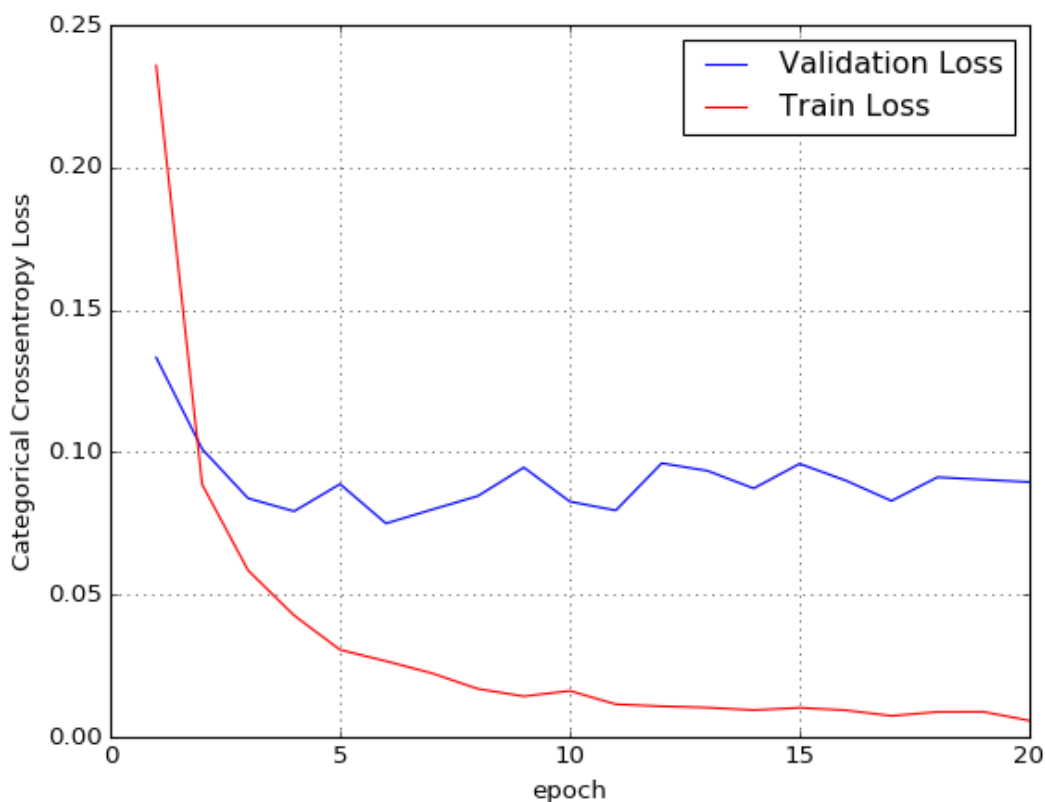
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# Loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08952085421274897

Test accuracy: 0.9792



In [21]:

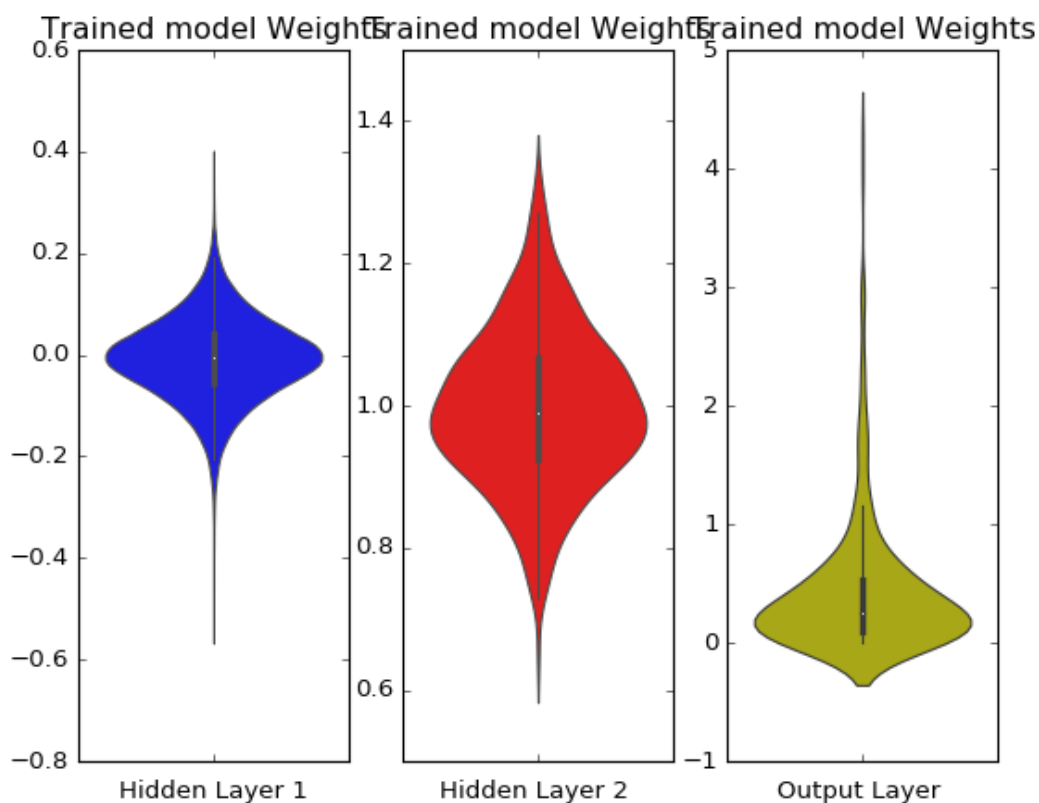
```
w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Dropout (rate = 0.5) + AdamOptimizer with 2 hidden layers

In [22]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-funct

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
=====		
dense_7 (Dense)	(None, 364)	285740
dropout_1 (Dropout)	(None, 364)	0
dense_8 (Dense)	(None, 52)	18980
dropout_2 (Dropout)	(None, 52)	0
dense_9 (Dense)	(None, 10)	530
=====		
Total params: 305,250		
Trainable params: 305,250		
Non-trainable params: 0		

In [23]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 97us/step - loss: 0.9070 - acc: 0.7097 - val_loss: 0.2468 - val_acc: 0.9309

Epoch 2/20

60000/60000 [=====] - 5s 82us/step - loss: 0.4398 - acc: 0.8680 - val_loss: 0.1933 - val_acc: 0.9446

Epoch 3/20

60000/60000 [=====] - 5s 83us/step - loss: 0.3459 - acc: 0.8991 - val_loss: 0.1679 - val_acc: 0.9530

Epoch 4/20

60000/60000 [=====] - 5s 83us/step - loss: 0.3020 - acc: 0.9151 - val_loss: 0.1473 - val_acc: 0.9577

Epoch 5/20

60000/60000 [=====] - 5s 82us/step - loss: 0.2697 - acc: 0.9248 - val_loss: 0.1346 - val_acc: 0.9635

Epoch 6/20

60000/60000 [=====] - 5s 83us/step - loss: 0.2425 - acc: 0.9330 - val_loss: 0.1258 - val_acc: 0.9637

Epoch 7/20

60000/60000 [=====] - 5s 81us/step - loss: 0.2166 - acc: 0.9393 - val_loss: 0.1239 - val_acc: 0.9664

Epoch 8/20

60000/60000 [=====] - 5s 84us/step - loss: 0.2087 - acc: 0.9416 - val_loss: 0.1113 - val_acc: 0.9706

Epoch 9/20

60000/60000 [=====] - 5s 81us/step - loss: 0.1947 - acc: 0.9464 - val_loss: 0.1106 - val_acc: 0.9709

Epoch 10/20

60000/60000 [=====] - 5s 80us/step - loss: 0.1828 - acc: 0.9490 - val_loss: 0.1110 - val_acc: 0.9702

Epoch 11/20

60000/60000 [=====] - 5s 79us/step - loss: 0.1713 - acc: 0.9524 - val_loss: 0.1102 - val_acc: 0.9718

Epoch 12/20

60000/60000 [=====] - 6s 104us/step - loss: 0.1651 - acc: 0.9533 - val_loss: 0.1041 - val_acc: 0.9725

Epoch 13/20

60000/60000 [=====] - 5s 82us/step - loss: 0.1594 - acc: 0.9558 - val_loss: 0.1029 - val_acc: 0.9735

Epoch 14/20

60000/60000 [=====] - 6s 92us/step - loss: 0.1524 - acc: 0.9574 - val_loss: 0.1051 - val_acc: 0.9737

Epoch 15/20

60000/60000 [=====] - 6s 100us/step - loss: 0.1457 - acc: 0.9592 - val_loss: 0.1000 - val_acc: 0.9746

Epoch 16/20

60000/60000 [=====] - 6s 100us/step - loss: 0.1384 - acc: 0.9602 - val_loss: 0.1041 - val_acc: 0.9763

Epoch 17/20

60000/60000 [=====] - 5s 90us/step - loss: 0.1372 - acc: 0.9613 - val_loss: 0.0975 - val_acc: 0.9759

Epoch 18/20

60000/60000 [=====] - 5s 82us/step - loss: 0.1306 - acc: 0.9638 - val_loss: 0.0993 - val_acc: 0.9753

Epoch 19/20

60000/60000 [=====] - 5s 87us/step - loss: 0.1241 -

acc: 0.9656 - val_loss: 0.0965 - val_acc: 0.9763

Epoch 20/20

60000/60000 [=====] - 6s 106us/step - loss: 0.1210

- acc: 0.9660 - val_loss: 0.0963 - val_acc: 0.9756

In [24]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

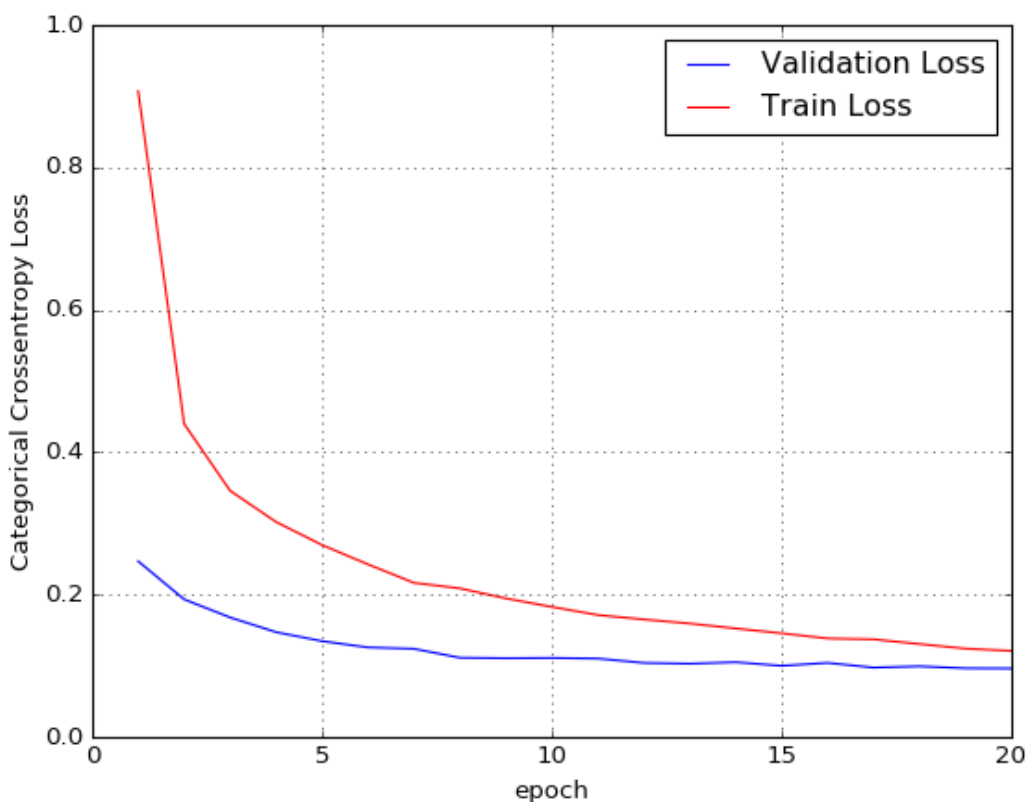
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# Loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09625575973058877

Test accuracy: 0.9756



In [25]:

```

w_after = model_drop.get_weights()

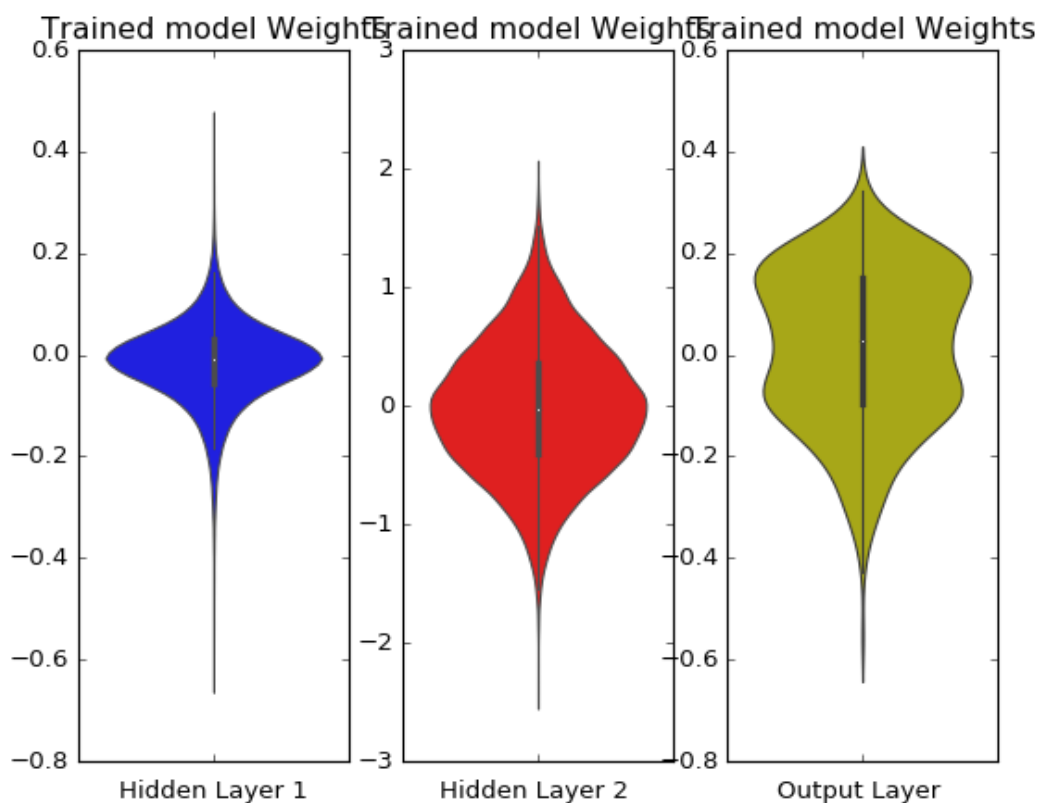
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



In []:

MLP + Dropout (rate = 0.3) + AdamOptimizer with 2 hidden layers

In [26]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-funct
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.3))

model_drop.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.3))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_10 (Dense)	(None, 364)	285740
dropout_3 (Dropout)	(None, 364)	0
dense_11 (Dense)	(None, 52)	18980
dropout_4 (Dropout)	(None, 52)	0
dense_12 (Dense)	(None, 10)	530
Total params: 305,250		
Trainable params: 305,250		
Non-trainable params: 0		

In [27]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 99us/step - loss: 0.5129 - acc: 0.8458 - val_loss: 0.1645 - val_acc: 0.9502

Epoch 2/20

60000/60000 [=====] - 5s 79us/step - loss: 0.2464 - acc: 0.9262 - val_loss: 0.1239 - val_acc: 0.9634

Epoch 3/20

60000/60000 [=====] - 5s 77us/step - loss: 0.1909 - acc: 0.9432 - val_loss: 0.1123 - val_acc: 0.9677

Epoch 4/20

60000/60000 [=====] - 5s 76us/step - loss: 0.1577 - acc: 0.9533 - val_loss: 0.0969 - val_acc: 0.9705

Epoch 5/20

60000/60000 [=====] - 5s 84us/step - loss: 0.1387 - acc: 0.9594 - val_loss: 0.0890 - val_acc: 0.9736

Epoch 6/20

60000/60000 [=====] - 5s 84us/step - loss: 0.1246 - acc: 0.9627 - val_loss: 0.0855 - val_acc: 0.9765

Epoch 7/20

60000/60000 [=====] - 5s 86us/step - loss: 0.1100 - acc: 0.9670 - val_loss: 0.0869 - val_acc: 0.9757

Epoch 8/20

60000/60000 [=====] - 5s 83us/step - loss: 0.1032 - acc: 0.9685 - val_loss: 0.0859 - val_acc: 0.9756

Epoch 9/20

60000/60000 [=====] - 5s 80us/step - loss: 0.0941 - acc: 0.9706 - val_loss: 0.0826 - val_acc: 0.9755

Epoch 10/20

60000/60000 [=====] - 5s 80us/step - loss: 0.0895 - acc: 0.9725 - val_loss: 0.0767 - val_acc: 0.9787

Epoch 11/20

60000/60000 [=====] - 5s 80us/step - loss: 0.0824 - acc: 0.9740 - val_loss: 0.0795 - val_acc: 0.9776

Epoch 12/20

60000/60000 [=====] - 5s 80us/step - loss: 0.0777 - acc: 0.9761 - val_loss: 0.0731 - val_acc: 0.9798

Epoch 13/20

60000/60000 [=====] - 5s 80us/step - loss: 0.0720 - acc: 0.9772 - val_loss: 0.0770 - val_acc: 0.9802

Epoch 14/20

60000/60000 [=====] - 5s 81us/step - loss: 0.0722 - acc: 0.9776 - val_loss: 0.0756 - val_acc: 0.9796

Epoch 15/20

60000/60000 [=====] - 5s 85us/step - loss: 0.0659 - acc: 0.9795 - val_loss: 0.0705 - val_acc: 0.9810

Epoch 16/20

60000/60000 [=====] - 5s 80us/step - loss: 0.0627 - acc: 0.9811 - val_loss: 0.0737 - val_acc: 0.9818

Epoch 17/20

60000/60000 [=====] - 6s 104us/step - loss: 0.0592 - acc: 0.9815 - val_loss: 0.0741 - val_acc: 0.9809

Epoch 18/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0537 - acc: 0.9834 - val_loss: 0.0764 - val_acc: 0.9798

Epoch 19/20

60000/60000 [=====] - 6s 103us/step - loss: 0.0539

- acc: 0.9831 - val_loss: 0.0730 - val_acc: 0.9814

Epoch 20/20

60000/60000 [=====] - 6s 93us/step - loss: 0.0542 -

acc: 0.9835 - val_loss: 0.0797 - val_acc: 0.9802

In [28]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

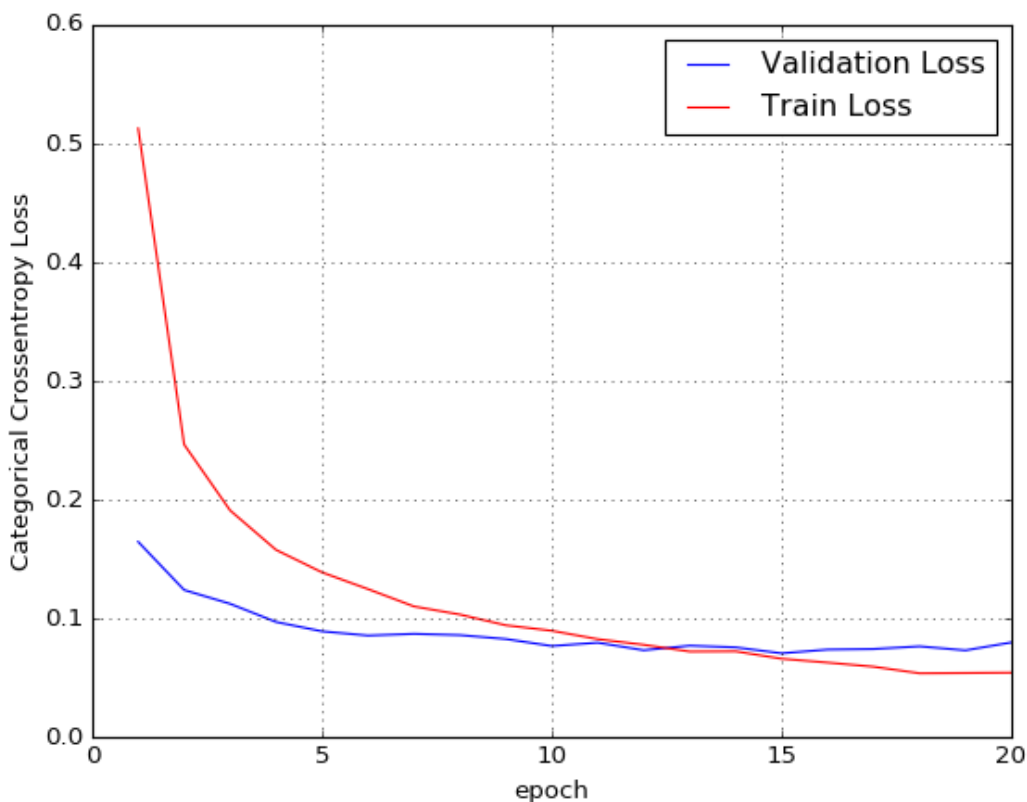
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.07972456046559565

Test accuracy: 0.9802



In [29]:

```

w_after = model_drop.get_weights()

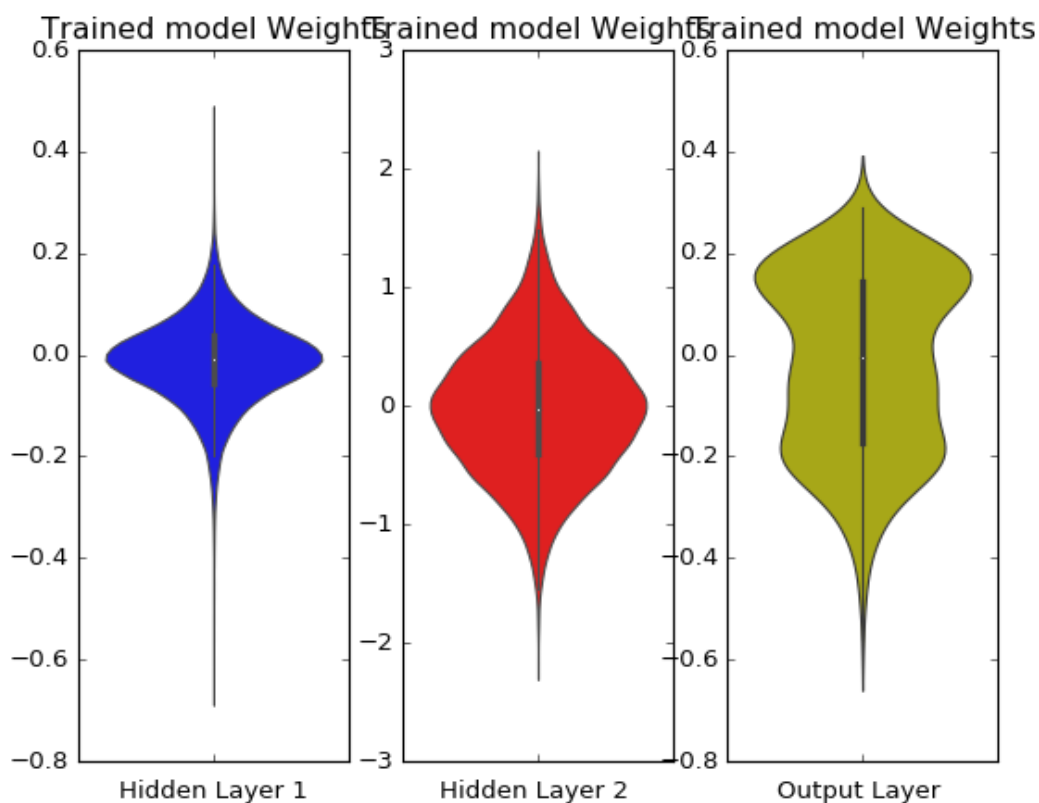
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



In []:

MLP + BatchNormalization + Dropout(rate = 0.5) + AdamOptimizer

with 2 hidden layers

In [30]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
=====		
dense_13 (Dense)	(None, 364)	285740
batch_normalization_3 (Batch Normalization)	(None, 364)	1456
dropout_5 (Dropout)	(None, 364)	0
dense_14 (Dense)	(None, 52)	18980
batch_normalization_4 (Batch Normalization)	(None, 52)	208
dropout_6 (Dropout)	(None, 52)	0
dense_15 (Dense)	(None, 10)	530
=====		
Total params: 306,914		
Trainable params: 306,082		
Non-trainable params: 832		

In [31]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 115us/step - loss: 0.6007
- acc: 0.8222 - val_loss: 0.1861 - val_acc: 0.9449

Epoch 2/20

60000/60000 [=====] - 6s 101us/step - loss: 0.3093
- acc: 0.9088 - val_loss: 0.1419 - val_acc: 0.9576

Epoch 3/20

60000/60000 [=====] - 6s 102us/step - loss: 0.2488
- acc: 0.9279 - val_loss: 0.1245 - val_acc: 0.9627

Epoch 4/20

60000/60000 [=====] - 6s 94us/step - loss: 0.2162 -
acc: 0.9373 - val_loss: 0.1098 - val_acc: 0.9646

Epoch 5/20

60000/60000 [=====] - 6s 96us/step - loss: 0.1912 -
acc: 0.9440 - val_loss: 0.0955 - val_acc: 0.9707

Epoch 6/20

60000/60000 [=====] - 6s 97us/step - loss: 0.1756 -
acc: 0.9489 - val_loss: 0.0954 - val_acc: 0.9714

Epoch 7/20

60000/60000 [=====] - 6s 100us/step - loss: 0.1659
- acc: 0.9525 - val_loss: 0.0879 - val_acc: 0.9751

Epoch 8/20

60000/60000 [=====] - 6s 93us/step - loss: 0.1504 -
acc: 0.9550 - val_loss: 0.0859 - val_acc: 0.9729

Epoch 9/20

60000/60000 [=====] - 6s 93us/step - loss: 0.1415 -
acc: 0.9587 - val_loss: 0.0836 - val_acc: 0.9736

Epoch 10/20

60000/60000 [=====] - 6s 97us/step - loss: 0.1365 -
acc: 0.9601 - val_loss: 0.0850 - val_acc: 0.9747

Epoch 11/20

60000/60000 [=====] - 7s 113us/step - loss: 0.1301
- acc: 0.9619 - val_loss: 0.0830 - val_acc: 0.9765

Epoch 12/20

60000/60000 [=====] - 6s 97us/step - loss: 0.1219 -
acc: 0.9637 - val_loss: 0.0820 - val_acc: 0.9755

Epoch 13/20

60000/60000 [=====] - 7s 115us/step - loss: 0.1196
- acc: 0.9650 - val_loss: 0.0756 - val_acc: 0.9775

Epoch 14/20

60000/60000 [=====] - 7s 114us/step - loss: 0.1112
- acc: 0.9674 - val_loss: 0.0745 - val_acc: 0.9778

Epoch 15/20

60000/60000 [=====] - 7s 114us/step - loss: 0.1071
- acc: 0.9686 - val_loss: 0.0714 - val_acc: 0.9795

Epoch 16/20

60000/60000 [=====] - 7s 110us/step - loss: 0.1020
- acc: 0.9700 - val_loss: 0.0752 - val_acc: 0.9776

Epoch 17/20

60000/60000 [=====] - 7s 110us/step - loss: 0.0977
- acc: 0.9708 - val_loss: 0.0815 - val_acc: 0.9775

Epoch 18/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0974
- acc: 0.9700 - val_loss: 0.0728 - val_acc: 0.9782

Epoch 19/20

60000/60000 [=====] - 6s 93us/step - loss: 0.0928 -

acc: 0.9723 - val_loss: 0.0688 - val_acc: 0.9795

Epoch 20/20

60000/60000 [=====] - 6s 97us/step - loss: 0.0912 -

acc: 0.9728 - val_loss: 0.0729 - val_acc: 0.9806

In [33]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

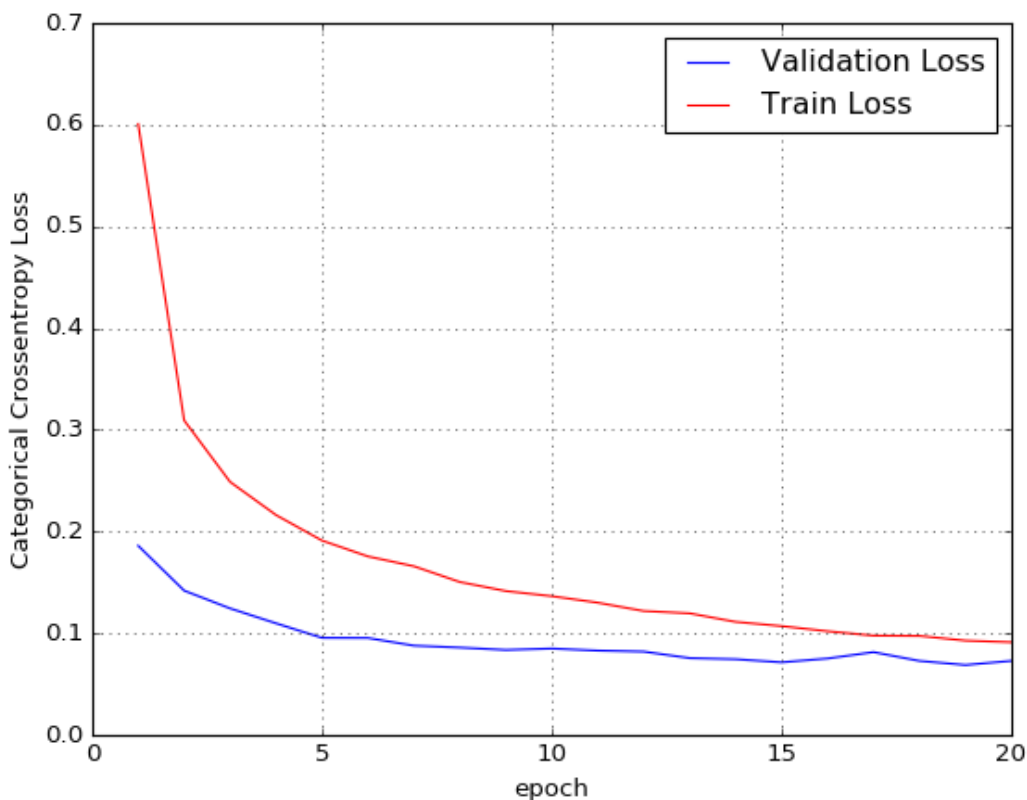
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.07286214664807195

Test accuracy: 0.9806



In [34]:

```

w_after = model_drop.get_weights()

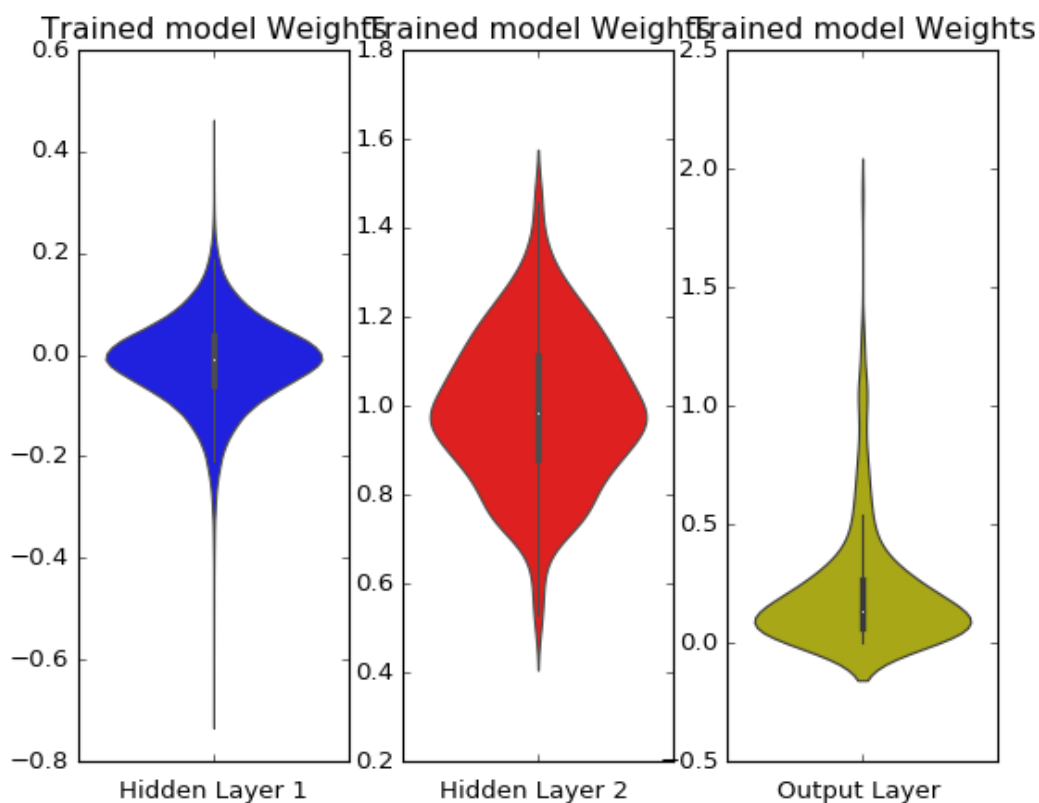
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



In []:

MLP + BatchNormalization + Dropout(rate = 0.3) + AdamOptimizer

with 2 hidden layers

In [36]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(364, activation='relu', input_shape=(input_dim,)), kernel_initializer=RandomNormal(mean=0.0, stddev=0.01))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.3))

model_drop.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.3))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
=====		
dense_16 (Dense)	(None, 364)	285740
batch_normalization_5 (Batch Normalization)	(None, 364)	1456
dropout_7 (Dropout)	(None, 364)	0
dense_17 (Dense)	(None, 52)	18980
batch_normalization_6 (Batch Normalization)	(None, 52)	208
dropout_8 (Dropout)	(None, 52)	0
dense_18 (Dense)	(None, 10)	530
=====		
Total params: 306,914		
Trainable params: 306,082		
Non-trainable params: 832		

In [37]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 114us/step - loss: 0.3831
- acc: 0.8881 - val_loss: 0.1457 - val_acc: 0.9558

Epoch 2/20

60000/60000 [=====] - 6s 105us/step - loss: 0.1861
- acc: 0.9443 - val_loss: 0.1056 - val_acc: 0.9669

Epoch 3/20

60000/60000 [=====] - 6s 98us/step - loss: 0.1469 -
acc: 0.9556 - val_loss: 0.0942 - val_acc: 0.9686

Epoch 4/20

60000/60000 [=====] - 5s 91us/step - loss: 0.1253 -
acc: 0.9615 - val_loss: 0.0818 - val_acc: 0.9751

Epoch 5/20

60000/60000 [=====] - 6s 97us/step - loss: 0.1083 -
acc: 0.9672 - val_loss: 0.0777 - val_acc: 0.9756

Epoch 6/20

60000/60000 [=====] - 5s 89us/step - loss: 0.0995 -
acc: 0.9694 - val_loss: 0.0744 - val_acc: 0.9778

Epoch 7/20

60000/60000 [=====] - 5s 87us/step - loss: 0.0892 -
acc: 0.9716 - val_loss: 0.0752 - val_acc: 0.9778

Epoch 8/20

60000/60000 [=====] - 5s 86us/step - loss: 0.0813 -
acc: 0.9748 - val_loss: 0.0697 - val_acc: 0.9776

Epoch 9/20

60000/60000 [=====] - 5s 90us/step - loss: 0.0758 -
acc: 0.9757 - val_loss: 0.0642 - val_acc: 0.9811

Epoch 10/20

60000/60000 [=====] - 6s 93us/step - loss: 0.0705 -
acc: 0.9784 - val_loss: 0.0641 - val_acc: 0.9812

Epoch 11/20

60000/60000 [=====] - 5s 91us/step - loss: 0.0664 -
acc: 0.9789 - val_loss: 0.0638 - val_acc: 0.9805

Epoch 12/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0636
- acc: 0.9799 - val_loss: 0.0686 - val_acc: 0.9791

Epoch 13/20

60000/60000 [=====] - 6s 94us/step - loss: 0.0583 -
acc: 0.9817 - val_loss: 0.0608 - val_acc: 0.9807

Epoch 14/20

60000/60000 [=====] - 6s 92us/step - loss: 0.0560 -
acc: 0.9823 - val_loss: 0.0638 - val_acc: 0.9823

Epoch 15/20

60000/60000 [=====] - 6s 93us/step - loss: 0.0535 -
acc: 0.9825 - val_loss: 0.0649 - val_acc: 0.9817

Epoch 16/20

60000/60000 [=====] - 5s 91us/step - loss: 0.0503 -
acc: 0.9842 - val_loss: 0.0667 - val_acc: 0.9808

Epoch 17/20

60000/60000 [=====] - 6s 104us/step - loss: 0.0509
- acc: 0.9833 - val_loss: 0.0618 - val_acc: 0.9816

Epoch 18/20

60000/60000 [=====] - 6s 98us/step - loss: 0.0484 -
acc: 0.9847 - val_loss: 0.0572 - val_acc: 0.9831

Epoch 19/20

60000/60000 [=====] - 6s 97us/step - loss: 0.0456 -

acc: 0.9844 - val_loss: 0.0655 - val_acc: 0.9818

Epoch 20/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0412 -

acc: 0.9859 - val_loss: 0.0644 - val_acc: 0.9825

In [38]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

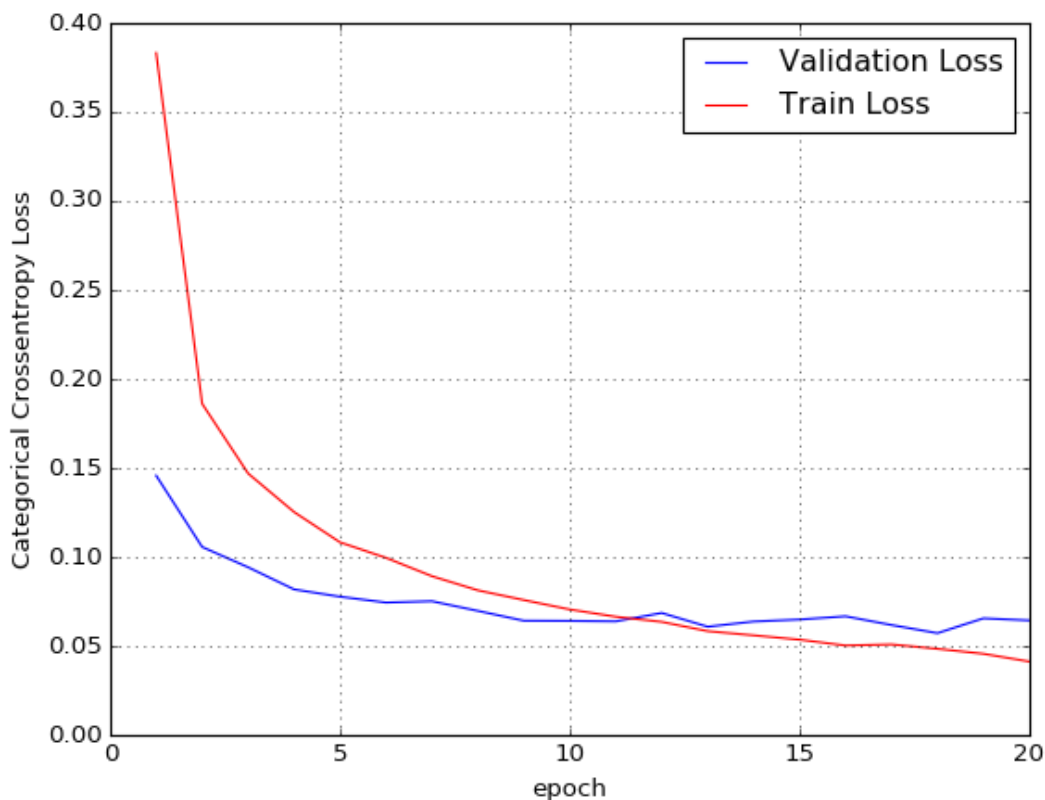
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06437753892600304

Test accuracy: 0.9825



In [39]:

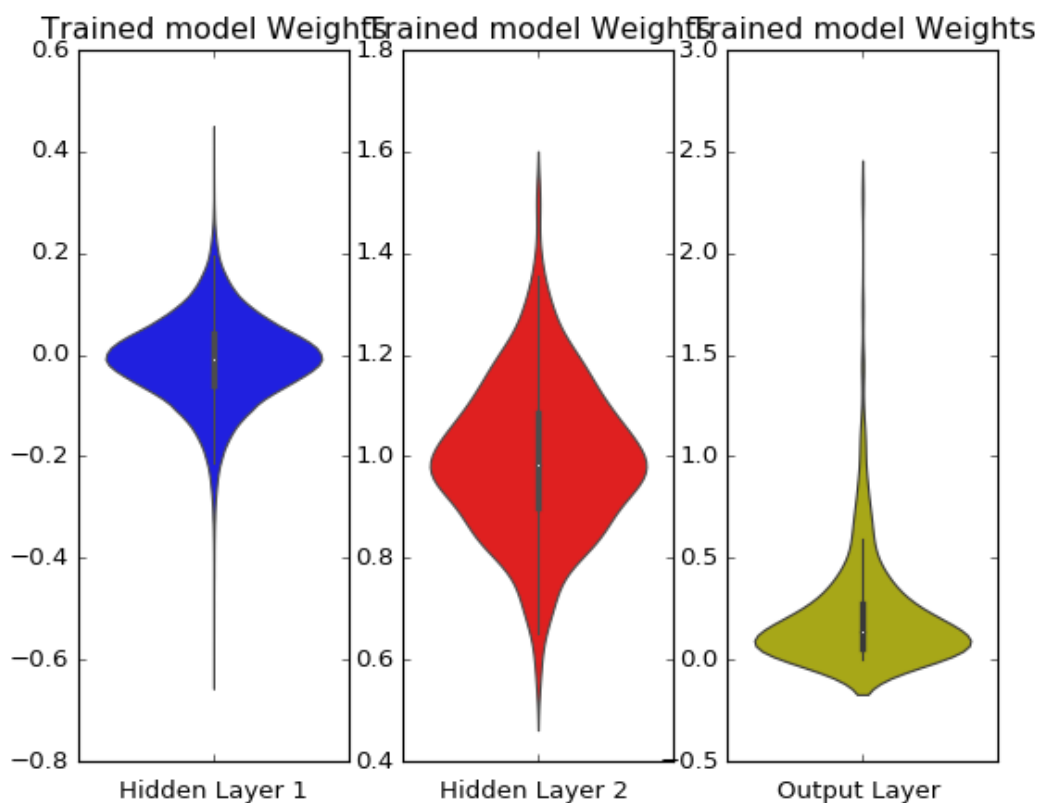
```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + ReLU + ADAM with 3 layers without Dropout and Batch Normalisation

In [40]:

```

model_relu = Sequential()
model_relu.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_relu.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)

```

Layer (type)	Output Shape	Param #
dense_19 (Dense)	(None, 364)	285740
dense_20 (Dense)	(None, 128)	46720
dense_21 (Dense)	(None, 52)	6708
dense_22 (Dense)	(None, 10)	530

=====
 Total params: 339,698
 Trainable params: 339,698
 Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 92us/step - loss: 0.2579 - acc: 0.9240 - val_loss: 0.1299 - val_acc: 0.9621

Epoch 2/20

60000/60000 [=====] - 5s 79us/step - loss: 0.0937 - acc: 0.9717 - val_loss: 0.0970 - val_acc: 0.9708

Epoch 3/20

60000/60000 [=====] - 4s 70us/step - loss: 0.0619 - acc: 0.9809 - val_loss: 0.0834 - val_acc: 0.9741

Epoch 4/20

60000/60000 [=====] - 4s 68us/step - loss: 0.0455 - acc: 0.9858 - val_loss: 0.0703 - val_acc: 0.9798

Epoch 5/20

60000/60000 [=====] - 5s 82us/step - loss: 0.0323 - acc: 0.9893 - val_loss: 0.0714 - val_acc: 0.9778

Epoch 6/20

60000/60000 [=====] - 5s 82us/step - loss: 0.0235 - acc: 0.9924 - val_loss: 0.0888 - val_acc: 0.9745

Epoch 7/20

60000/60000 [=====] - 5s 78us/step - loss: 0.0218 - acc: 0.9928 - val_loss: 0.0952 - val_acc: 0.9740

Epoch 8/20

60000/60000 [=====] - 5s 77us/step - loss: 0.0196 - acc: 0.9939 - val_loss: 0.0866 - val_acc: 0.9768

Epoch 9/20

60000/60000 [=====] - 5s 78us/step - loss: 0.0147 - acc: 0.9951 - val_loss: 0.0988 - val_acc: 0.9772

Epoch 10/20

60000/60000 [=====] - 6s 92us/step - loss: 0.0155 -

```
acc: 0.9947 - val_loss: 0.0893 - val_acc: 0.9792
Epoch 11/20
60000/60000 [=====] - 6s 106us/step - loss: 0.0159
- acc: 0.9945 - val_loss: 0.1069 - val_acc: 0.9745
Epoch 12/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0129 -
acc: 0.9958 - val_loss: 0.0797 - val_acc: 0.9806
Epoch 13/20
60000/60000 [=====] - 5s 77us/step - loss: 0.0122 -
acc: 0.9958 - val_loss: 0.0905 - val_acc: 0.9810
Epoch 14/20
60000/60000 [=====] - 5s 77us/step - loss: 0.0128 -
acc: 0.9958 - val_loss: 0.1108 - val_acc: 0.9765
Epoch 15/20
60000/60000 [=====] - 5s 77us/step - loss: 0.0087 -
acc: 0.9971 - val_loss: 0.0922 - val_acc: 0.9805
Epoch 16/20
60000/60000 [=====] - 5s 77us/step - loss: 0.0075 -
acc: 0.9977 - val_loss: 0.1271 - val_acc: 0.9749
Epoch 17/20
60000/60000 [=====] - 5s 77us/step - loss: 0.0132 -
acc: 0.9959 - val_loss: 0.0947 - val_acc: 0.9799
Epoch 18/20
60000/60000 [=====] - 5s 79us/step - loss: 0.0098 -
acc: 0.9967 - val_loss: 0.0982 - val_acc: 0.9793
Epoch 19/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0089 -
acc: 0.9970 - val_loss: 0.0940 - val_acc: 0.9806
Epoch 20/20
60000/60000 [=====] - 5s 80us/step - loss: 0.0095 -
acc: 0.9970 - val_loss: 0.1108 - val_acc: 0.9788
```

In [41]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

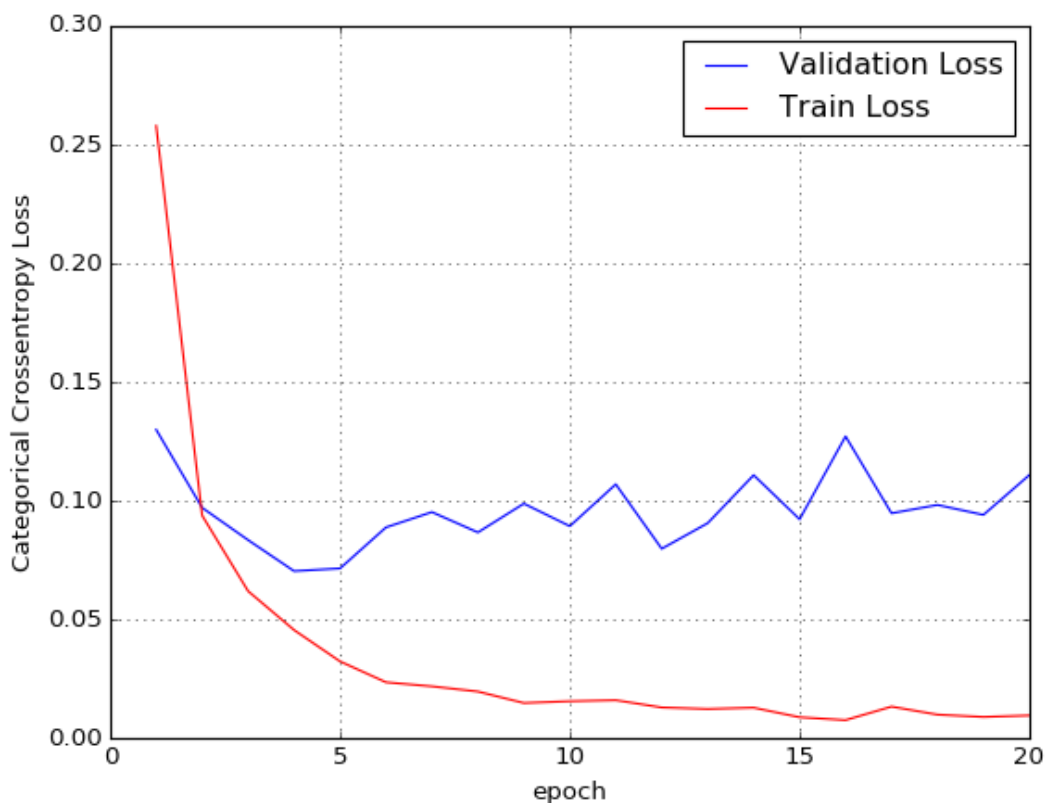
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# Loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.1108168175617167

Test accuracy: 0.9788



In [43]:

```
w_after = model_relu.get_weights()

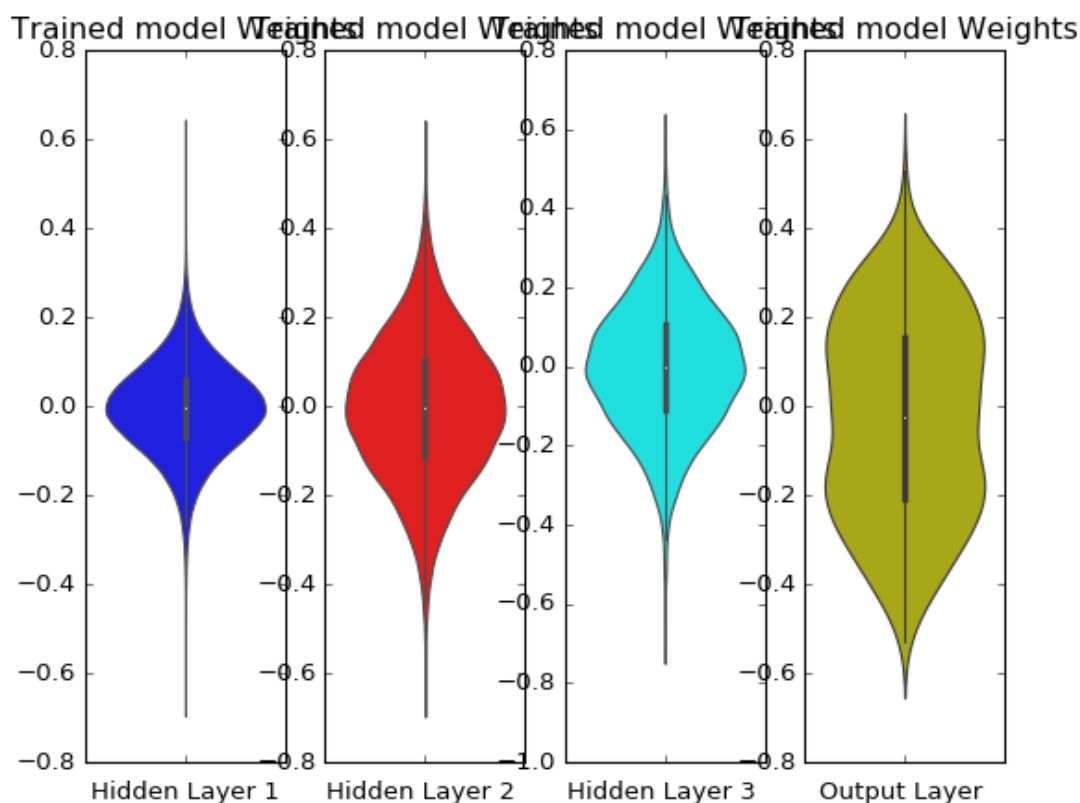
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='cyan')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Batch-Norm on 3 hidden Layers + AdamOptimizer

In [44]:

```
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i + n_{i+1})}$ 
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(0, \sigma) = N(0, 0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(0, \sigma) = N(0, 0.055)$ 
# h3 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, std=0.039)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std=0.055)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std=0.120)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Layer (type)	Output Shape	Param #
dense_23 (Dense)	(None, 364)	285740
batch_normalization_7 (Batch Normalization)	(None, 364)	1456
dense_24 (Dense)	(None, 128)	46720
batch_normalization_8 (Batch Normalization)	(None, 128)	512
dense_25 (Dense)	(None, 52)	6708
batch_normalization_9 (Batch Normalization)	(None, 52)	208
dense_26 (Dense)	(None, 10)	530
Total params: 341,874		
Trainable params: 340,786		
Non-trainable params: 1,088		

In [45]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 110us/step - loss: 0.2334
- acc: 0.9338 - val_loss: 0.1146 - val_acc: 0.9651

Epoch 2/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0870
- acc: 0.9744 - val_loss: 0.0894 - val_acc: 0.9718

Epoch 3/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0557
- acc: 0.9826 - val_loss: 0.0884 - val_acc: 0.9715

Epoch 4/20

60000/60000 [=====] - 6s 97us/step - loss: 0.0400 -
acc: 0.9876 - val_loss: 0.0823 - val_acc: 0.9760

Epoch 5/20

60000/60000 [=====] - 6s 92us/step - loss: 0.0318 -
acc: 0.9901 - val_loss: 0.0830 - val_acc: 0.9758

Epoch 6/20

60000/60000 [=====] - 7s 111us/step - loss: 0.0281
- acc: 0.9907 - val_loss: 0.0779 - val_acc: 0.9770

Epoch 7/20

60000/60000 [=====] - 6s 102us/step - loss: 0.0226
- acc: 0.9930 - val_loss: 0.0812 - val_acc: 0.9746

Epoch 8/20

60000/60000 [=====] - 7s 124us/step - loss: 0.0202
- acc: 0.9934 - val_loss: 0.0760 - val_acc: 0.9788

Epoch 9/20

60000/60000 [=====] - 7s 122us/step - loss: 0.0170
- acc: 0.9944 - val_loss: 0.0818 - val_acc: 0.9777

Epoch 10/20

60000/60000 [=====] - 6s 108us/step - loss: 0.0150
- acc: 0.9949 - val_loss: 0.0788 - val_acc: 0.9778

Epoch 11/20

60000/60000 [=====] - 6s 105us/step - loss: 0.0130
- acc: 0.9959 - val_loss: 0.0798 - val_acc: 0.9788

Epoch 12/20

60000/60000 [=====] - 6s 102us/step - loss: 0.0159
- acc: 0.9948 - val_loss: 0.0847 - val_acc: 0.9777

Epoch 13/20

60000/60000 [=====] - 5s 91us/step - loss: 0.0121 -
acc: 0.9961 - val_loss: 0.0773 - val_acc: 0.9794

Epoch 14/20

60000/60000 [=====] - 5s 90us/step - loss: 0.0115 -
acc: 0.9960 - val_loss: 0.0802 - val_acc: 0.9798

Epoch 15/20

60000/60000 [=====] - 6s 95us/step - loss: 0.0091 -
acc: 0.9970 - val_loss: 0.0752 - val_acc: 0.9803

Epoch 16/20

60000/60000 [=====] - 6s 92us/step - loss: 0.0124 -
acc: 0.9956 - val_loss: 0.1012 - val_acc: 0.9752

Epoch 17/20

60000/60000 [=====] - 5s 89us/step - loss: 0.0112 -
acc: 0.9961 - val_loss: 0.0907 - val_acc: 0.9775

Epoch 18/20

60000/60000 [=====] - 6s 106us/step - loss: 0.0101
- acc: 0.9966 - val_loss: 0.0792 - val_acc: 0.9799

Epoch 19/20

60000/60000 [=====] - 7s 115us/step - loss: 0.0064

- acc: 0.9980 - val_loss: 0.0808 - val_acc: 0.9787

Epoch 20/20

60000/60000 [=====] - 7s 115us/step - loss: 0.0085

- acc: 0.9972 - val_loss: 0.0867 - val_acc: 0.9795

In [46]:

```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

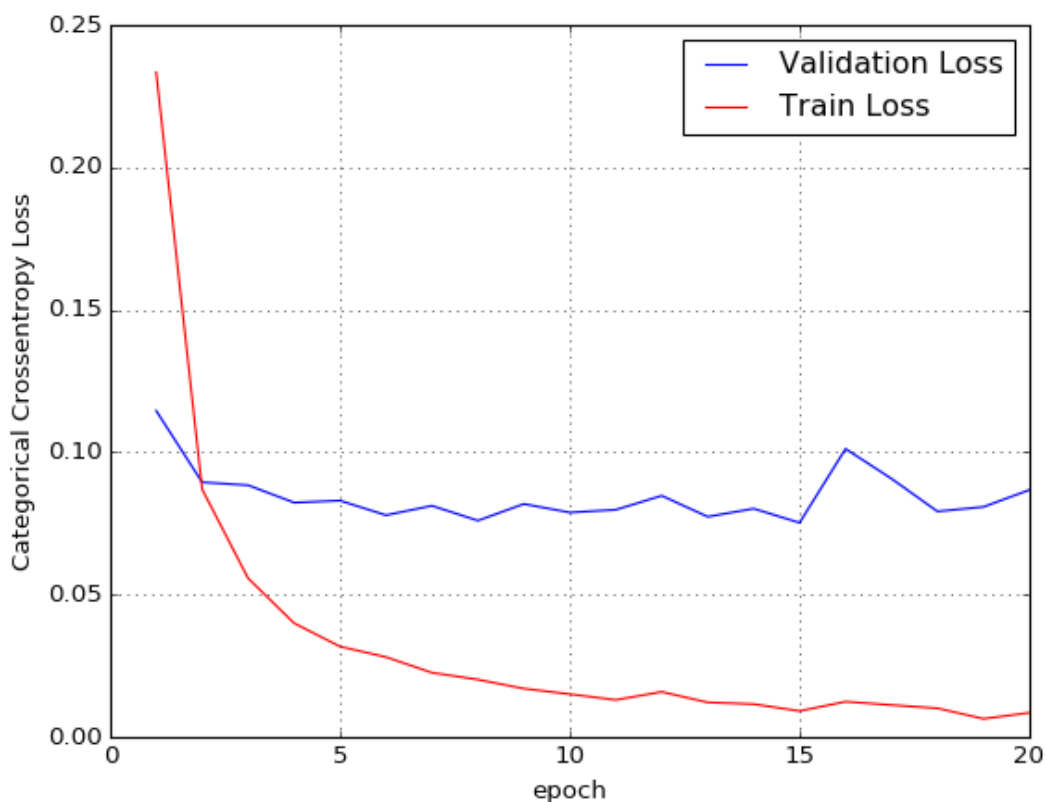
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# Loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08674843371730385

Test accuracy: 0.9795



In [47]:

```
w_after = model_batch.get_weights()

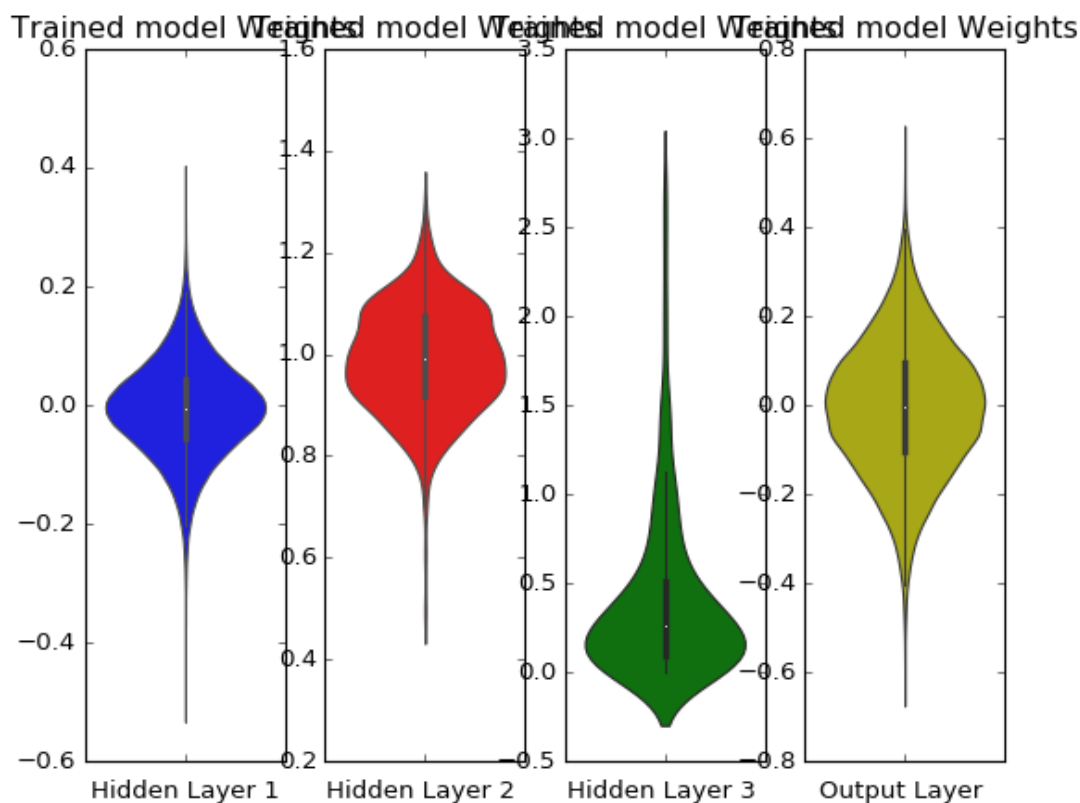
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Dropout(rate = 0.5) + AdamOptimizer with 3 hidden layers

In [48]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-funct
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=R
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stdc
model_drop.add(Dropout(0.5))

model_drop.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stdde
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_27 (Dense)	(None, 364)	285740
dropout_9 (Dropout)	(None, 364)	0
dense_28 (Dense)	(None, 128)	46720
dropout_10 (Dropout)	(None, 128)	0
dense_29 (Dense)	(None, 52)	6708
dropout_11 (Dropout)	(None, 52)	0
dense_30 (Dense)	(None, 10)	530
=====	=====	=====
Total params: 339,698		
Trainable params: 339,698		
Non-trainable params: 0		

In [49]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 100us/step - loss: 1.3778
- acc: 0.5752 - val_loss: 0.3179 - val_acc: 0.9180

Epoch 2/20

60000/60000 [=====] - 5s 91us/step - loss: 0.5697 -
acc: 0.8389 - val_loss: 0.2357 - val_acc: 0.9406

Epoch 3/20

60000/60000 [=====] - 6s 106us/step - loss: 0.4437
- acc: 0.8838 - val_loss: 0.2002 - val_acc: 0.9455

Epoch 4/20

60000/60000 [=====] - 7s 121us/step - loss: 0.3701
- acc: 0.9034 - val_loss: 0.1784 - val_acc: 0.9543

Epoch 5/20

60000/60000 [=====] - 6s 103us/step - loss: 0.3346
- acc: 0.9151 - val_loss: 0.1610 - val_acc: 0.9578

Epoch 6/20

60000/60000 [=====] - 6s 94us/step - loss: 0.3017 -
acc: 0.9235 - val_loss: 0.1637 - val_acc: 0.9576

Epoch 7/20

60000/60000 [=====] - 6s 93us/step - loss: 0.2769 -
acc: 0.9303 - val_loss: 0.1512 - val_acc: 0.9607

Epoch 8/20

60000/60000 [=====] - 6s 100us/step - loss: 0.2577
- acc: 0.9356 - val_loss: 0.1434 - val_acc: 0.9621

Epoch 9/20

60000/60000 [=====] - 6s 101us/step - loss: 0.2455
- acc: 0.9385 - val_loss: 0.1340 - val_acc: 0.9660

Epoch 10/20

60000/60000 [=====] - 6s 98us/step - loss: 0.2290 -
acc: 0.9429 - val_loss: 0.1325 - val_acc: 0.9661

Epoch 11/20

60000/60000 [=====] - 6s 106us/step - loss: 0.2208
- acc: 0.9442 - val_loss: 0.1257 - val_acc: 0.9683

Epoch 12/20

60000/60000 [=====] - 6s 106us/step - loss: 0.2071
- acc: 0.9477 - val_loss: 0.1286 - val_acc: 0.9702

Epoch 13/20

60000/60000 [=====] - 6s 101us/step - loss: 0.2008
- acc: 0.9485 - val_loss: 0.1258 - val_acc: 0.9699

Epoch 14/20

60000/60000 [=====] - 6s 99us/step - loss: 0.1956 -
acc: 0.9515 - val_loss: 0.1284 - val_acc: 0.9693

Epoch 15/20

60000/60000 [=====] - 7s 118us/step - loss: 0.1862
- acc: 0.9537 - val_loss: 0.1253 - val_acc: 0.9703

Epoch 16/20

60000/60000 [=====] - 6s 107us/step - loss: 0.1797
- acc: 0.9547 - val_loss: 0.1183 - val_acc: 0.9694

Epoch 17/20

60000/60000 [=====] - 7s 114us/step - loss: 0.1766
- acc: 0.9546 - val_loss: 0.1206 - val_acc: 0.9718

Epoch 18/20

60000/60000 [=====] - 7s 117us/step - loss: 0.1697
- acc: 0.9570 - val_loss: 0.1203 - val_acc: 0.9703

Epoch 19/20

60000/60000 [=====] - 6s 105us/step - loss: 0.1644

- acc: 0.9590 - val_loss: 0.1134 - val_acc: 0.9734

Epoch 20/20

60000/60000 [=====] - 6s 97us/step - loss: 0.1611 -

acc: 0.9606 - val_loss: 0.1137 - val_acc: 0.9741

In [50]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

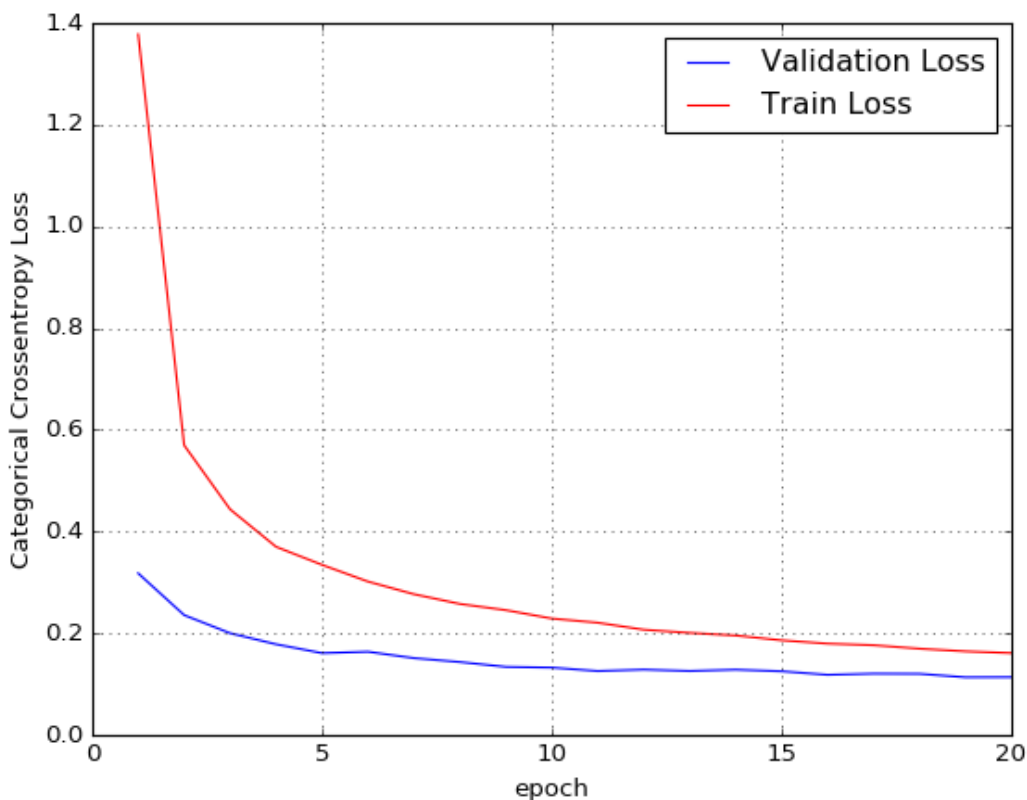
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# Loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.11365028675964568

Test accuracy: 0.9741



In [51]:

```
w_after = model_drop.get_weights()

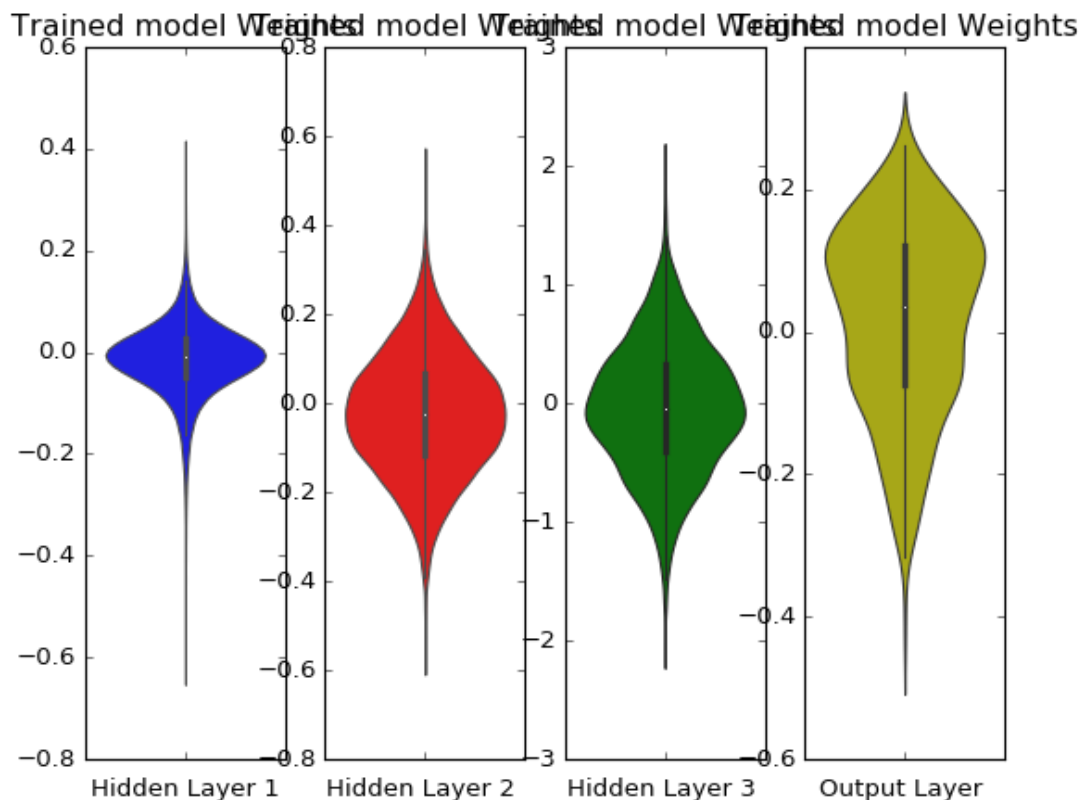
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Dropout(rate = 0.3) + AdamOptimizer with 3 hidden layers

In [52]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-funct
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=R
model_drop.add(Dropout(0.3))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stdde
model_drop.add(Dropout(0.3))

model_drop.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stdde
model_drop.add(Dropout(0.3))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_31 (Dense)	(None, 364)	285740
dropout_12 (Dropout)	(None, 364)	0
dense_32 (Dense)	(None, 128)	46720
dropout_13 (Dropout)	(None, 128)	0
dense_33 (Dense)	(None, 52)	6708
dropout_14 (Dropout)	(None, 52)	0
dense_34 (Dense)	(None, 10)	530
=====	=====	=====
Total params: 339,698		
Trainable params: 339,698		
Non-trainable params: 0		

In [53]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 113us/step - loss: 0.6964
- acc: 0.7865 - val_loss: 0.1897 - val_acc: 0.9470

Epoch 2/20

60000/60000 [=====] - 6s 108us/step - loss: 0.2802
- acc: 0.9218 - val_loss: 0.1425 - val_acc: 0.9601

Epoch 3/20

60000/60000 [=====] - 6s 96us/step - loss: 0.2134 -
acc: 0.9400 - val_loss: 0.1135 - val_acc: 0.9674

Epoch 4/20

60000/60000 [=====] - 5s 81us/step - loss: 0.1782 -
acc: 0.9506 - val_loss: 0.1015 - val_acc: 0.9721

Epoch 5/20

60000/60000 [=====] - 5s 79us/step - loss: 0.1472 -
acc: 0.9601 - val_loss: 0.0982 - val_acc: 0.9723

Epoch 6/20

60000/60000 [=====] - 6s 95us/step - loss: 0.1285 -
acc: 0.9646 - val_loss: 0.0883 - val_acc: 0.9757

Epoch 7/20

60000/60000 [=====] - 6s 95us/step - loss: 0.1164 -
acc: 0.9681 - val_loss: 0.0880 - val_acc: 0.9763

Epoch 8/20

60000/60000 [=====] - 6s 93us/step - loss: 0.1076 -
acc: 0.9702 - val_loss: 0.0868 - val_acc: 0.9761

Epoch 9/20

60000/60000 [=====] - 6s 98us/step - loss: 0.0982 -
acc: 0.9726 - val_loss: 0.0881 - val_acc: 0.9772

Epoch 10/20

60000/60000 [=====] - 6s 105us/step - loss: 0.0911
- acc: 0.9751 - val_loss: 0.0777 - val_acc: 0.9785

Epoch 11/20

60000/60000 [=====] - 6s 94us/step - loss: 0.0851 -
acc: 0.9759 - val_loss: 0.0810 - val_acc: 0.9799

Epoch 12/20

60000/60000 [=====] - 6s 102us/step - loss: 0.0790
- acc: 0.9782 - val_loss: 0.0853 - val_acc: 0.9786

Epoch 13/20

60000/60000 [=====] - 6s 94us/step - loss: 0.0719 -
acc: 0.9791 - val_loss: 0.0756 - val_acc: 0.9799

Epoch 14/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0716 -
acc: 0.9803 - val_loss: 0.0721 - val_acc: 0.9794

Epoch 15/20

60000/60000 [=====] - 6s 101us/step - loss: 0.0699
- acc: 0.9804 - val_loss: 0.0834 - val_acc: 0.9789

Epoch 16/20

60000/60000 [=====] - 6s 94us/step - loss: 0.0635 -
acc: 0.9818 - val_loss: 0.0845 - val_acc: 0.9799

Epoch 17/20

60000/60000 [=====] - 6s 93us/step - loss: 0.0625 -
acc: 0.9828 - val_loss: 0.0773 - val_acc: 0.9803

Epoch 18/20

60000/60000 [=====] - 6s 104us/step - loss: 0.0570
- acc: 0.9835 - val_loss: 0.0743 - val_acc: 0.9823

Epoch 19/20

60000/60000 [=====] - 6s 96us/step - loss: 0.0564 -

acc: 0.9842 - val_loss: 0.0719 - val_acc: 0.9811

Epoch 20/20

60000/60000 [=====] - 5s 82us/step - loss: 0.0509 -

acc: 0.9853 - val_loss: 0.0763 - val_acc: 0.9806

In [55]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

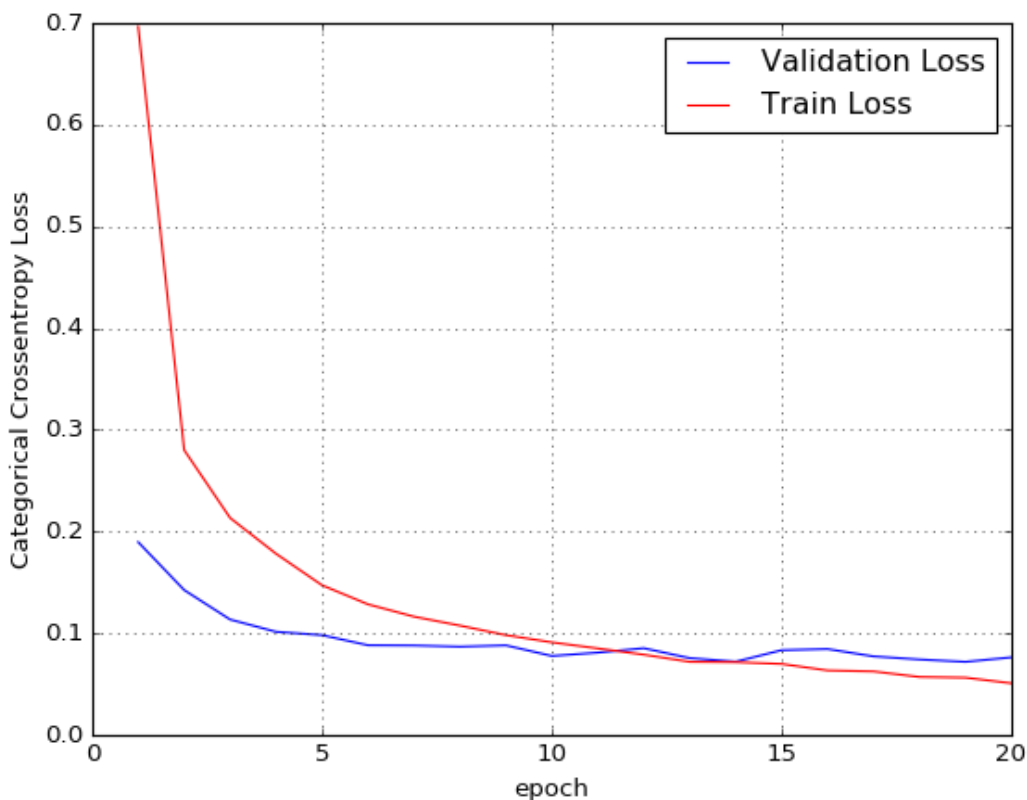
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.0763443440160543

Test accuracy: 0.9806



In [56]:

```

w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

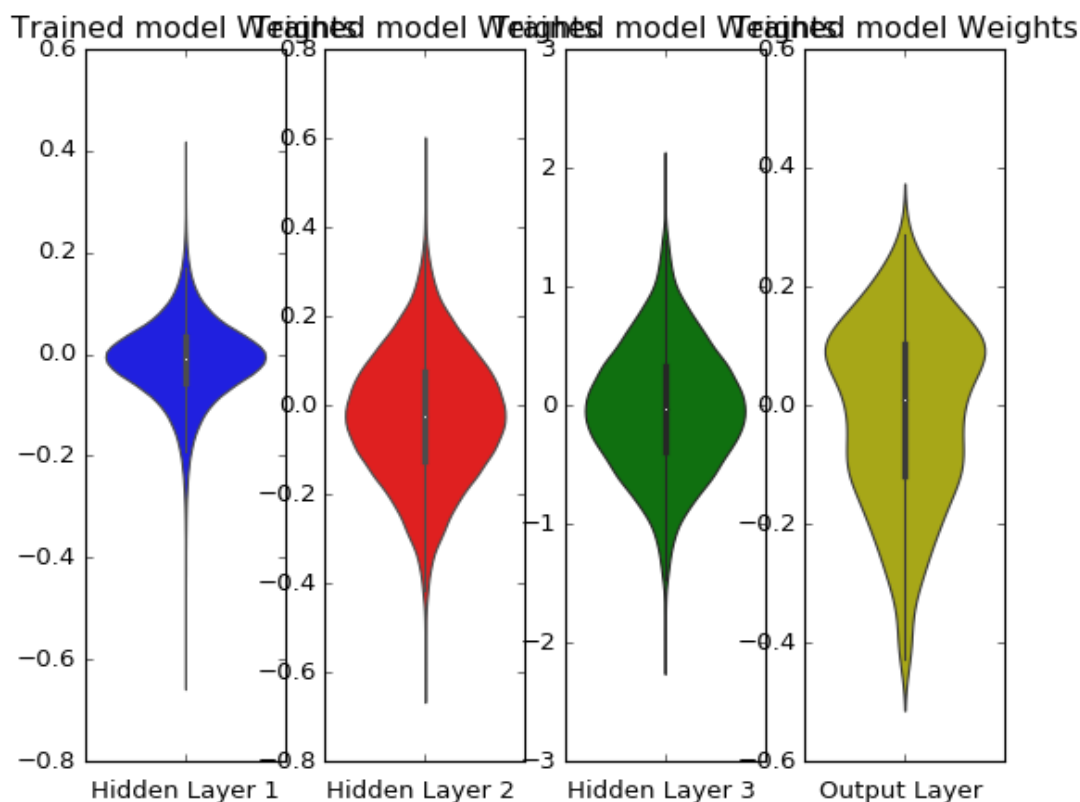
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



MLP + BatchNormalization + Dropout(rate = 0.5) + AdamOptimizer with 3 hidden layers

In [57]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-funct
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=R
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stdc
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stdde
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_35 (Dense)	(None, 364)	285740
batch_normalization_10 (Batch Normalization)	(None, 364)	1456
dropout_15 (Dropout)	(None, 364)	0
dense_36 (Dense)	(None, 128)	46720
batch_normalization_11 (Batch Normalization)	(None, 128)	512
dropout_16 (Dropout)	(None, 128)	0
dense_37 (Dense)	(None, 52)	6708
batch_normalization_12 (Batch Normalization)	(None, 52)	208
dropout_17 (Dropout)	(None, 52)	0
dense_38 (Dense)	(None, 10)	530
Total params: 341,874		
Trainable params: 340,786		
Non-trainable params: 1,088		

In [58]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 133us/step - loss: 0.9536
- acc: 0.7001 - val_loss: 0.2521 - val_acc: 0.9256

Epoch 2/20

60000/60000 [=====] - 6s 107us/step - loss: 0.4609
- acc: 0.8654 - val_loss: 0.1837 - val_acc: 0.9456

Epoch 3/20

60000/60000 [=====] - 7s 108us/step - loss: 0.3600
- acc: 0.8969 - val_loss: 0.1557 - val_acc: 0.9529

Epoch 4/20

60000/60000 [=====] - 6s 105us/step - loss: 0.2992
- acc: 0.9157 - val_loss: 0.1426 - val_acc: 0.9560

Epoch 5/20

60000/60000 [=====] - 7s 110us/step - loss: 0.2653
- acc: 0.9260 - val_loss: 0.1260 - val_acc: 0.9624

Epoch 6/20

60000/60000 [=====] - 7s 113us/step - loss: 0.2374
- acc: 0.9345 - val_loss: 0.1145 - val_acc: 0.9660

Epoch 7/20

60000/60000 [=====] - 7s 111us/step - loss: 0.2146
- acc: 0.9404 - val_loss: 0.1085 - val_acc: 0.9691

Epoch 8/20

60000/60000 [=====] - 7s 111us/step - loss: 0.1971
- acc: 0.9455 - val_loss: 0.1016 - val_acc: 0.9689

Epoch 9/20

60000/60000 [=====] - 7s 111us/step - loss: 0.1864
- acc: 0.9488 - val_loss: 0.0927 - val_acc: 0.9721

Epoch 10/20

60000/60000 [=====] - 7s 112us/step - loss: 0.1710
- acc: 0.9527 - val_loss: 0.0948 - val_acc: 0.9714

Epoch 11/20

60000/60000 [=====] - 7s 115us/step - loss: 0.1629
- acc: 0.9552 - val_loss: 0.0933 - val_acc: 0.9722

Epoch 12/20

60000/60000 [=====] - 7s 112us/step - loss: 0.1631
- acc: 0.9558 - val_loss: 0.0903 - val_acc: 0.9742

Epoch 13/20

60000/60000 [=====] - 7s 111us/step - loss: 0.1505
- acc: 0.9584 - val_loss: 0.0852 - val_acc: 0.9749

Epoch 14/20

60000/60000 [=====] - 8s 136us/step - loss: 0.1406
- acc: 0.9608 - val_loss: 0.0831 - val_acc: 0.9766

Epoch 15/20

60000/60000 [=====] - 8s 130us/step - loss: 0.1397
- acc: 0.9609 - val_loss: 0.0859 - val_acc: 0.9755

Epoch 16/20

60000/60000 [=====] - 9s 151us/step - loss: 0.1312
- acc: 0.9641 - val_loss: 0.0792 - val_acc: 0.9766

Epoch 17/20

60000/60000 [=====] - 8s 137us/step - loss: 0.1305
- acc: 0.9632 - val_loss: 0.0803 - val_acc: 0.9778

Epoch 18/20

60000/60000 [=====] - 8s 128us/step - loss: 0.1257
- acc: 0.9645 - val_loss: 0.0776 - val_acc: 0.9787

Epoch 19/20

60000/60000 [=====] - 8s 126us/step - loss: 0.1207

- acc: 0.9661 - val_loss: 0.0778 - val_acc: 0.9782

Epoch 20/20

60000/60000 [=====] - 9s 148us/step - loss: 0.1140

- acc: 0.9693 - val_loss: 0.0768 - val_acc: 0.9797

In [59]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

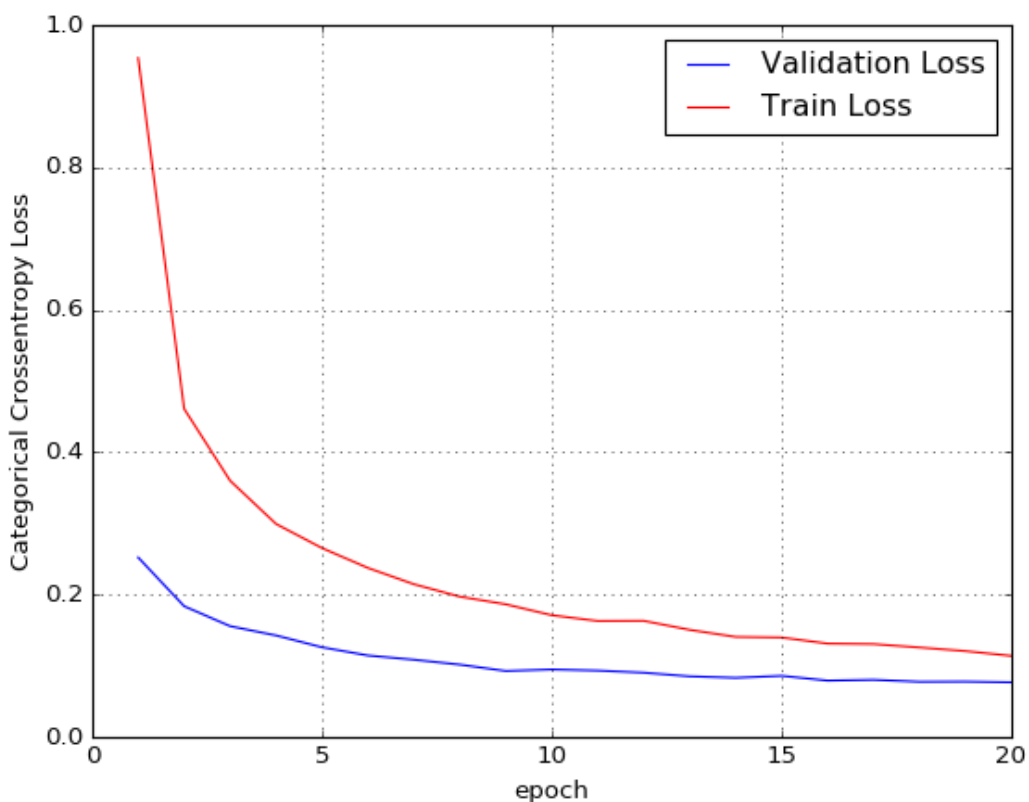
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.07683449132365058

Test accuracy: 0.9797



In [60]:

```
w_after = model_drop.get_weights()

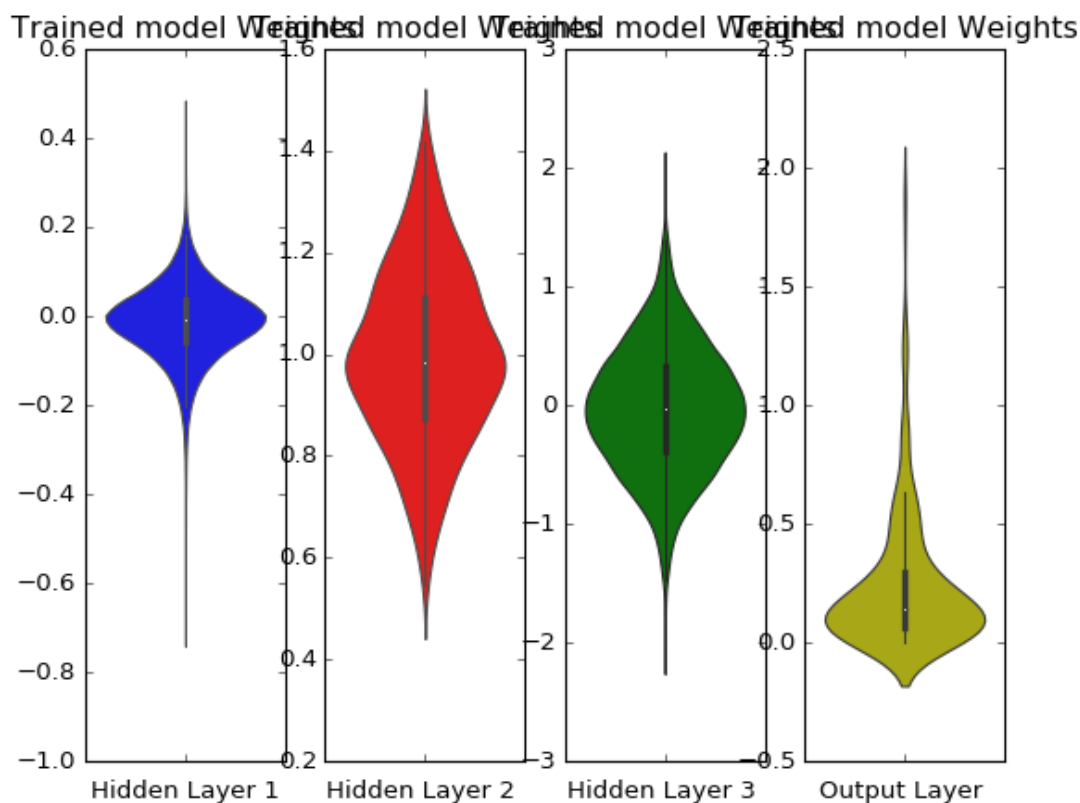
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + BatchNormalization + Dropout(rate = 0.3) + AdamOptimizer with 3 hidden layers

In [61]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-funct
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=R
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.3))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stdc
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.3))

model_drop.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stdde
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.3))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_39 (Dense)	(None, 364)	285740
batch_normalization_13 (Batch Normalization)	(None, 364)	1456
dropout_18 (Dropout)	(None, 364)	0
dense_40 (Dense)	(None, 128)	46720
batch_normalization_14 (Batch Normalization)	(None, 128)	512
dropout_19 (Dropout)	(None, 128)	0
dense_41 (Dense)	(None, 52)	6708
batch_normalization_15 (Batch Normalization)	(None, 52)	208
dropout_20 (Dropout)	(None, 52)	0
dense_42 (Dense)	(None, 10)	530
Total params: 341,874		
Trainable params: 340,786		
Non-trainable params: 1,088		

In [62]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 9s 145us/step - loss: 0.5354
- acc: 0.8428 - val_loss: 0.1646 - val_acc: 0.9491

Epoch 2/20

60000/60000 [=====] - 8s 131us/step - loss: 0.2464
- acc: 0.9281 - val_loss: 0.1239 - val_acc: 0.9613

Epoch 3/20

60000/60000 [=====] - 9s 155us/step - loss: 0.1866
- acc: 0.9460 - val_loss: 0.1056 - val_acc: 0.9687

Epoch 4/20

60000/60000 [=====] - 9s 156us/step - loss: 0.1574
- acc: 0.9528 - val_loss: 0.0967 - val_acc: 0.9694

Epoch 5/20

60000/60000 [=====] - 9s 158us/step - loss: 0.1442
- acc: 0.9573 - val_loss: 0.0887 - val_acc: 0.9729

Epoch 6/20

60000/60000 [=====] - 9s 158us/step - loss: 0.1263
- acc: 0.9619 - val_loss: 0.0780 - val_acc: 0.9776

Epoch 7/20

60000/60000 [=====] - 8s 138us/step - loss: 0.1125
- acc: 0.9672 - val_loss: 0.0806 - val_acc: 0.9757

Epoch 8/20

60000/60000 [=====] - 9s 145us/step - loss: 0.1040
- acc: 0.9682 - val_loss: 0.0790 - val_acc: 0.9764

Epoch 9/20

60000/60000 [=====] - 9s 144us/step - loss: 0.0973
- acc: 0.9709 - val_loss: 0.0759 - val_acc: 0.9773

Epoch 10/20

60000/60000 [=====] - 9s 142us/step - loss: 0.0921
- acc: 0.9727 - val_loss: 0.0733 - val_acc: 0.9782

Epoch 11/20

60000/60000 [=====] - 9s 150us/step - loss: 0.0845
- acc: 0.9739 - val_loss: 0.0677 - val_acc: 0.9794

Epoch 12/20

60000/60000 [=====] - 10s 162us/step - loss: 0.0815
- acc: 0.9748 - val_loss: 0.0687 - val_acc: 0.9799

Epoch 13/20

60000/60000 [=====] - 9s 148us/step - loss: 0.0783
- acc: 0.9760 - val_loss: 0.0685 - val_acc: 0.9801

Epoch 14/20

60000/60000 [=====] - 10s 159us/step - loss: 0.0727
- acc: 0.9775 - val_loss: 0.0618 - val_acc: 0.9808

Epoch 15/20

60000/60000 [=====] - 9s 153us/step - loss: 0.0645
- acc: 0.9798 - val_loss: 0.0645 - val_acc: 0.9813

Epoch 16/20

60000/60000 [=====] - 9s 151us/step - loss: 0.0656
- acc: 0.9798 - val_loss: 0.0663 - val_acc: 0.9823

Epoch 17/20

60000/60000 [=====] - 9s 156us/step - loss: 0.0613
- acc: 0.9808 - val_loss: 0.0606 - val_acc: 0.9823

Epoch 18/20

60000/60000 [=====] - 9s 152us/step - loss: 0.0582
- acc: 0.9815 - val_loss: 0.0671 - val_acc: 0.9814

Epoch 19/20

60000/60000 [=====] - 9s 145us/step - loss: 0.0568

- acc: 0.9829 - val_loss: 0.0642 - val_acc: 0.9817

Epoch 20/20

60000/60000 [=====] - 9s 152us/step - loss: 0.0535

- acc: 0.9835 - val_loss: 0.0632 - val_acc: 0.9835

In [63]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

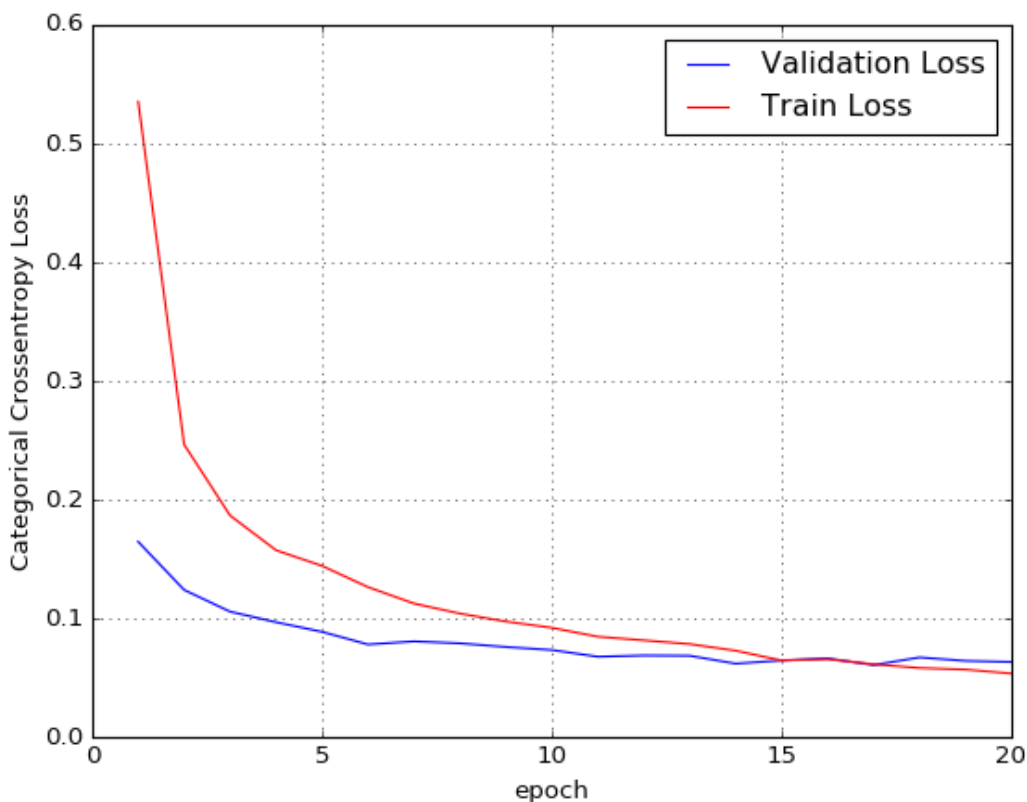
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06322702438611887

Test accuracy: 0.9835



In [64]:

```

w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

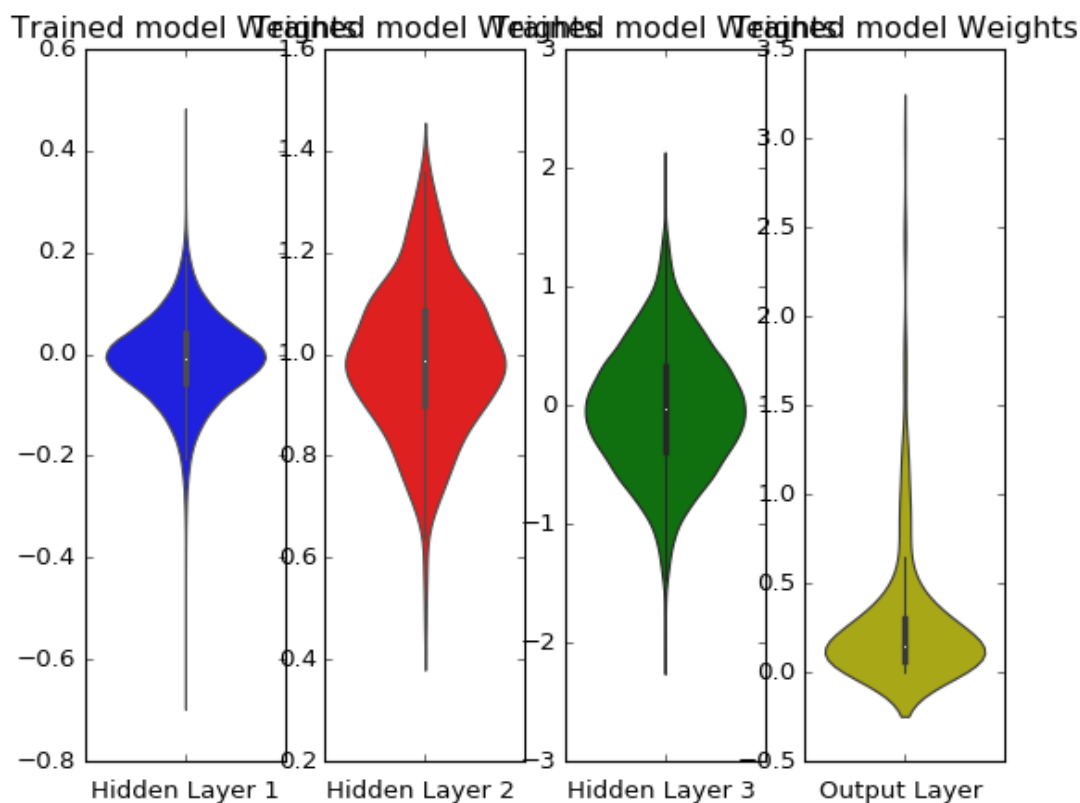
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



MLP + ReLU + ADAM with 5 layers without Dropout and Batch Normalisation

In [65]:

```

model_relu = Sequential()
model_relu.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_relu.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_relu.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_relu.add(Dense(16, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

```

Layer (type)	Output Shape	Param #
dense_43 (Dense)	(None, 256)	200960
dense_44 (Dense)	(None, 128)	32896
dense_45 (Dense)	(None, 64)	8256
dense_46 (Dense)	(None, 32)	2080
dense_47 (Dense)	(None, 16)	528
dense_48 (Dense)	(None, 10)	170
Total params: 244,890		
Trainable params: 244,890		
Non-trainable params: 0		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 70us/step - loss: 0.5227 -
 acc: 0.8537 - val_loss: 0.1814 - val_acc: 0.9456

Epoch 2/20

60000/60000 [=====] - 4s 61us/step - loss: 0.1297 -
 acc: 0.9628 - val_loss: 0.1234 - val_acc: 0.9640

Epoch 3/20

60000/60000 [=====] - 3s 47us/step - loss: 0.0845 -
 acc: 0.9744 - val_loss: 0.0956 - val_acc: 0.9720

Epoch 4/20

60000/60000 [=====] - 3s 47us/step - loss: 0.0605 -
 acc: 0.9817 - val_loss: 0.0815 - val_acc: 0.9738

Epoch 5/20

60000/60000 [=====] - 3s 44us/step - loss: 0.0440 -
 acc: 0.9865 - val_loss: 0.0836 - val_acc: 0.9751

Epoch 6/20

60000/60000 [=====] - 3s 45us/step - loss: 0.0388 -
 acc: 0.9875 - val_loss: 0.0808 - val_acc: 0.9774

Epoch 7/20

60000/60000 [=====] - 3s 44us/step - loss: 0.0312 -
 acc: 0.9898 - val_loss: 0.0765 - val_acc: 0.9802

Epoch 8/20

60000/60000 [=====] - 3s 44us/step - loss: 0.0244 -

```
acc: 0.9919 - val_loss: 0.0875 - val_acc: 0.9764
Epoch 9/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0205 -
acc: 0.9934 - val_loss: 0.0840 - val_acc: 0.9784
Epoch 10/20
60000/60000 [=====] - 3s 46us/step - loss: 0.0191 -
acc: 0.9937 - val_loss: 0.1042 - val_acc: 0.9761
Epoch 11/20
60000/60000 [=====] - 3s 45us/step - loss: 0.0200 -
acc: 0.9933 - val_loss: 0.1077 - val_acc: 0.9743
Epoch 12/20
60000/60000 [=====] - 3s 45us/step - loss: 0.0161 -
acc: 0.9947 - val_loss: 0.1308 - val_acc: 0.9662
Epoch 13/20
60000/60000 [=====] - 3s 45us/step - loss: 0.0145 -
acc: 0.9954 - val_loss: 0.1026 - val_acc: 0.9762
Epoch 14/20
60000/60000 [=====] - 3s 47us/step - loss: 0.0125 -
acc: 0.9960 - val_loss: 0.1330 - val_acc: 0.9746
Epoch 15/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0139 -
acc: 0.9951 - val_loss: 0.1005 - val_acc: 0.9766
Epoch 16/20
60000/60000 [=====] - 3s 43us/step - loss: 0.0105 -
acc: 0.9965 - val_loss: 0.1085 - val_acc: 0.9771
Epoch 17/20
60000/60000 [=====] - 3s 43us/step - loss: 0.0136 -
acc: 0.9957 - val_loss: 0.1148 - val_acc: 0.9746
Epoch 18/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0118 -
acc: 0.9963 - val_loss: 0.0982 - val_acc: 0.9776
Epoch 19/20
60000/60000 [=====] - 3s 45us/step - loss: 0.0096 -
acc: 0.9968 - val_loss: 0.0994 - val_acc: 0.9792
Epoch 20/20
60000/60000 [=====] - 3s 47us/step - loss: 0.0074 -
acc: 0.9976 - val_loss: 0.1014 - val_acc: 0.9809
```

In [67]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

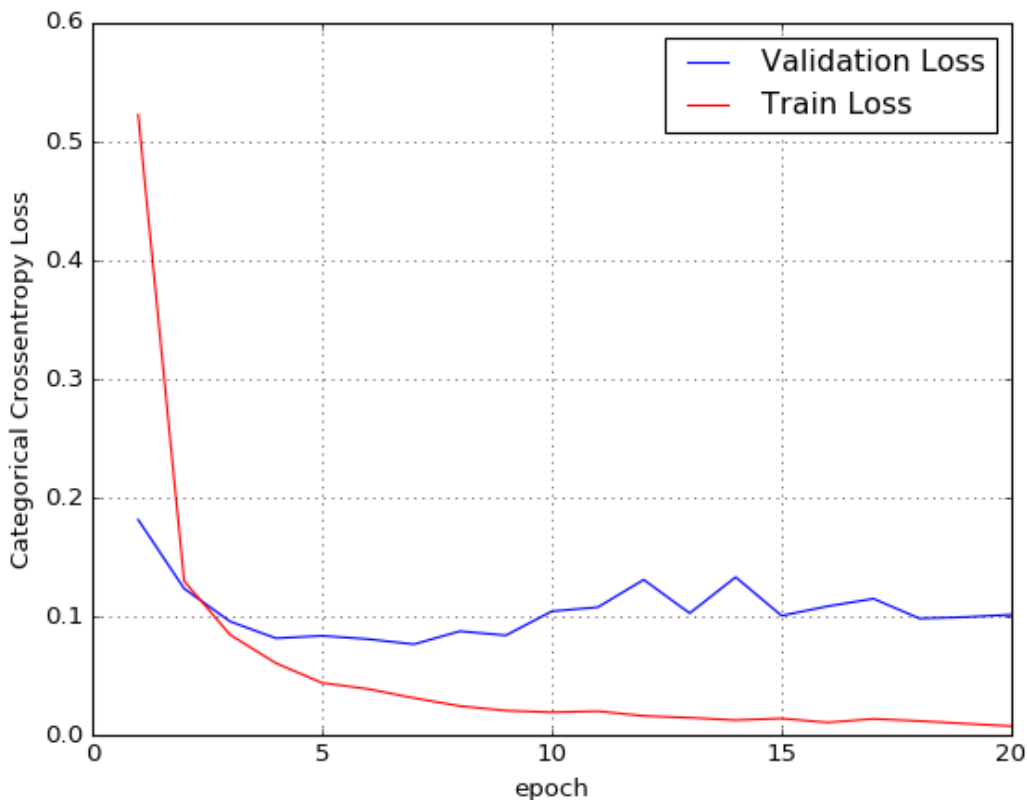
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# Loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10143251653150177

Test accuracy: 0.9809



In [68]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

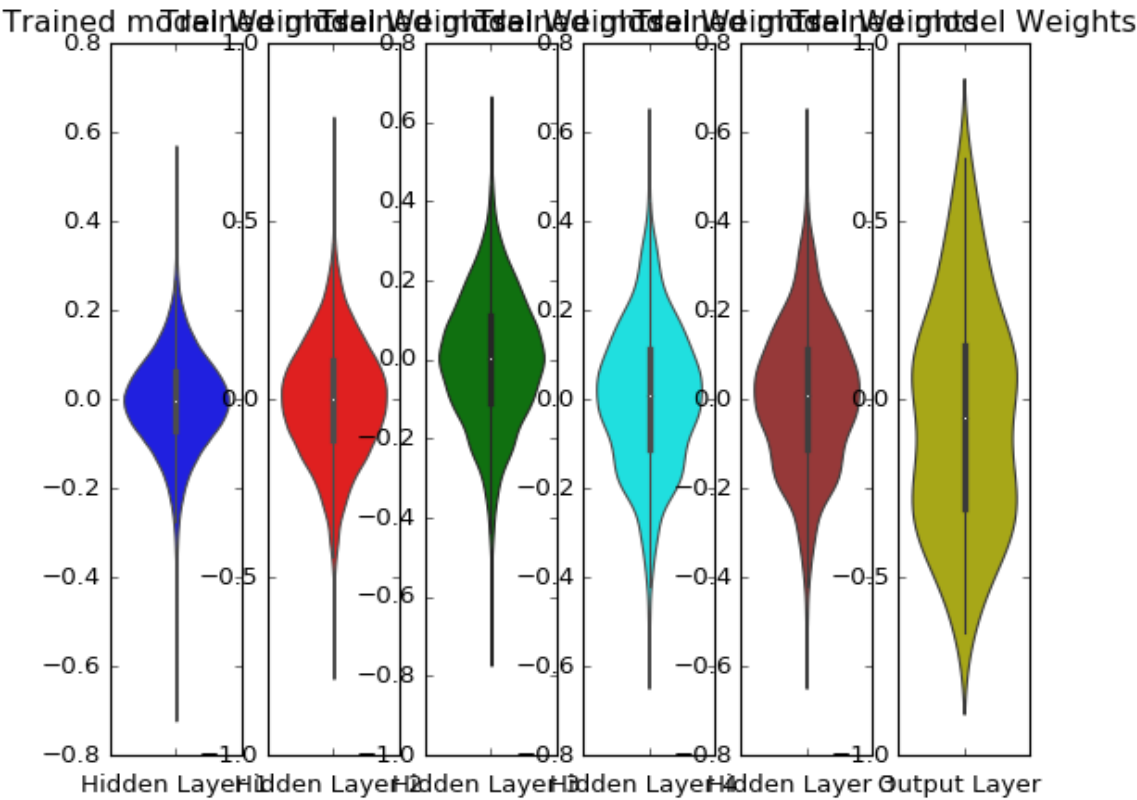
plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='cyan')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='brown')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Batch-Norm on 5 hidden Layers + AdamOptimizer

In [69]:

```

# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i + n_{i+1})}$ 
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(\theta, \sigma) = N(\theta, 0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(\theta, \sigma) = N(\theta, 0.055)$ 
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(\theta, \sigma) = N(\theta, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, std=0.039)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(132, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std=0.055)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std=0.120)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std=0.120)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(16, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std=0.120)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()

```

Layer (type)	Output Shape	Param #
dense_49 (Dense)	(None, 256)	200960
batch_normalization_16 (Batch Normalization)	(None, 256)	1024
dense_50 (Dense)	(None, 132)	33924
batch_normalization_17 (Batch Normalization)	(None, 132)	528
dense_51 (Dense)	(None, 64)	8512
batch_normalization_18 (Batch Normalization)	(None, 64)	256
dense_52 (Dense)	(None, 32)	2080
batch_normalization_19 (Batch Normalization)	(None, 32)	128
dense_53 (Dense)	(None, 16)	528
batch_normalization_20 (Batch Normalization)	(None, 16)	64
dense_54 (Dense)	(None, 10)	170
Total params: 248,174		

Trainable params: 247,174

Non-trainable params: 1,000

In [70]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 99us/step - loss: 0.3940 - acc: 0.9010 - val_loss: 0.1705 - val_acc: 0.9525

Epoch 2/20

60000/60000 [=====] - 4s 72us/step - loss: 0.1194 - acc: 0.9662 - val_loss: 0.1321 - val_acc: 0.9620

Epoch 3/20

60000/60000 [=====] - 4s 71us/step - loss: 0.0810 - acc: 0.9761 - val_loss: 0.1025 - val_acc: 0.9690

Epoch 4/20

60000/60000 [=====] - 4s 70us/step - loss: 0.0632 - acc: 0.9812 - val_loss: 0.0927 - val_acc: 0.9718

Epoch 5/20

60000/60000 [=====] - 4s 71us/step - loss: 0.0488 - acc: 0.9851 - val_loss: 0.0833 - val_acc: 0.9759

Epoch 6/20

60000/60000 [=====] - 4s 72us/step - loss: 0.0433 - acc: 0.9864 - val_loss: 0.0792 - val_acc: 0.9765

Epoch 7/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0349 - acc: 0.9892 - val_loss: 0.0881 - val_acc: 0.9750

Epoch 8/20

60000/60000 [=====] - 5s 78us/step - loss: 0.0321 - acc: 0.9900 - val_loss: 0.0759 - val_acc: 0.9786

Epoch 9/20

60000/60000 [=====] - 4s 69us/step - loss: 0.0287 - acc: 0.9909 - val_loss: 0.0840 - val_acc: 0.9771

Epoch 10/20

60000/60000 [=====] - 4s 70us/step - loss: 0.0261 - acc: 0.9915 - val_loss: 0.0810 - val_acc: 0.9778

Epoch 11/20

60000/60000 [=====] - 4s 72us/step - loss: 0.0237 - acc: 0.9920 - val_loss: 0.0764 - val_acc: 0.9797

Epoch 12/20

60000/60000 [=====] - 5s 78us/step - loss: 0.0234 - acc: 0.9924 - val_loss: 0.0791 - val_acc: 0.9785

Epoch 13/20

60000/60000 [=====] - 4s 71us/step - loss: 0.0198 - acc: 0.9934 - val_loss: 0.0763 - val_acc: 0.9801

Epoch 14/20

60000/60000 [=====] - 5s 75us/step - loss: 0.0177 - acc: 0.9942 - val_loss: 0.0852 - val_acc: 0.9785

Epoch 15/20

60000/60000 [=====] - 4s 71us/step - loss: 0.0176 - acc: 0.9940 - val_loss: 0.0802 - val_acc: 0.9781

Epoch 16/20

60000/60000 [=====] - 4s 72us/step - loss: 0.0162 - acc: 0.9945 - val_loss: 0.0783 - val_acc: 0.9801

Epoch 17/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0130 - acc: 0.9955 - val_loss: 0.0743 - val_acc: 0.9811

Epoch 18/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0138 - acc: 0.9954 - val_loss: 0.0809 - val_acc: 0.9813

Epoch 19/20

60000/60000 [=====] - 4s 74us/step - loss: 0.0126 -

acc: 0.9958 - val_loss: 0.0898 - val_acc: 0.9785

Epoch 20/20

60000/60000 [=====] - 4s 71us/step - loss: 0.0119 -

acc: 0.9964 - val_loss: 0.0782 - val_acc: 0.9812

In [71]:

```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

Test score: 0.07821065115529928

Test accuracy: 0.9812

In [72]:

```
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

In [73]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

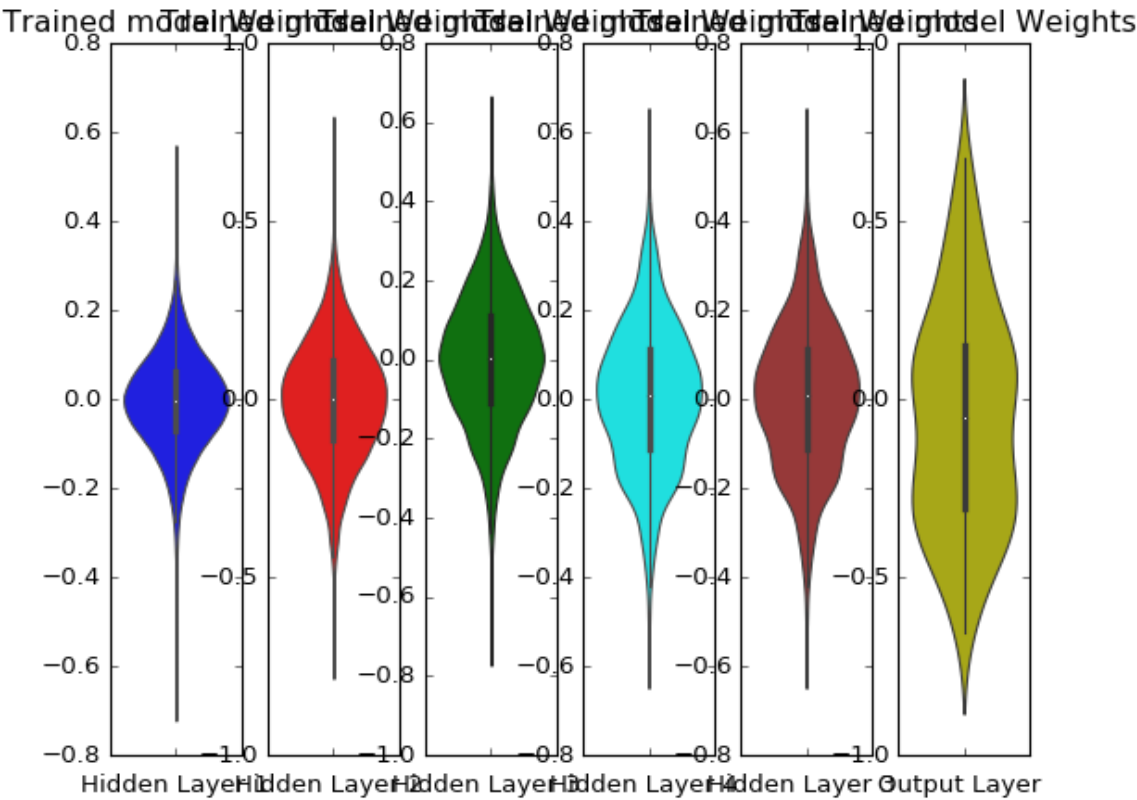
plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='cyan')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='brown')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Dropout(rate = 0.5) + AdamOptimizer with 5 hidden layers

In [74]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-funct
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(16, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_55 (Dense)	(None, 256)	200960
dropout_21 (Dropout)	(None, 256)	0
dense_56 (Dense)	(None, 128)	32896
dropout_22 (Dropout)	(None, 128)	0
dense_57 (Dense)	(None, 64)	8256
dropout_23 (Dropout)	(None, 64)	0
dense_58 (Dense)	(None, 32)	2080
dropout_24 (Dropout)	(None, 32)	0
dense_59 (Dense)	(None, 16)	528
dropout_25 (Dropout)	(None, 16)	0
dense_60 (Dense)	(None, 10)	170
Total params: 244,890		
Trainable params: 244,890		
Non-trainable params: 0		

In [75]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 77us/step - loss: 2.0919 - acc: 0.2437 - val_loss: 1.4719 - val_acc: 0.5675

Epoch 2/20

60000/60000 [=====] - 5s 76us/step - loss: 1.5121 - acc: 0.4698 - val_loss: 1.0775 - val_acc: 0.6651

Epoch 3/20

60000/60000 [=====] - 4s 65us/step - loss: 1.2592 - acc: 0.5563 - val_loss: 0.9283 - val_acc: 0.7135

Epoch 4/20

60000/60000 [=====] - 4s 60us/step - loss: 1.1274 - acc: 0.5964 - val_loss: 0.8500 - val_acc: 0.7431

Epoch 5/20

60000/60000 [=====] - 4s 64us/step - loss: 1.0385 - acc: 0.6244 - val_loss: 0.7729 - val_acc: 0.7404

Epoch 6/20

60000/60000 [=====] - 4s 69us/step - loss: 0.9854 - acc: 0.6372 - val_loss: 0.7186 - val_acc: 0.7485

Epoch 7/20

60000/60000 [=====] - 4s 62us/step - loss: 0.9358 - acc: 0.6564 - val_loss: 0.7074 - val_acc: 0.7570

Epoch 8/20

60000/60000 [=====] - 4s 61us/step - loss: 0.9070 - acc: 0.6706 - val_loss: 0.6708 - val_acc: 0.7626

Epoch 9/20

60000/60000 [=====] - 4s 64us/step - loss: 0.8758 - acc: 0.6771 - val_loss: 0.6378 - val_acc: 0.7682

Epoch 10/20

60000/60000 [=====] - 4s 68us/step - loss: 0.8545 - acc: 0.6867 - val_loss: 0.6129 - val_acc: 0.7741

Epoch 11/20

60000/60000 [=====] - 4s 61us/step - loss: 0.8247 - acc: 0.6961 - val_loss: 0.5988 - val_acc: 0.7782

Epoch 12/20

60000/60000 [=====] - 4s 70us/step - loss: 0.8060 - acc: 0.7010 - val_loss: 0.5926 - val_acc: 0.7778

Epoch 13/20

60000/60000 [=====] - 4s 60us/step - loss: 0.7989 - acc: 0.7045 - val_loss: 0.5832 - val_acc: 0.7754

Epoch 14/20

60000/60000 [=====] - 4s 63us/step - loss: 0.7900 - acc: 0.7084 - val_loss: 0.5665 - val_acc: 0.7786

Epoch 15/20

60000/60000 [=====] - 4s 69us/step - loss: 0.7791 - acc: 0.7119 - val_loss: 0.5853 - val_acc: 0.7789

Epoch 16/20

60000/60000 [=====] - 4s 65us/step - loss: 0.7677 - acc: 0.7154 - val_loss: 0.5975 - val_acc: 0.7814

Epoch 17/20

60000/60000 [=====] - 4s 64us/step - loss: 0.7564 - acc: 0.7220 - val_loss: 0.5574 - val_acc: 0.7841

Epoch 18/20

60000/60000 [=====] - 4s 70us/step - loss: 0.7600 - acc: 0.7284 - val_loss: 0.5639 - val_acc: 0.7812

Epoch 19/20

60000/60000 [=====] - 4s 65us/step - loss: 0.7450 -

acc: 0.7352 - val_loss: 0.5385 - val_acc: 0.8084

Epoch 20/20

60000/60000 [=====] - 4s 64us/step - loss: 0.7221 -

acc: 0.7454 - val_loss: 0.5309 - val_acc: 0.8095

In [76]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

Test score: 0.5309434900760651

Test accuracy: 0.8095

In [77]:

```
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

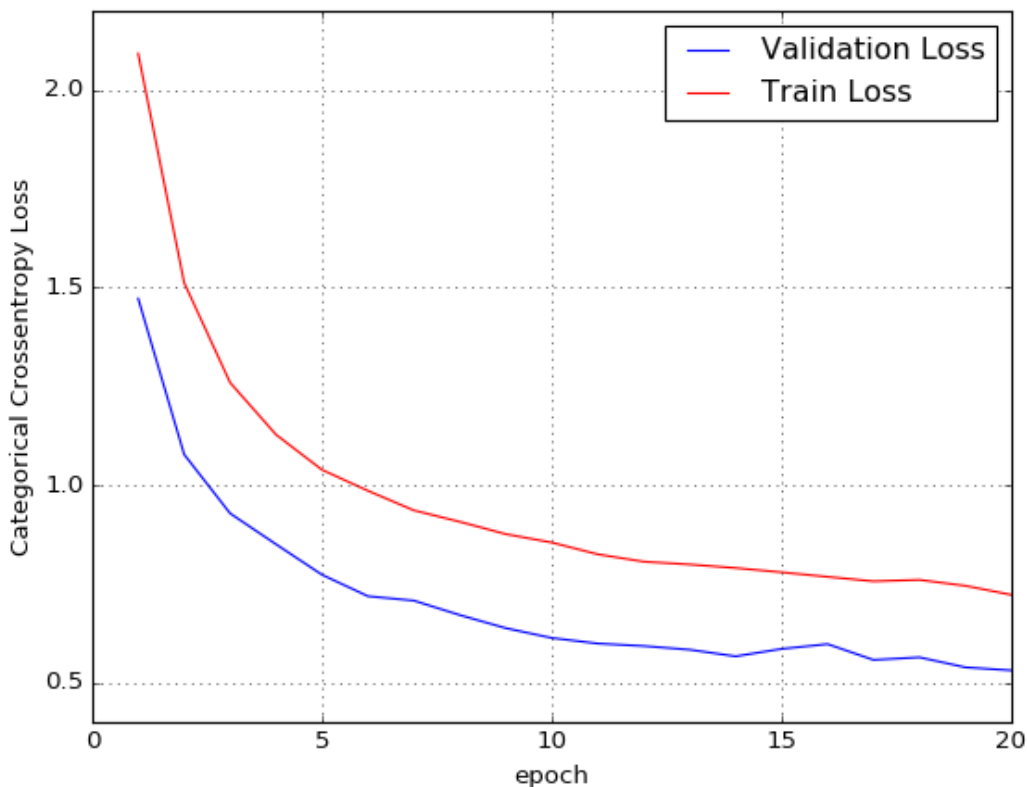
# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```



In [78]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

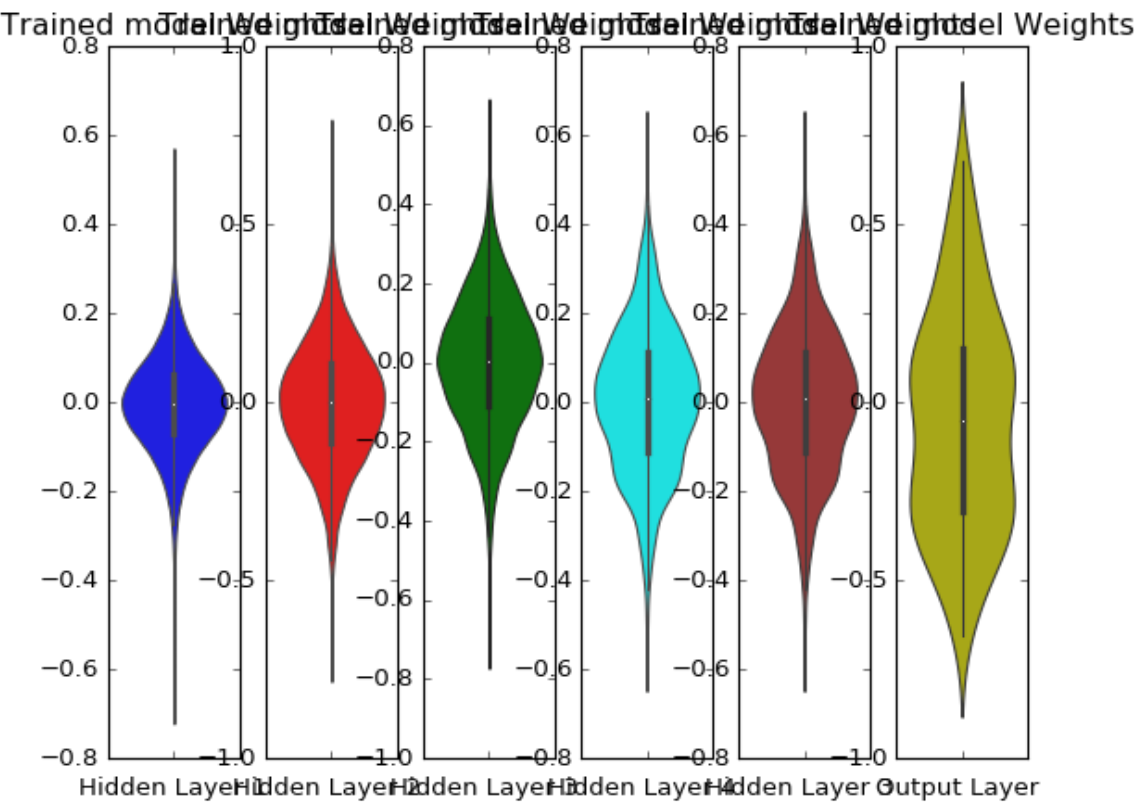
plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='cyan')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='brown')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

MLP + Dropout(rate = 0.3) + AdamOptimizer with 5 hidden layers

In [79]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-funct
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.3))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.3))

model_drop.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.3))

model_drop.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.3))

model_drop.add(Dense(16, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.3))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_61 (Dense)	(None, 256)	200960
dropout_26 (Dropout)	(None, 256)	0
dense_62 (Dense)	(None, 128)	32896
dropout_27 (Dropout)	(None, 128)	0
dense_63 (Dense)	(None, 64)	8256
dropout_28 (Dropout)	(None, 64)	0
dense_64 (Dense)	(None, 32)	2080
dropout_29 (Dropout)	(None, 32)	0
dense_65 (Dense)	(None, 16)	528
dropout_30 (Dropout)	(None, 16)	0
dense_66 (Dense)	(None, 10)	170
Total params: 244,890		
Trainable params: 244,890		
Non-trainable params: 0		

In [80]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 79us/step - loss: 1.2953 - acc: 0.5550 - val_loss: 0.3444 - val_acc: 0.9196

Epoch 2/20

60000/60000 [=====] - 4s 62us/step - loss: 0.6133 - acc: 0.8158 - val_loss: 0.2037 - val_acc: 0.9518

Epoch 3/20

60000/60000 [=====] - 4s 68us/step - loss: 0.4686 - acc: 0.8674 - val_loss: 0.1706 - val_acc: 0.9584

Epoch 4/20

60000/60000 [=====] - 4s 69us/step - loss: 0.3907 - acc: 0.8955 - val_loss: 0.1666 - val_acc: 0.9600

Epoch 5/20

60000/60000 [=====] - 4s 73us/step - loss: 0.3488 - acc: 0.9061 - val_loss: 0.1574 - val_acc: 0.9633

Epoch 6/20

60000/60000 [=====] - 4s 68us/step - loss: 0.3226 - acc: 0.9148 - val_loss: 0.1406 - val_acc: 0.9700

Epoch 7/20

60000/60000 [=====] - 4s 60us/step - loss: 0.2941 - acc: 0.9220 - val_loss: 0.1450 - val_acc: 0.9679

Epoch 8/20

60000/60000 [=====] - 4s 62us/step - loss: 0.2711 - acc: 0.9304 - val_loss: 0.1344 - val_acc: 0.9682

Epoch 9/20

60000/60000 [=====] - 4s 63us/step - loss: 0.2582 - acc: 0.9318 - val_loss: 0.1334 - val_acc: 0.9721

Epoch 10/20

60000/60000 [=====] - 4s 65us/step - loss: 0.2399 - acc: 0.9374 - val_loss: 0.1381 - val_acc: 0.9713

Epoch 11/20

60000/60000 [=====] - 4s 67us/step - loss: 0.2325 - acc: 0.9389 - val_loss: 0.1267 - val_acc: 0.9736

Epoch 12/20

60000/60000 [=====] - 4s 61us/step - loss: 0.2199 - acc: 0.9426 - val_loss: 0.1407 - val_acc: 0.9708

Epoch 13/20

60000/60000 [=====] - 4s 64us/step - loss: 0.2073 - acc: 0.9456 - val_loss: 0.1249 - val_acc: 0.9734

Epoch 14/20

60000/60000 [=====] - 4s 67us/step - loss: 0.2039 - acc: 0.9464 - val_loss: 0.1340 - val_acc: 0.9731

Epoch 15/20

60000/60000 [=====] - 4s 72us/step - loss: 0.2019 - acc: 0.9470 - val_loss: 0.1302 - val_acc: 0.9746

Epoch 16/20

60000/60000 [=====] - 5s 75us/step - loss: 0.1932 - acc: 0.9503 - val_loss: 0.1149 - val_acc: 0.9756

Epoch 17/20

60000/60000 [=====] - 4s 73us/step - loss: 0.1823 - acc: 0.9527 - val_loss: 0.1123 - val_acc: 0.9785

Epoch 18/20

60000/60000 [=====] - 4s 74us/step - loss: 0.1789 - acc: 0.9540 - val_loss: 0.1136 - val_acc: 0.9766

Epoch 19/20

60000/60000 [=====] - 4s 66us/step - loss: 0.1710 -

acc: 0.9562 - val_loss: 0.1113 - val_acc: 0.9788

Epoch 20/20

60000/60000 [=====] - 4s 70us/step - loss: 0.1696 -

acc: 0.9562 - val_loss: 0.1158 - val_acc: 0.9794

In [81]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

Test score: 0.1157619362520054

Test accuracy: 0.9794

In [82]:

```
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

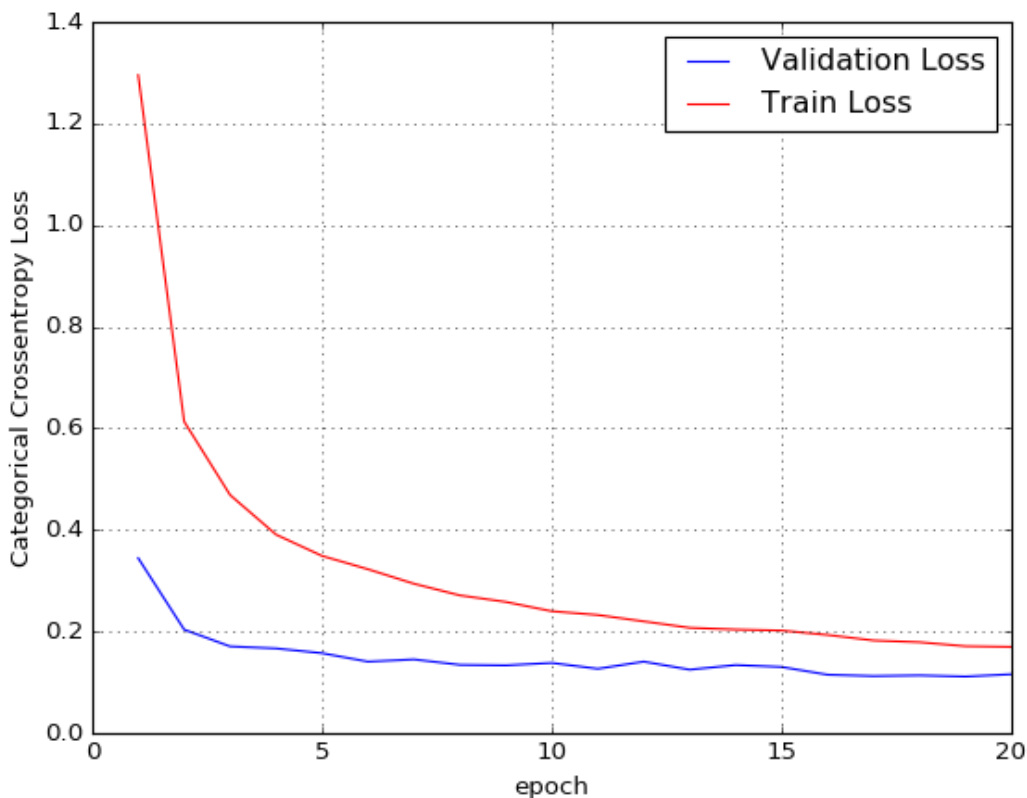
# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```



In [83]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

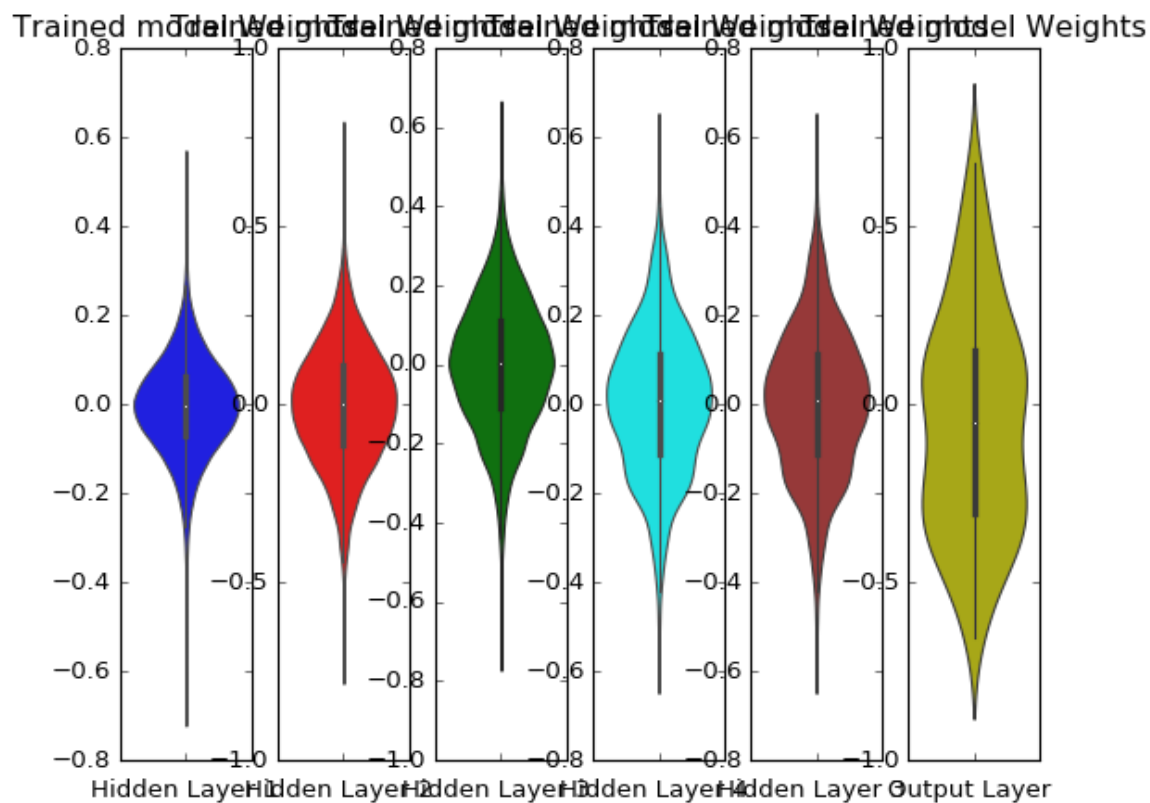
plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='cyan')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='brown')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



**MLP + BatchNormalization + Dropout(rate = 0.5) + AdamOptimizer
with 3 hidden layers**

In [84]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-funct
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=R
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stdde
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stdde
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stdde
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(16, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stdde
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_67 (Dense)	(None, 256)	200960
batch_normalization_21 (Batch Normalization)	(None, 256)	1024
dropout_31 (Dropout)	(None, 256)	0
dense_68 (Dense)	(None, 128)	32896
batch_normalization_22 (Batch Normalization)	(None, 128)	512
dropout_32 (Dropout)	(None, 128)	0
dense_69 (Dense)	(None, 64)	8256
batch_normalization_23 (Batch Normalization)	(None, 64)	256
dropout_33 (Dropout)	(None, 64)	0
dense_70 (Dense)	(None, 32)	2080
batch_normalization_24 (Batch Normalization)	(None, 32)	128
dropout_34 (Dropout)	(None, 32)	0

dense_71 (Dense)	(None, 16)	528
batch_normalization_25 (Batch Normalization)	(None, 16)	64
dropout_35 (Dropout)	(None, 16)	0
dense_72 (Dense)	(None, 10)	170
=====		
Total params: 246,874		
Trainable params: 245,882		
Non-trainable params: 992		

In [85]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 121us/step - loss: 2.2357
- acc: 0.2311 - val_loss: 1.3396 - val_acc: 0.7074

Epoch 2/20

60000/60000 [=====] - 5s 83us/step - loss: 1.5868 -
acc: 0.4290 - val_loss: 0.8430 - val_acc: 0.7752

Epoch 3/20

60000/60000 [=====] - 5s 83us/step - loss: 1.2666 -
acc: 0.5369 - val_loss: 0.6421 - val_acc: 0.8368

Epoch 4/20

60000/60000 [=====] - 5s 87us/step - loss: 1.0828 -
acc: 0.6060 - val_loss: 0.5345 - val_acc: 0.8755

Epoch 5/20

60000/60000 [=====] - 6s 96us/step - loss: 0.9650 -
acc: 0.6517 - val_loss: 0.4411 - val_acc: 0.9000

Epoch 6/20

60000/60000 [=====] - 5s 89us/step - loss: 0.8813 -
acc: 0.6906 - val_loss: 0.3809 - val_acc: 0.9190

Epoch 7/20

60000/60000 [=====] - 5s 85us/step - loss: 0.7947 -
acc: 0.7292 - val_loss: 0.3086 - val_acc: 0.9374

Epoch 8/20

60000/60000 [=====] - 5s 89us/step - loss: 0.7309 -
acc: 0.7568 - val_loss: 0.2643 - val_acc: 0.9387

Epoch 9/20

60000/60000 [=====] - 5s 87us/step - loss: 0.6728 -
acc: 0.7874 - val_loss: 0.2188 - val_acc: 0.9487

Epoch 10/20

60000/60000 [=====] - 5s 90us/step - loss: 0.6153 -
acc: 0.8104 - val_loss: 0.2023 - val_acc: 0.9529

Epoch 11/20

60000/60000 [=====] - 5s 85us/step - loss: 0.5722 -
acc: 0.8287 - val_loss: 0.1780 - val_acc: 0.9557

Epoch 12/20

60000/60000 [=====] - 5s 85us/step - loss: 0.5389 -
acc: 0.8404 - val_loss: 0.1658 - val_acc: 0.9583

Epoch 13/20

60000/60000 [=====] - 5s 90us/step - loss: 0.5108 -
acc: 0.8528 - val_loss: 0.1583 - val_acc: 0.9609

Epoch 14/20

60000/60000 [=====] - 5s 86us/step - loss: 0.4867 -
acc: 0.8607 - val_loss: 0.1645 - val_acc: 0.9599

Epoch 15/20

60000/60000 [=====] - 5s 92us/step - loss: 0.4716 -
acc: 0.8661 - val_loss: 0.1575 - val_acc: 0.9626

Epoch 16/20

60000/60000 [=====] - 5s 87us/step - loss: 0.4561 -
acc: 0.8706 - val_loss: 0.1447 - val_acc: 0.9651

Epoch 17/20

60000/60000 [=====] - 5s 82us/step - loss: 0.4360 -
acc: 0.8786 - val_loss: 0.1449 - val_acc: 0.9658

Epoch 18/20

60000/60000 [=====] - 5s 84us/step - loss: 0.4191 -
acc: 0.8827 - val_loss: 0.1466 - val_acc: 0.9661

Epoch 19/20

60000/60000 [=====] - 5s 82us/step - loss: 0.4146 -

acc: 0.8872 - val_loss: 0.1389 - val_acc: 0.9674

Epoch 20/20

60000/60000 [=====] - 5s 82us/step - loss: 0.3964 -

acc: 0.8903 - val_loss: 0.1381 - val_acc: 0.9688

In [86]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

Test score: 0.13805706812888383

Test accuracy: 0.9688

In [87]:

```
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

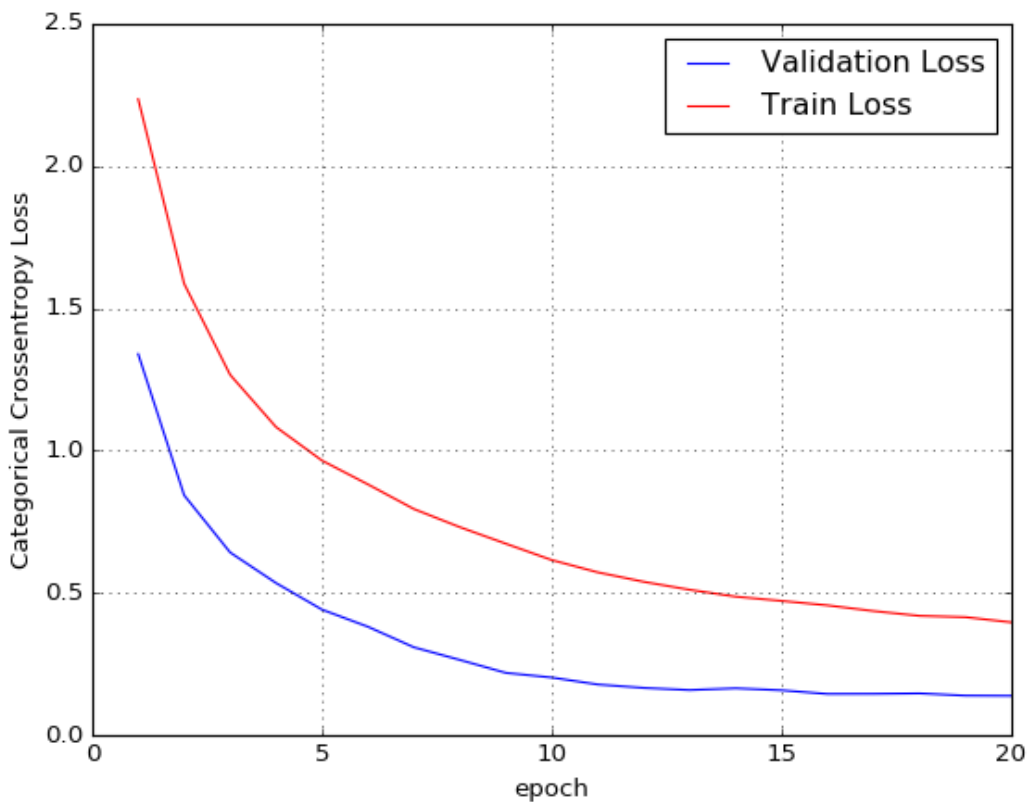
# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```



In [88]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

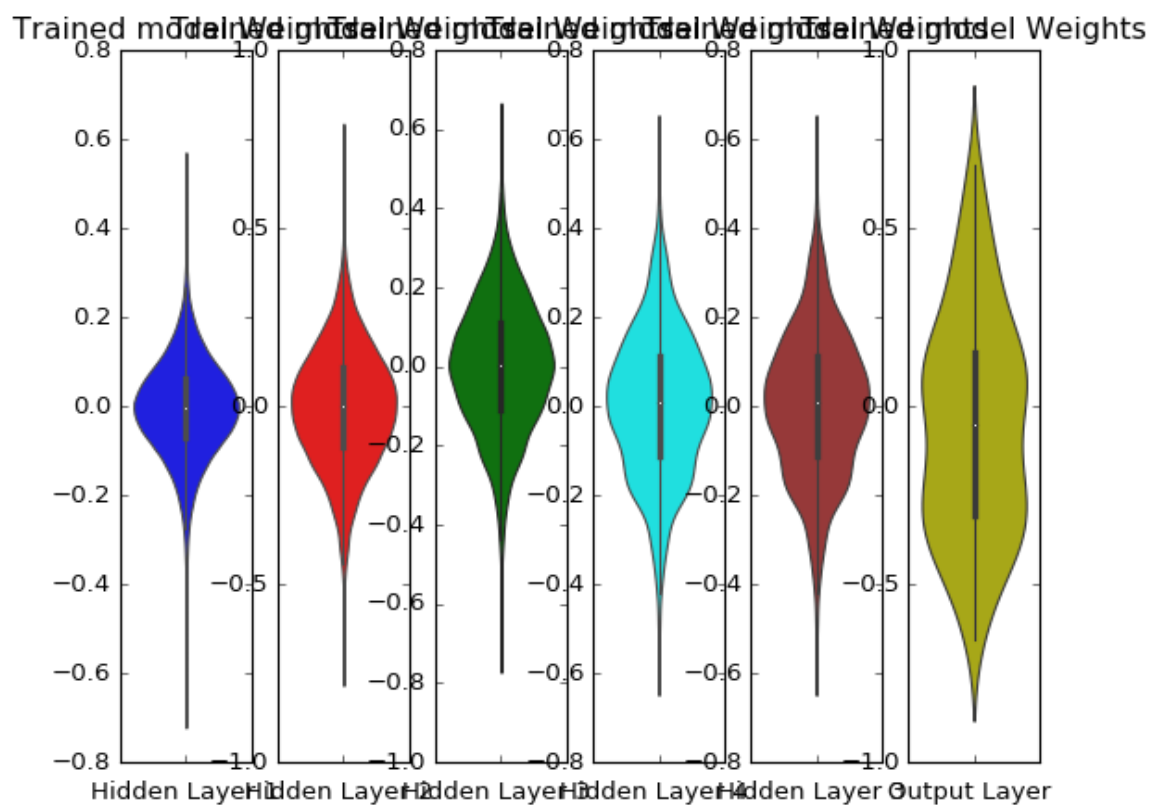
plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='cyan')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='brown')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



**MLP + BatchNormalization + Dropout(rate = 0.3) + AdamOptimizer
with 3 hidden layers**

In [89]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-funct
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=R
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.3))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stdde
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.3))

model_drop.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stdde
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.3))

model_drop.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stdde
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.3))

model_drop.add(Dense(16, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stdde
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.3))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_73 (Dense)	(None, 256)	200960
batch_normalization_26 (Batch Normalization)	(None, 256)	1024
dropout_36 (Dropout)	(None, 256)	0
dense_74 (Dense)	(None, 128)	32896
batch_normalization_27 (Batch Normalization)	(None, 128)	512
dropout_37 (Dropout)	(None, 128)	0
dense_75 (Dense)	(None, 64)	8256
batch_normalization_28 (Batch Normalization)	(None, 64)	256
dropout_38 (Dropout)	(None, 64)	0
dense_76 (Dense)	(None, 32)	2080
batch_normalization_29 (Batch Normalization)	(None, 32)	128
dropout_39 (Dropout)	(None, 32)	0

dense_77 (Dense)	(None, 16)	528
batch_normalization_30 (Batch Normalization)	(None, 16)	64
dropout_40 (Dropout)	(None, 16)	0
dense_78 (Dense)	(None, 10)	170
=====		
Total params: 246,874		
Trainable params: 245,882		
Non-trainable params: 992		

In [90]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 127us/step - loss: 1.3488
- acc: 0.5626 - val_loss: 0.3823 - val_acc: 0.9094

Epoch 2/20

60000/60000 [=====] - 5s 87us/step - loss: 0.6504 -
acc: 0.8141 - val_loss: 0.2155 - val_acc: 0.9436

Epoch 3/20

60000/60000 [=====] - 5s 84us/step - loss: 0.4596 -
acc: 0.8756 - val_loss: 0.1916 - val_acc: 0.9499

Epoch 4/20

60000/60000 [=====] - 5s 84us/step - loss: 0.3727 -
acc: 0.9042 - val_loss: 0.1556 - val_acc: 0.9577

Epoch 5/20

60000/60000 [=====] - 5s 86us/step - loss: 0.3184 -
acc: 0.9225 - val_loss: 0.1419 - val_acc: 0.9614

Epoch 6/20

60000/60000 [=====] - 5s 85us/step - loss: 0.2840 -
acc: 0.9301 - val_loss: 0.1332 - val_acc: 0.9652

Epoch 7/20

60000/60000 [=====] - 5s 84us/step - loss: 0.2605 -
acc: 0.9364 - val_loss: 0.1287 - val_acc: 0.9680

Epoch 8/20

60000/60000 [=====] - 5s 85us/step - loss: 0.2333 -
acc: 0.9447 - val_loss: 0.1204 - val_acc: 0.9717

Epoch 9/20

60000/60000 [=====] - 5s 86us/step - loss: 0.2152 -
acc: 0.9466 - val_loss: 0.1132 - val_acc: 0.9721

Epoch 10/20

60000/60000 [=====] - 5s 85us/step - loss: 0.2051 -
acc: 0.9512 - val_loss: 0.1213 - val_acc: 0.9713

Epoch 11/20

60000/60000 [=====] - 5s 84us/step - loss: 0.1973 -
acc: 0.9532 - val_loss: 0.1044 - val_acc: 0.9735

Epoch 12/20

60000/60000 [=====] - 5s 86us/step - loss: 0.1844 -
acc: 0.9558 - val_loss: 0.1013 - val_acc: 0.9764

Epoch 13/20

60000/60000 [=====] - 5s 89us/step - loss: 0.1732 -
acc: 0.9595 - val_loss: 0.0936 - val_acc: 0.9770

Epoch 14/20

60000/60000 [=====] - 5s 88us/step - loss: 0.1682 -
acc: 0.9610 - val_loss: 0.0920 - val_acc: 0.9782

Epoch 15/20

60000/60000 [=====] - 6s 92us/step - loss: 0.1559 -
acc: 0.9627 - val_loss: 0.0934 - val_acc: 0.9777

Epoch 16/20

60000/60000 [=====] - 7s 113us/step - loss: 0.1485
- acc: 0.9651 - val_loss: 0.0998 - val_acc: 0.9765

Epoch 17/20

60000/60000 [=====] - 5s 91us/step - loss: 0.1447 -
acc: 0.9668 - val_loss: 0.1004 - val_acc: 0.9765

Epoch 18/20

60000/60000 [=====] - 6s 93us/step - loss: 0.1361 -
acc: 0.9673 - val_loss: 0.0935 - val_acc: 0.9793

Epoch 19/20

60000/60000 [=====] - 6s 98us/step - loss: 0.1354 -

acc: 0.9694 - val_loss: 0.0916 - val_acc: 0.9789

Epoch 20/20

60000/60000 [=====] - 6s 100us/step - loss: 0.1322

- acc: 0.9685 - val_loss: 0.0917 - val_acc: 0.9776

In [91]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

Test score: 0.09172347799413838

Test accuracy: 0.9776

In [92]:

```
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

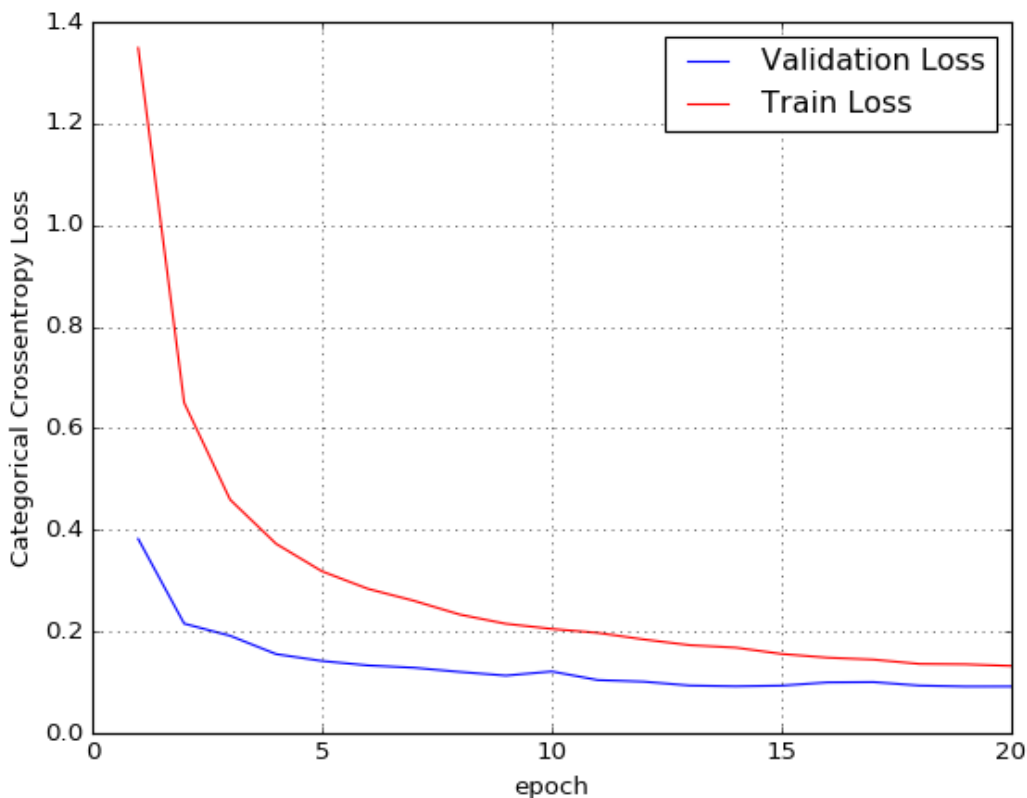
# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```



In [93]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

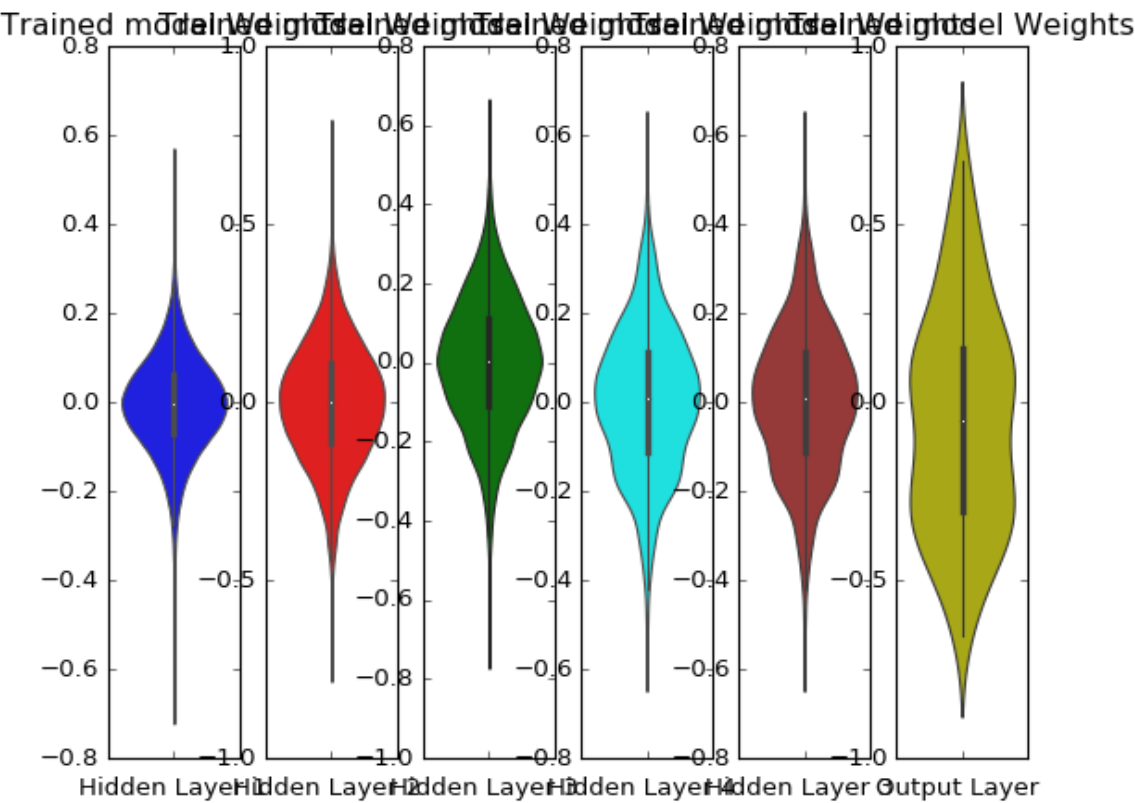
plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='cyan')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='brown')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Comparisons

In [100]:

```
from prettytable import PrettyTable

x = PrettyTable()
x.field_names = ["Architecture", "parameters", "Accuracy"]

x.add_row(["2 layer", "without Dropout and Batch Normalization", 97.9])
x.add_row(["2 layer", "with Dropuot(rate = 0.5)", 97.44])
x.add_row(["2 layer", "with Dropuot(rate = 0.3)", 98.02])
x.add_row(["2 layer", "with Batch Normalization", 97.98])
x.add_row(["2 layer", "with Dropuot(rate = 0.5) and Batch Normalization ", 98.08])
x.add_row(["2 layer", "with Dropuot(rate = 0.3) and Batch Normalization ", 98.25])

x.add_row(["3 layer", "without Dropout and Batch Normalization", 97.87])
x.add_row(["3 layer", "with Dropuot", 97.38])
x.add_row(["3 layer", "with Dropuot(rate = 0.3)", 98.06])
x.add_row(["3 layer", "with Batch Normalization", 97.59])
x.add_row(["3 layer", "with Dropuot(rate = 0.5) and Batch Normalization ", 97.84])
x.add_row(["3 layer", "with Dropuot(rate = 0.3) and Batch Normalization ", 98.35])

x.add_row(["5 layer", "without Dropout and Batch Normalization", 98.09])
x.add_row(["5 layer", "with Dropuot(rate = 0.5)", 85.8])
x.add_row(["5 layer", "with Dropuot(rate = 0.3)", 97.54])
x.add_row(["5 layer", "with Batch Normalization", 98.14])
x.add_row(["5 layer", "with Dropuot(rate = 0.5) and Batch Normalization ", 95.9])
x.add_row(["5 layer", "with Dropuot(rate = 0.3) and Batch Normalization ", 97.76])

print(x)
```

```
+-----+-----+-----+
--+
| Architecture |           parameters           | Accuracy |
+-----+-----+-----+
--+
| 2 layer      | without Dropout and Batch Normalization | 97.9      |
| 2 layer      |           with Dropuot(rate = 0.5)      | 97.44     |
| 2 layer      |           with Dropuot(rate = 0.3)      | 98.02     |
| 2 layer      |           with Batch Normalization      | 97.98     |
| 2 layer      | with Dropuot(rate = 0.5) and Batch Normalization | 98.08     |
| 2 layer      | with Dropuot(rate = 0.3) and Batch Normalization | 98.25     |
| 3 layer      | without Dropout and Batch Normalization | 97.87     |
| 3 layer      |           with Dropuot                  | 97.38     |
| 3 layer      |           with Dropuot(rate = 0.3)      | 98.06     |
| 3 layer      |           with Batch Normalization      | 97.59     |
| 3 layer      | with Dropuot(rate = 0.5) and Batch Normalization | 97.84     |
| 3 layer      | with Dropuot(rate = 0.3) and Batch Normalization | 98.35     |
| 5 layer      | without Dropout and Batch Normalization | 98.09     |
```

5 layer		with Dropout(rate = 0.5)		85.8
5 layer		with Dropout(rate = 0.3)		97.54
5 layer		with Batch Normalization		98.14
5 layer		with Dropout(rate = 0.5) and Batch Normalization		95.9
5 layer		with Dropout(rate = 0.3) and Batch Normalization		97.76
+-----+-----+-----+-----+-----+				
--+				

Procedure followed

1. Flattened the 28*28 dimensional MNIST data to 784
2. Normalized the data
3. Used a softmax classifier of output dimensions = 10
4. Created multiple models in Keras with various parameter combinations like activation function = 'relu', optimizer = 'Adam', with/without dropout of different rates, with/without Batch normalization
5. Plotted the epoch vs Train/Test loss of each model
6. Plotted the weights of each model